

# What does blurring remove?



Now let's add it back.....

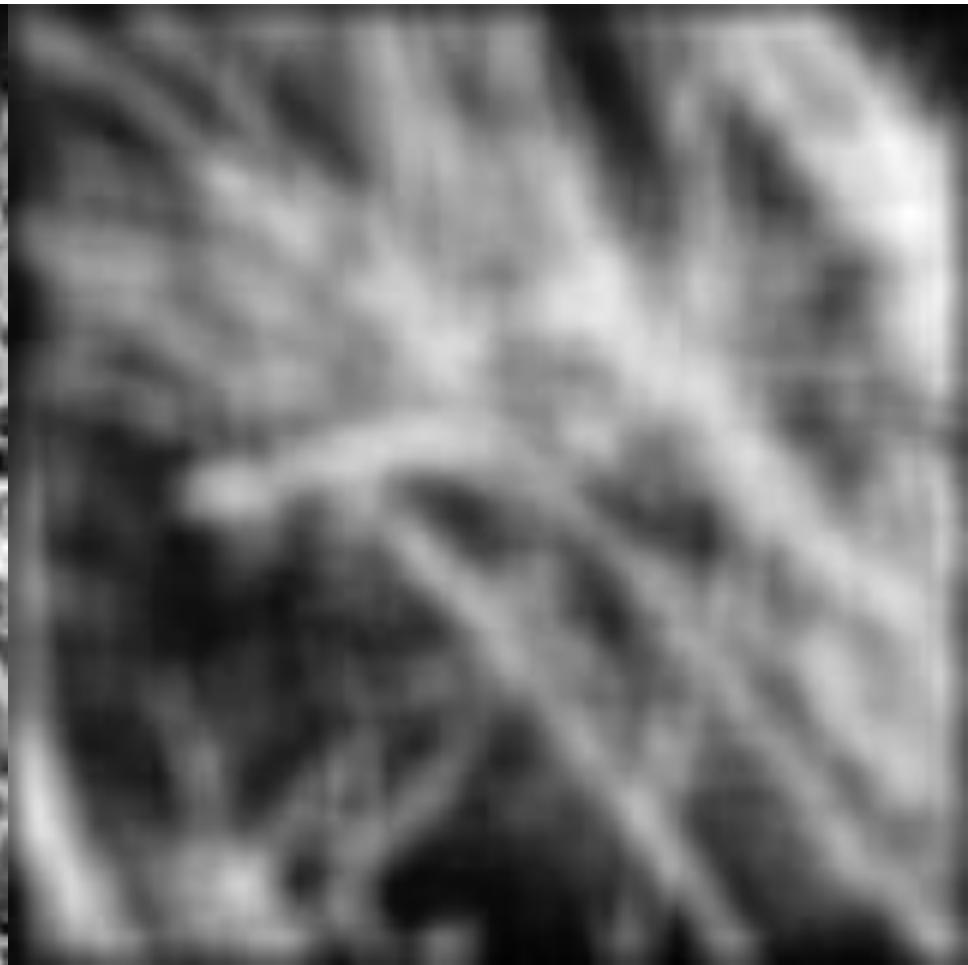


# Smoothing with box filter

$$f[\cdot, \cdot]$$

$$\frac{1}{9}$$

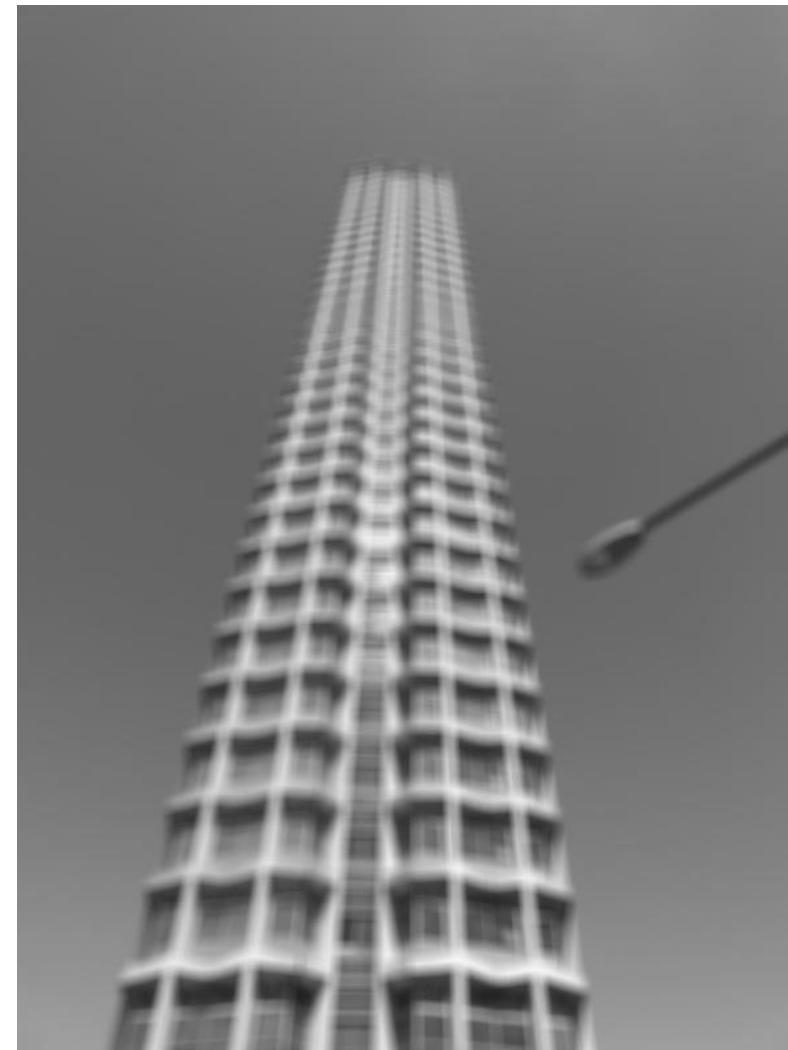
1	1	1
1	1	1
1	1	1



# A few more filters



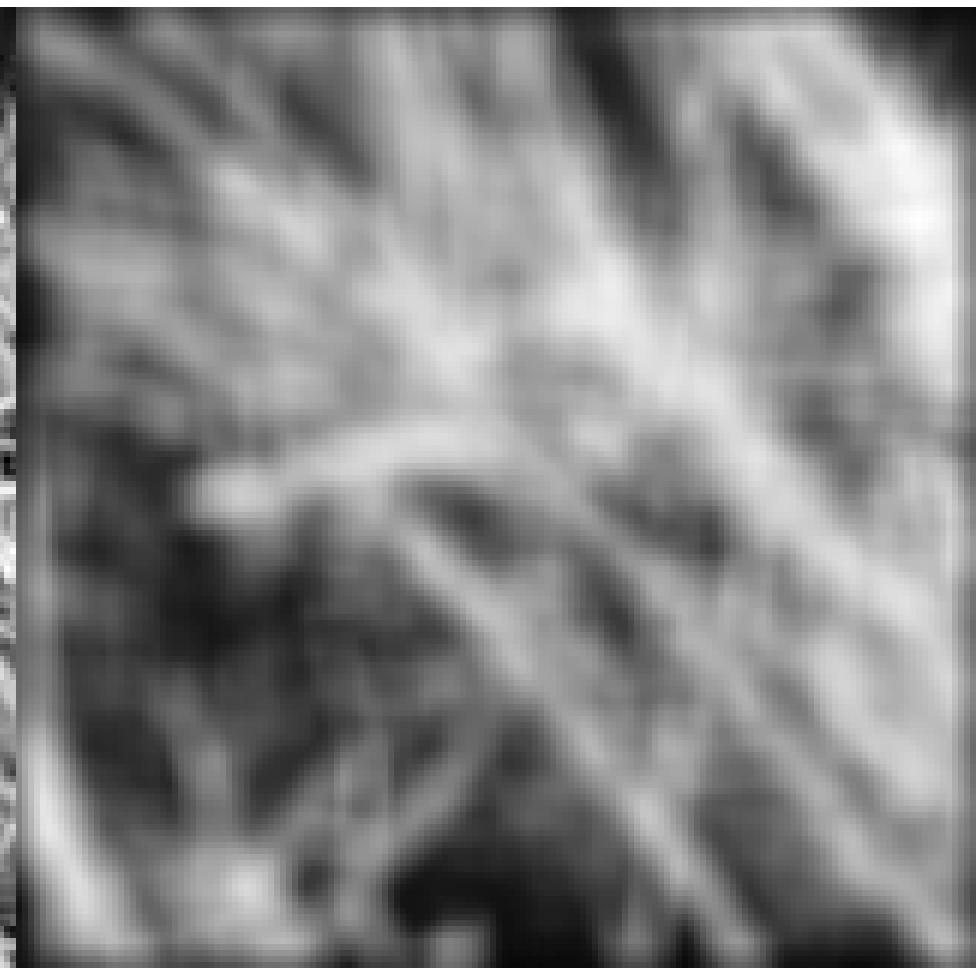
original



3x3 box filter

do you see  
any problems  
in this image?

# Example: Smoothing by Averaging

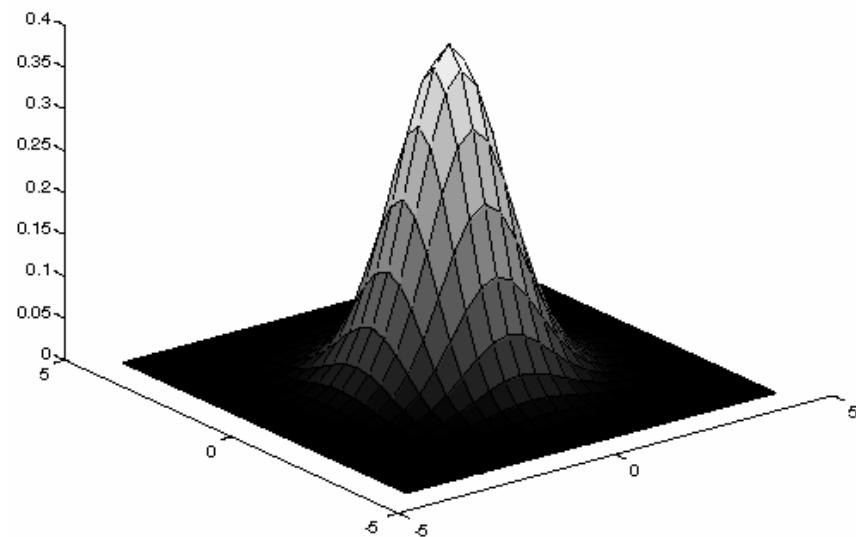


# Gaussian filters

- Remove “high-frequency” components from the image (a low-pass filter)
  - Images become more smooth
- Gaussian convolved with Gaussian...  
...is another Gaussian
  - So can smooth with small-width kernel, repeat, and get same result as larger-width kernel would have
  - Convolving twice with Gaussian kernel of width  $\sigma$  is same as convolving once with kernel of width  $\sigma\sqrt{2}$
- *Separable* kernel
  - Factors into product of two 1D Gaussians

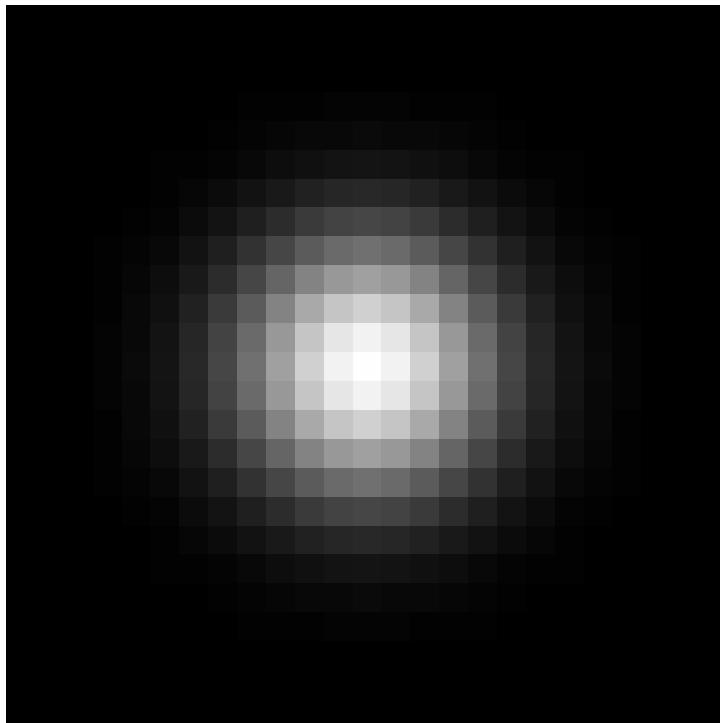
# Gaussian Averaging

- Rotationally symmetric.
- Weights nearby pixels more than distant ones.
  - This makes sense as probabilistic inference.



- A Gaussian gives a good model of a fuzzy blob

# An Isotropic Gaussian



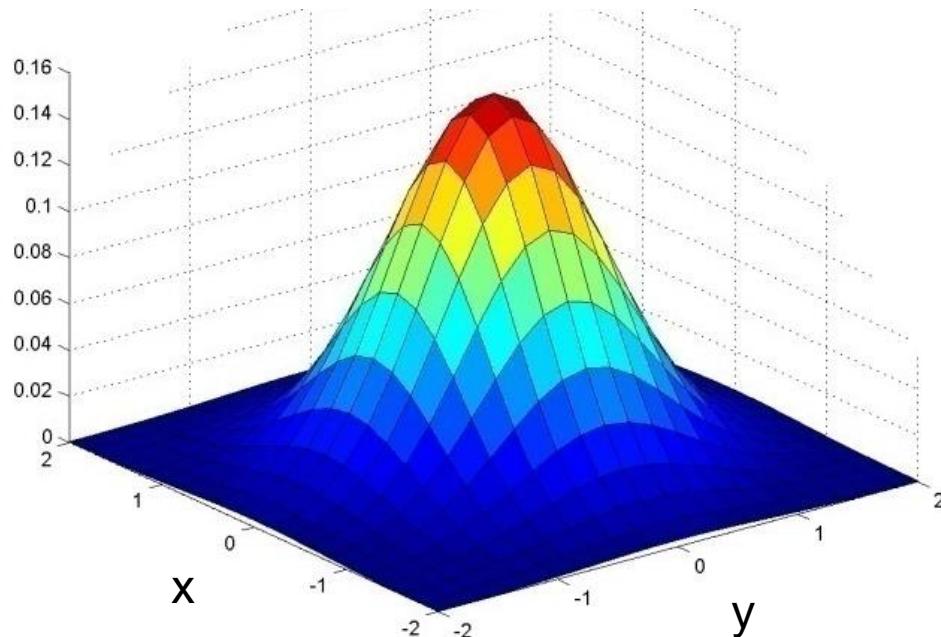
- The picture shows a smoothing kernel proportional to

$$\exp\left(-\left(\frac{x^2 + y^2}{2\sigma^2}\right)\right)$$

(which is a reasonable model of a circularly symmetric fuzzy blob)

# Important filter: Gaussian

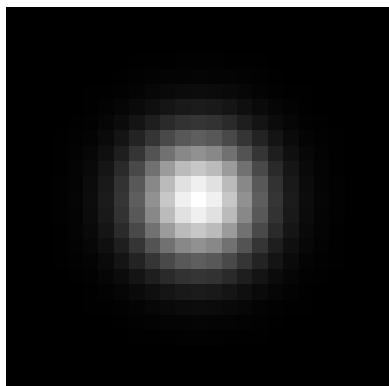
Weight contributions of neighboring pixels by nearness



x	0.003	0.013	0.022	0.013	0.003
y	0.013	0.059	0.097	0.059	0.013
x	0.022	0.097	0.159	0.097	0.022
y	0.013	0.059	0.097	0.059	0.013
x	0.003	0.013	0.022	0.013	0.003

Kernel size  $5 \times 5$ ,  
Standard deviation  $\sigma = 1$

Viewed  
from top



$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

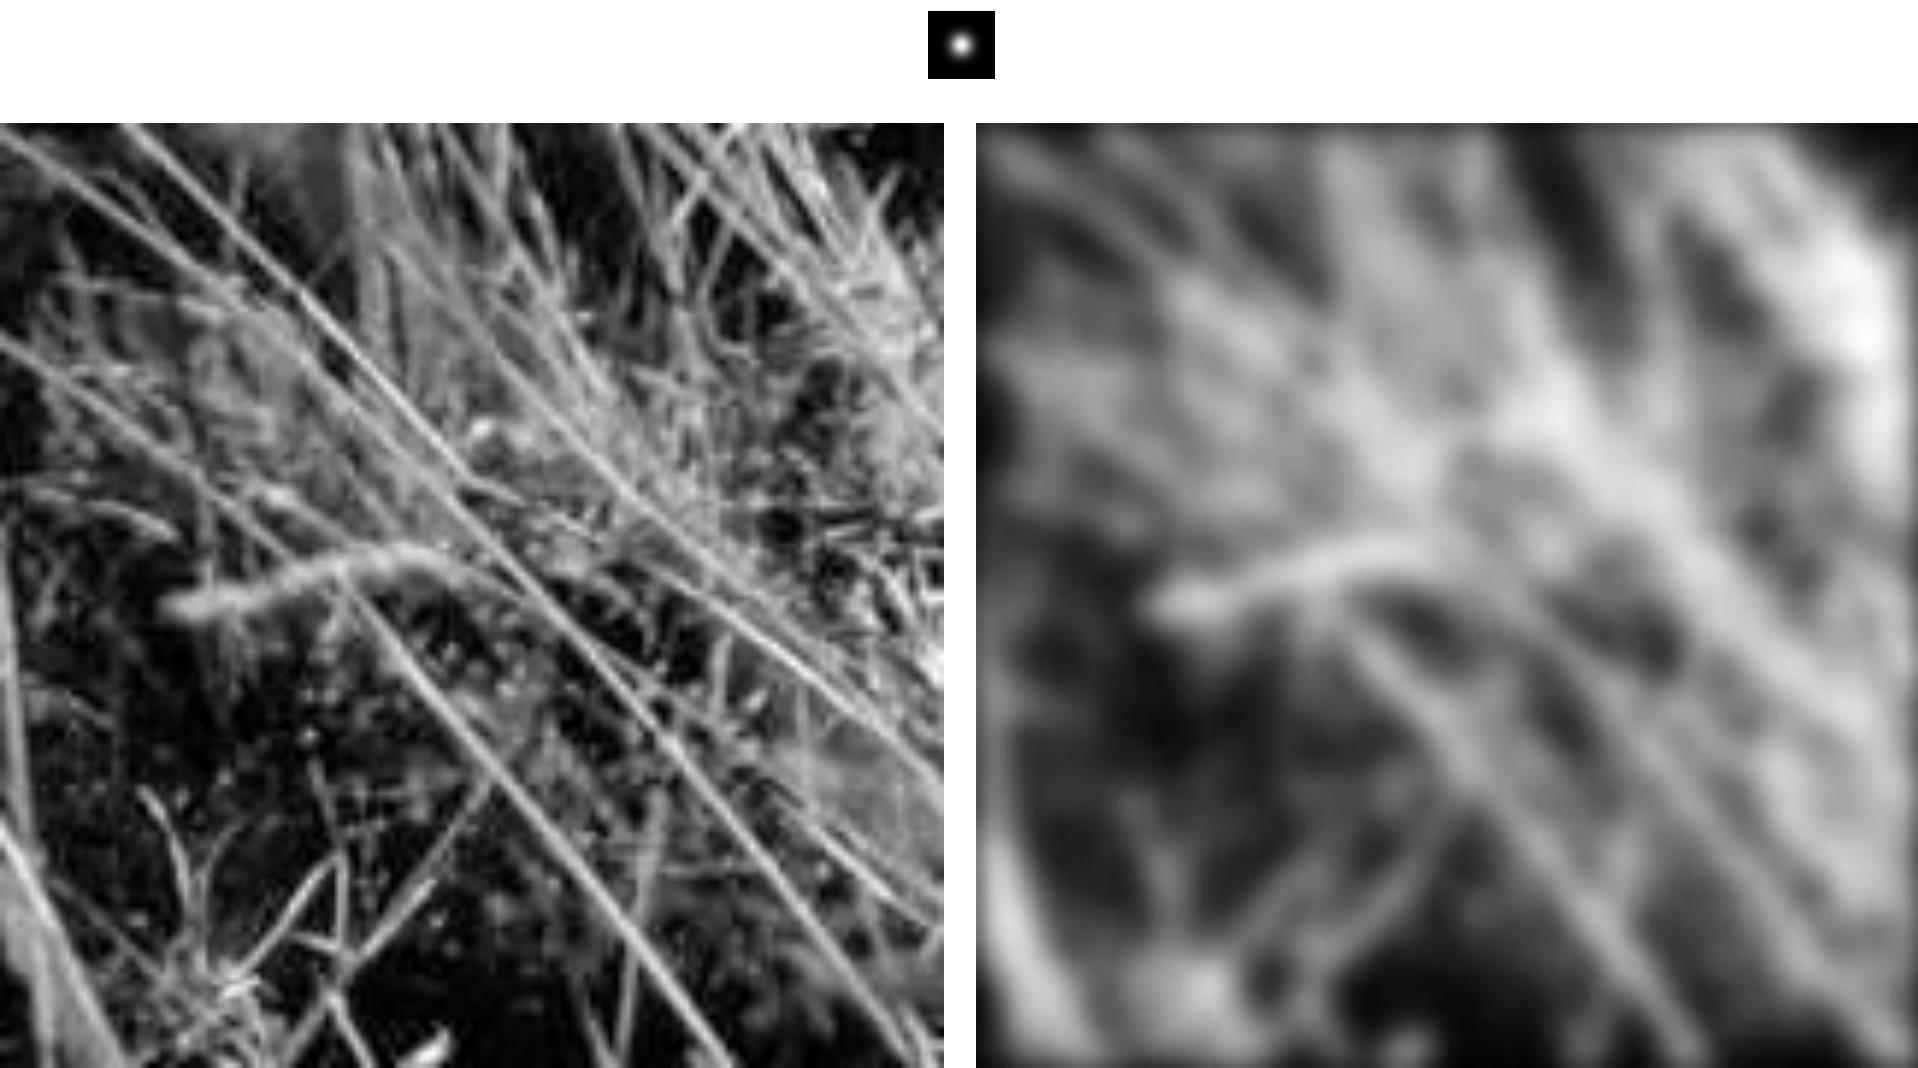
# Separability of the Gaussian filter

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

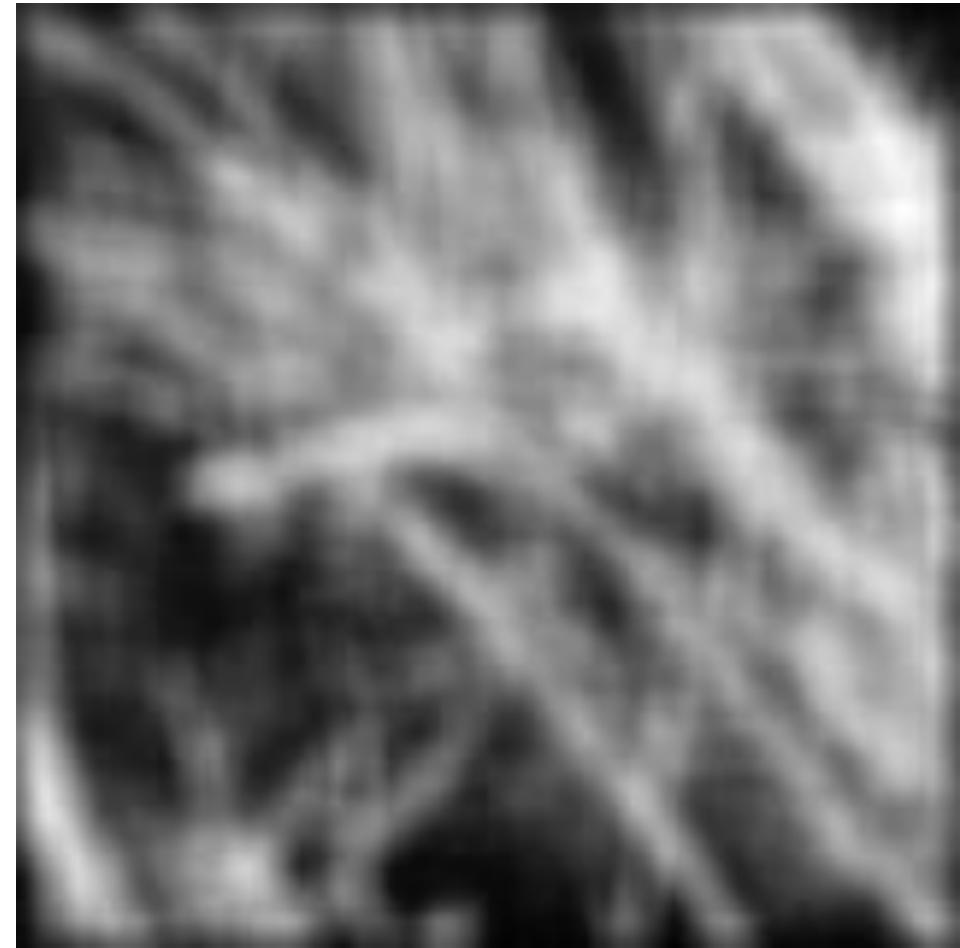
The 2D Gaussian can be expressed as the product of two functions, one a function of  $x$  and the other a function of  $y$

In this case, the two functions are the (identical) 1D Gaussian

# Smoothing with Gaussian filter



# Smoothing with box filter

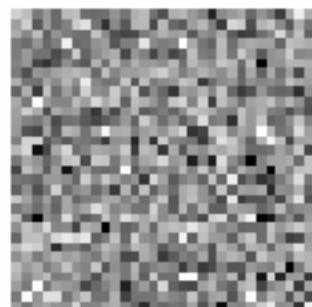
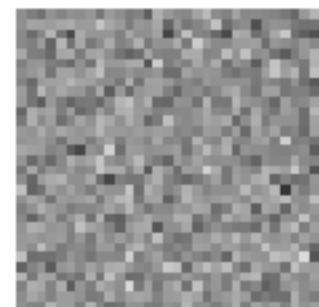


$\sigma=0.05$

$\sigma=0.1$

$\sigma=0.2$

no  
smoothing



$\sigma=1$  pixel

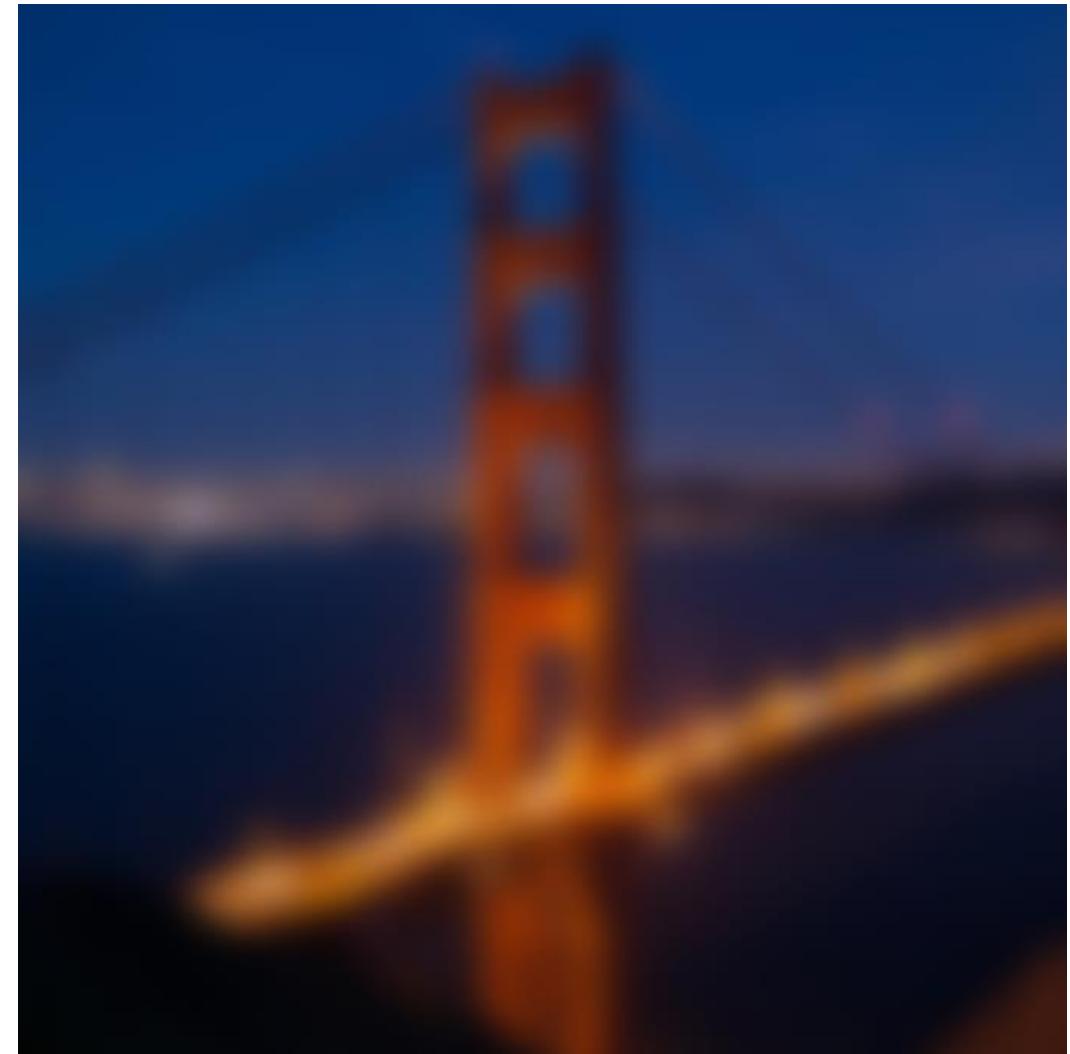


$\sigma=2$  pixels

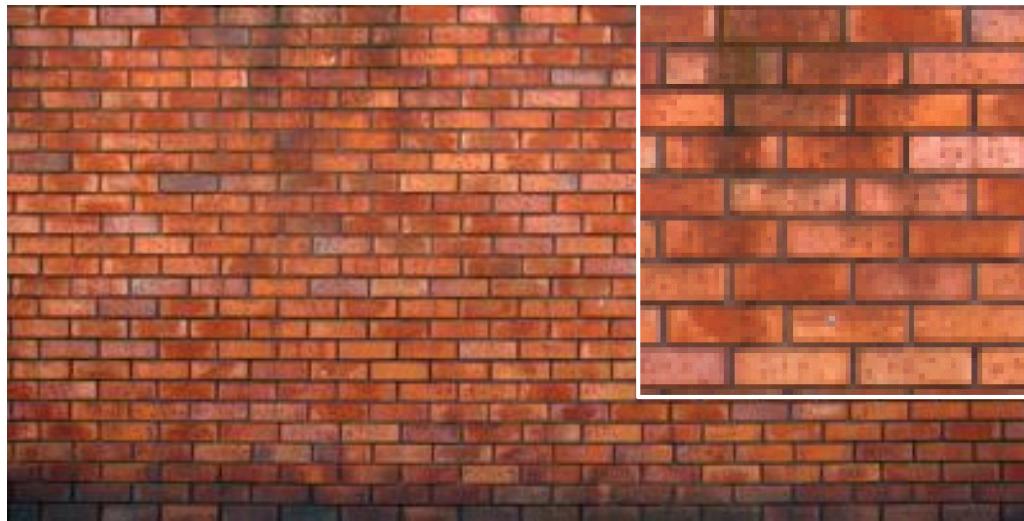
## The effects of smoothing

Each row shows smoothing with gaussians of different width; each column shows different realizations of an image of gaussian noise.

# Gaussian filtering example

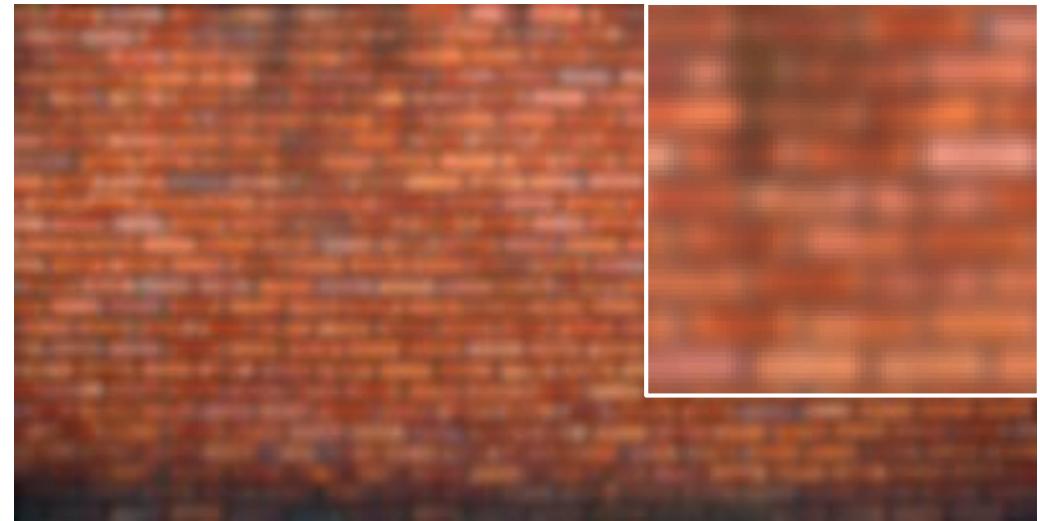


# Gaussian vs box filtering

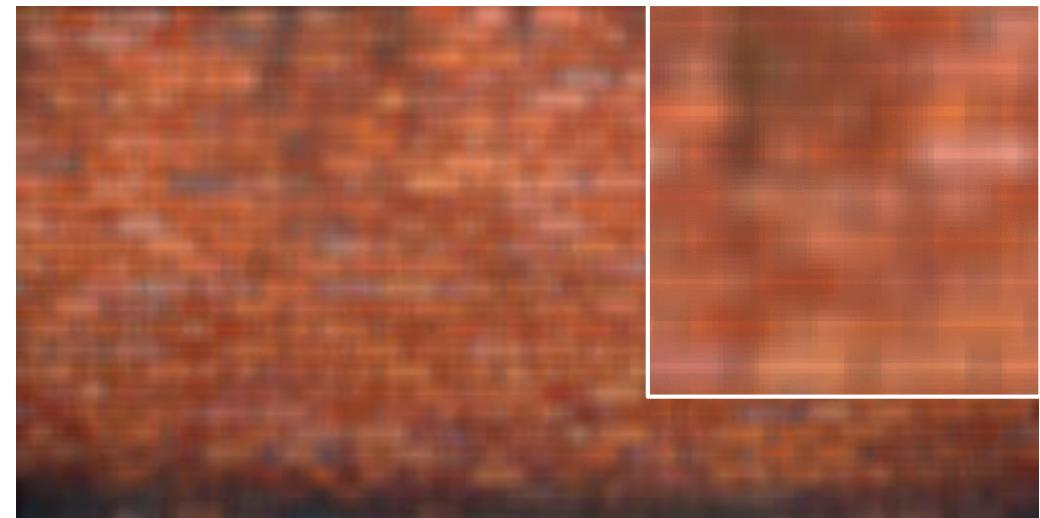


original

Which blur do you like better?



7x7 Gaussian

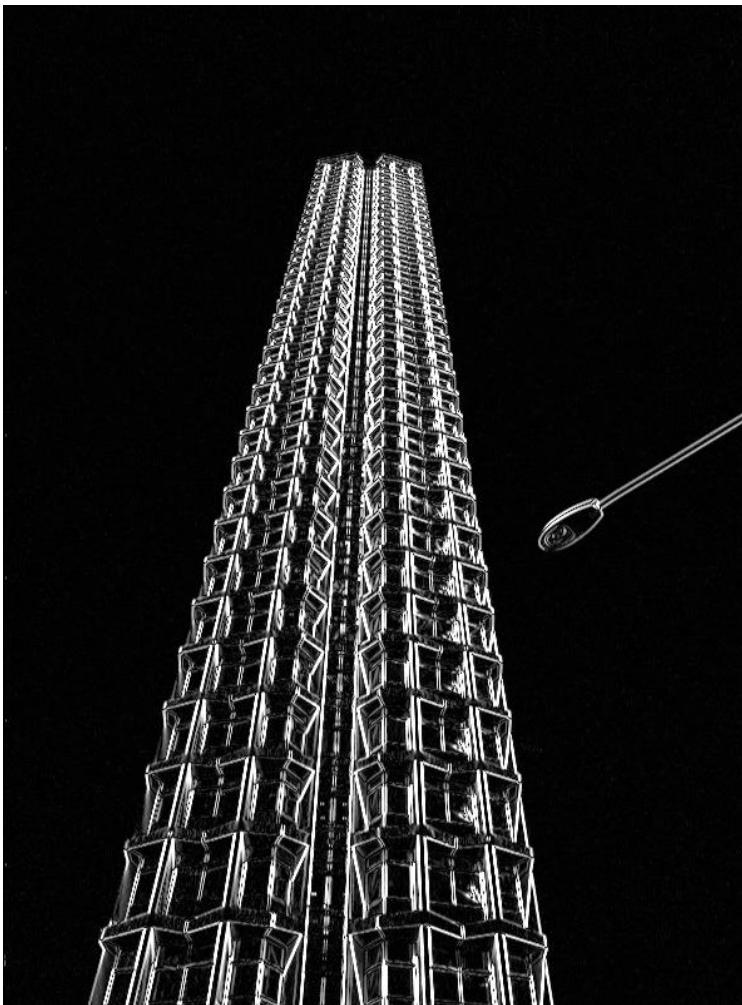


7x7 box

# Sobel filter example



original



which Sobel filter?

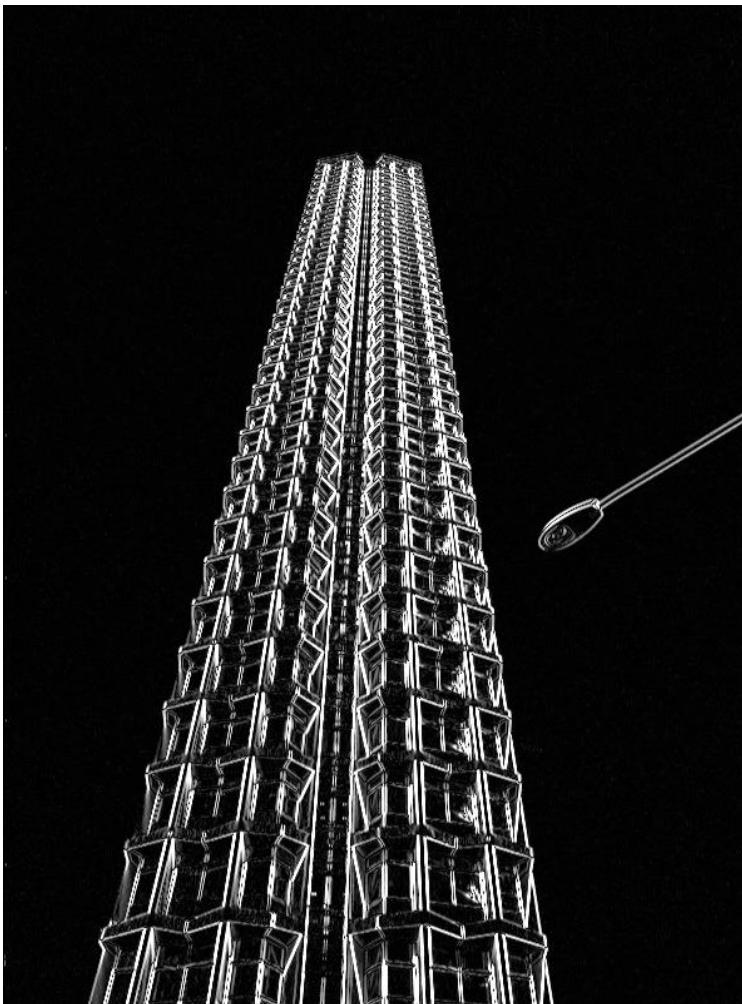


which Sobel filter?

# Sobel filter example



original

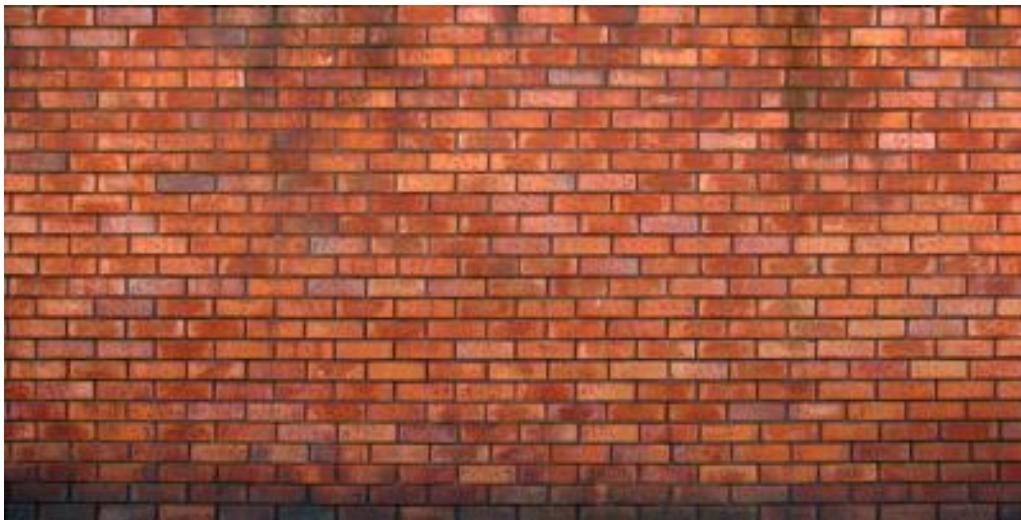


horizontal Sobel filter

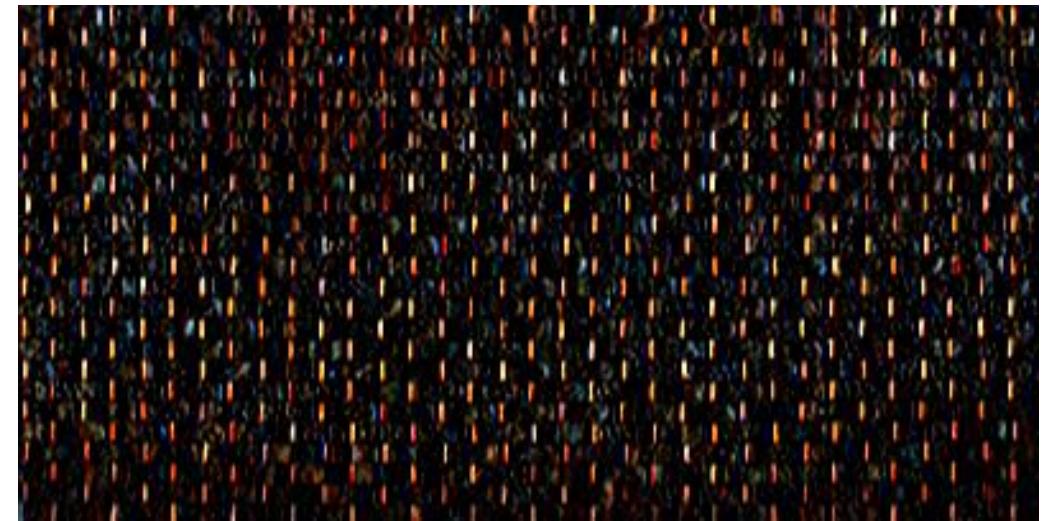


vertical Sobel filter

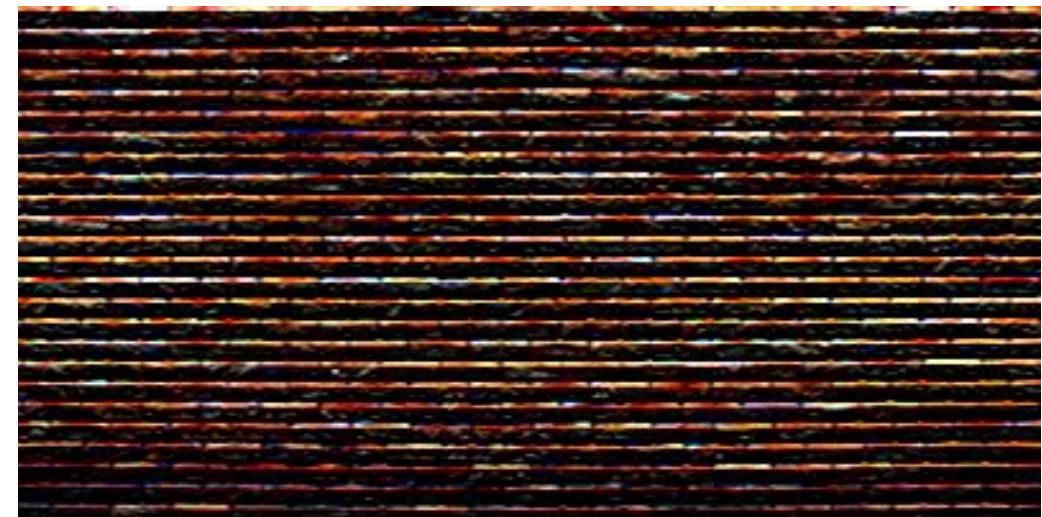
# Sobel filter example



original



horizontal Sobel filter



vertical Sobel filter

# Computing image gradients

1. Select your favorite derivative filters.

$$\mathbf{S}_x = \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array}$$

$$\mathbf{S}_y = \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

# Computing image gradients

1. Select your favorite derivative filters.

$$\mathbf{S}_x = \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array} \quad \mathbf{S}_y = \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

2. Convolve with the image to compute derivatives.

$$\frac{\partial \mathbf{f}}{\partial x} = \mathbf{S}_x \otimes \mathbf{f}$$

$$\frac{\partial \mathbf{f}}{\partial y} = \mathbf{S}_y \otimes \mathbf{f}$$

# Computing image gradients

1. Select your favorite derivative filters.

$$\mathbf{S}_x = \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array}$$

$$\mathbf{S}_y = \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

2. Convolve with the image to compute derivatives.

$$\frac{\partial \mathbf{f}}{\partial x} = \mathbf{S}_x \otimes \mathbf{f}$$

$$\frac{\partial \mathbf{f}}{\partial y} = \mathbf{S}_y \otimes \mathbf{f}$$

3. Form the image gradient, and compute its direction and amplitude.

$$\nabla \mathbf{f} = \left[ \frac{\partial \mathbf{f}}{\partial x}, \frac{\partial \mathbf{f}}{\partial y} \right]$$

gradient

$$\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

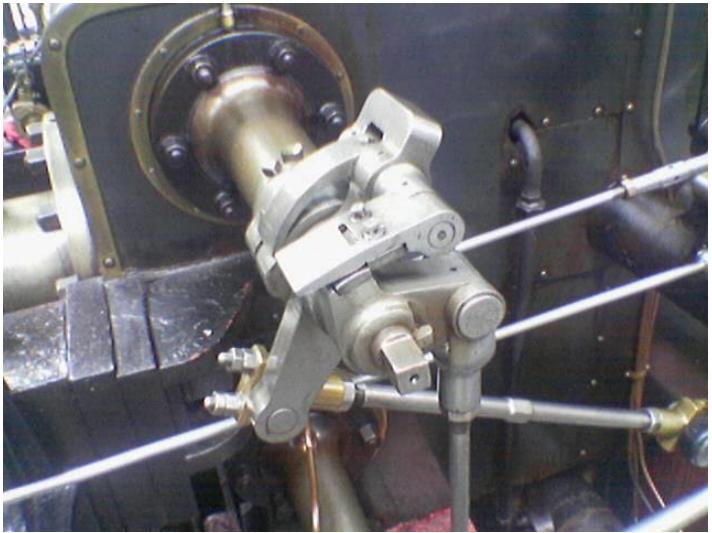
direction

$$\|\nabla f\| = \sqrt{\left( \frac{\partial f}{\partial x} \right)^2 + \left( \frac{\partial f}{\partial y} \right)^2}$$

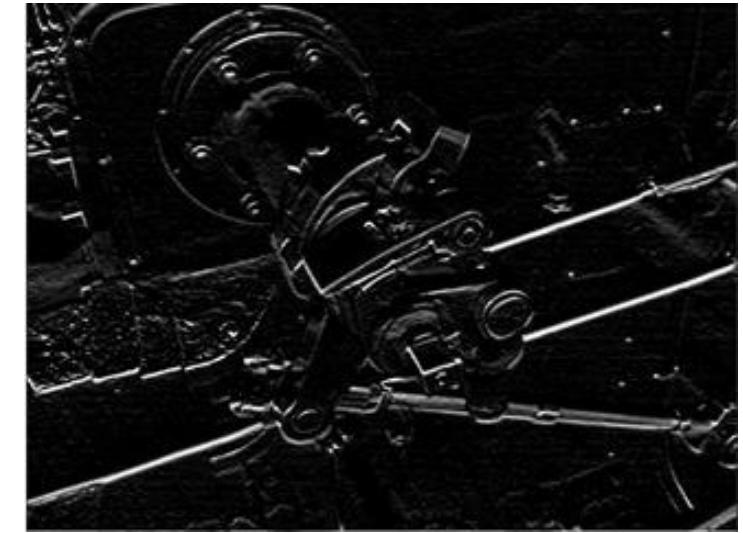
amplitude

# Image gradient example

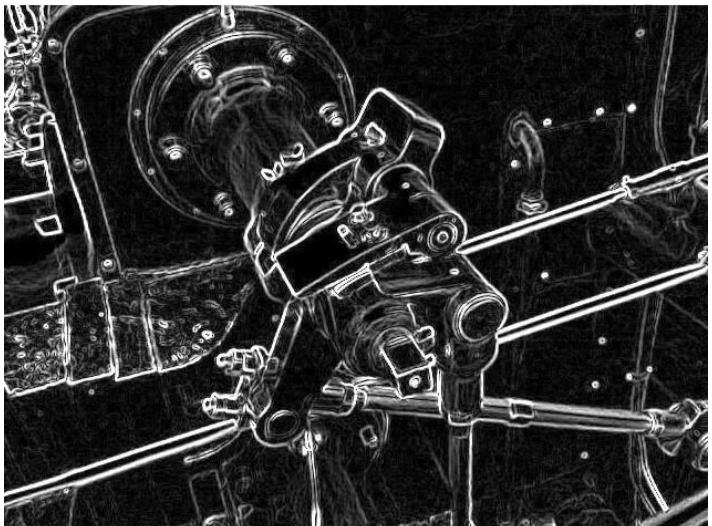
original



vertical derivative



gradient amplitude

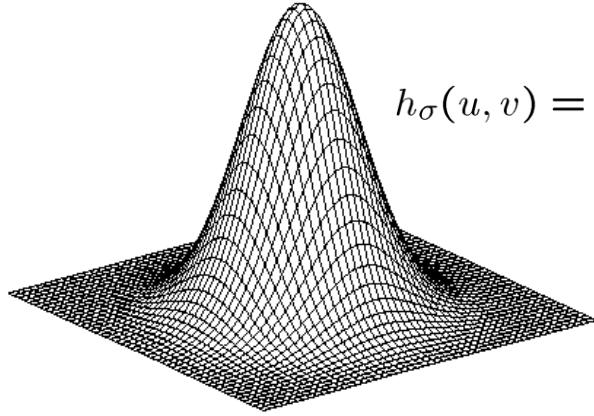


horizontal derivative



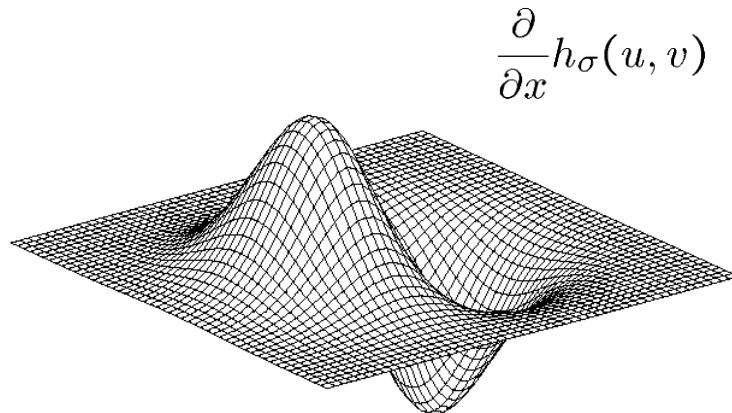
How does the gradient direction relate to these edges?

# 2D Gaussian filters



Gaussian

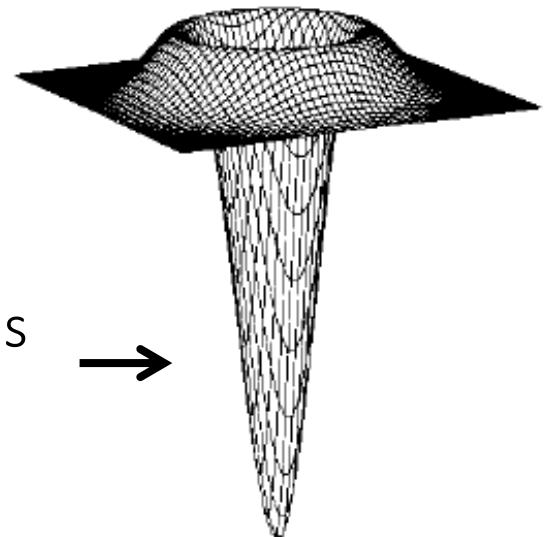
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



Derivative of Gaussian

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

$$\nabla^2 h_\sigma(u, v)$$



Laplacian of Gaussian

how does this relate to this  
lecture's cover picture?

# Laplace and LoG filtering examples



Laplacian of Gaussian filtering



Laplace filtering

# Laplacian of Gaussian vs Derivative of Gaussian

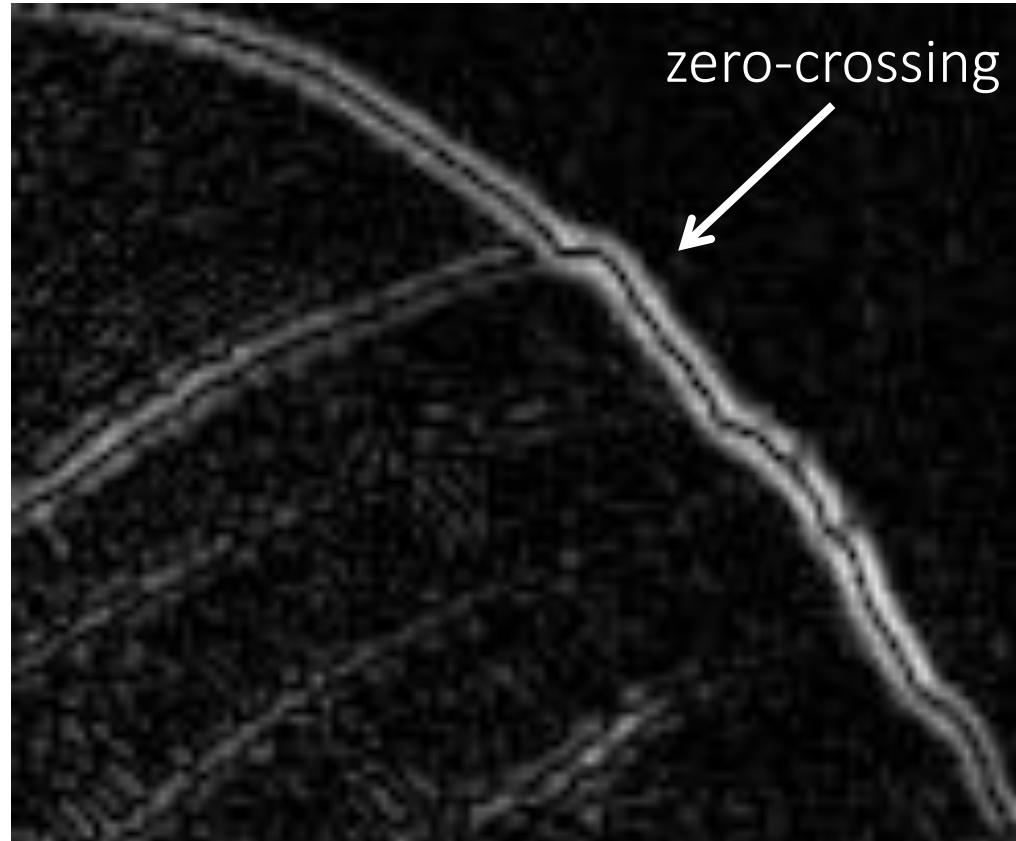


Laplacian of Gaussian filtering

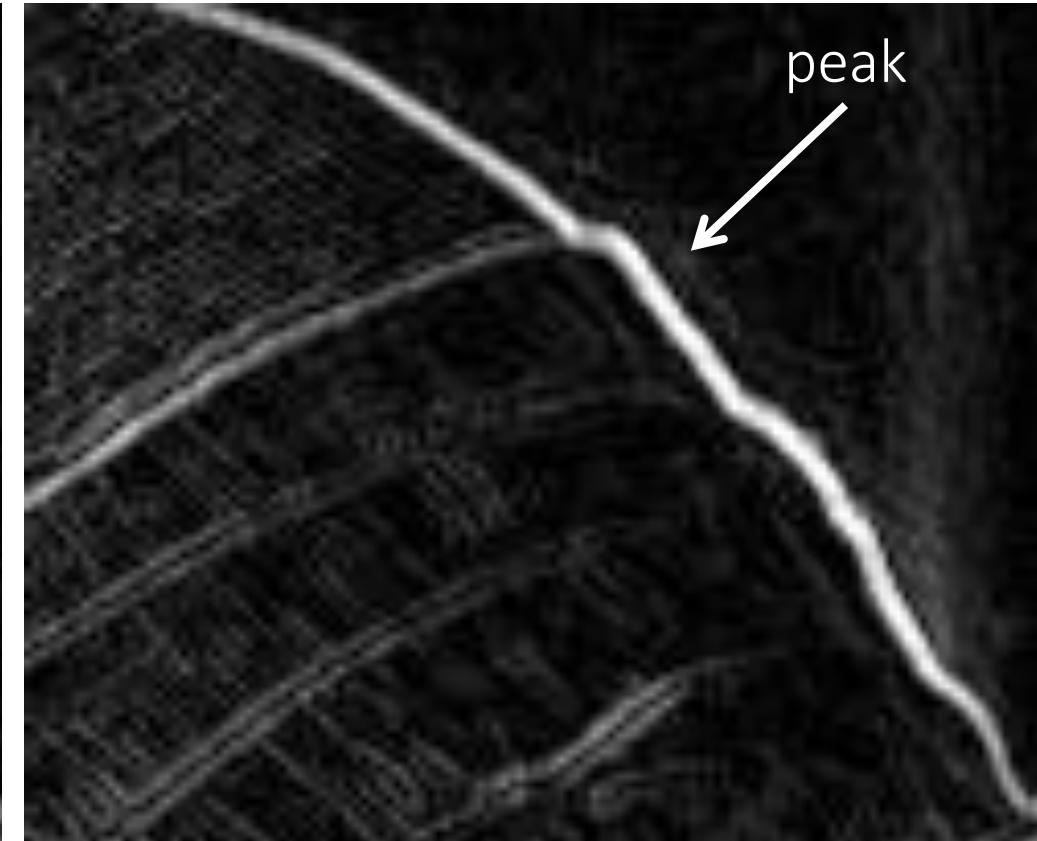


Derivative of Gaussian filtering

# Laplacian of Gaussian vs Derivative of Gaussian



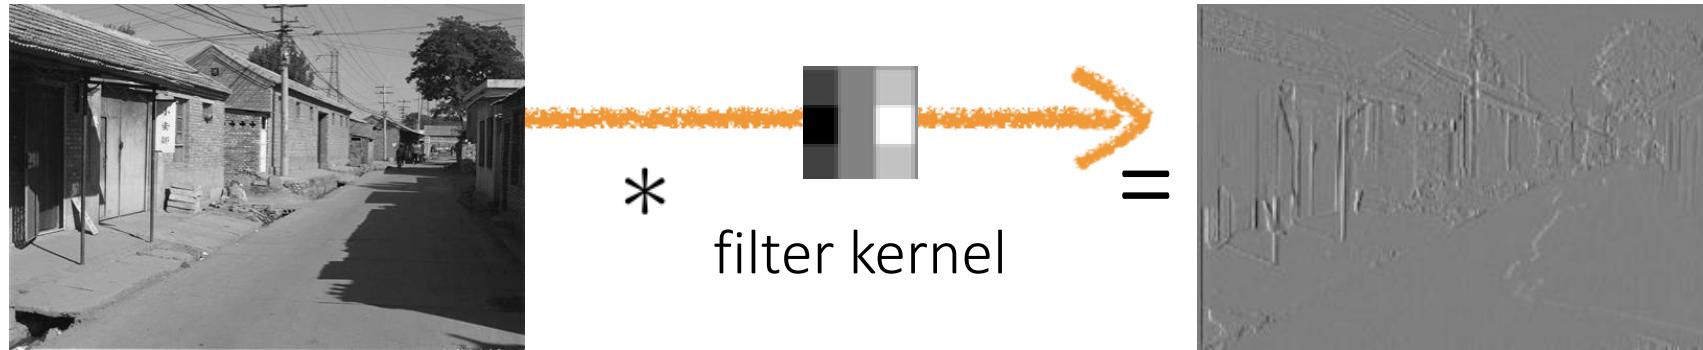
Laplacian of Gaussian filtering



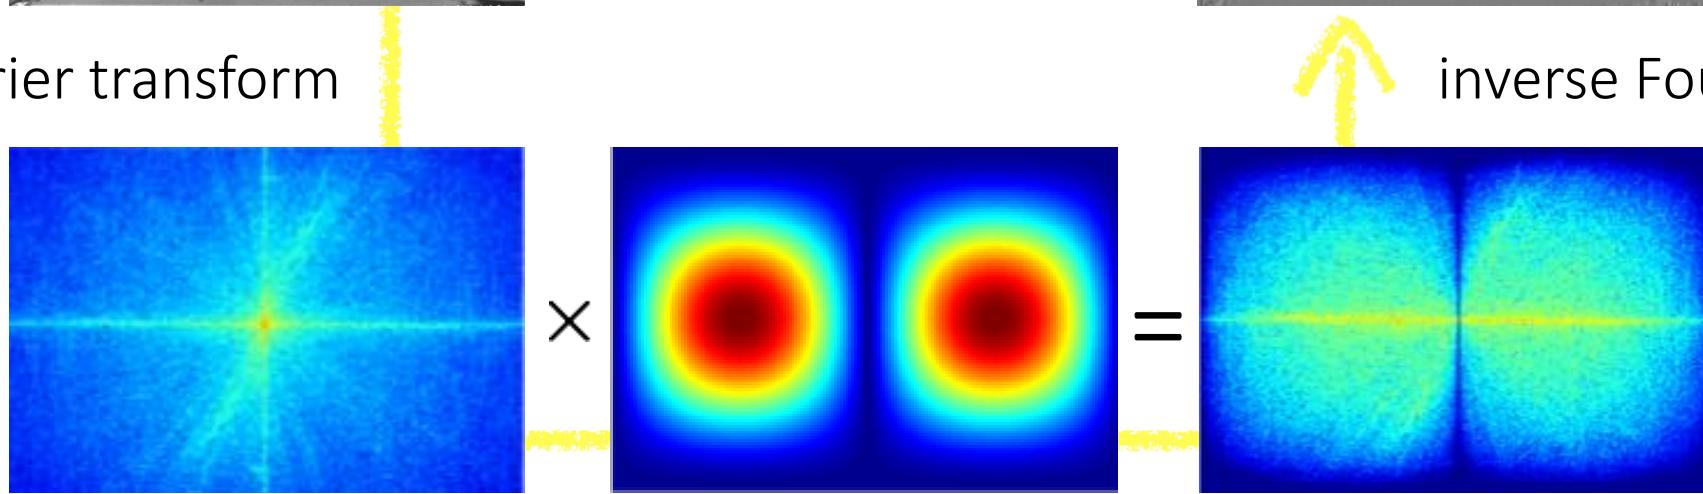
Derivative of Gaussian filtering

Zero crossings are more accurate at localizing edges (but not very convenient).

# Spatial domain filtering

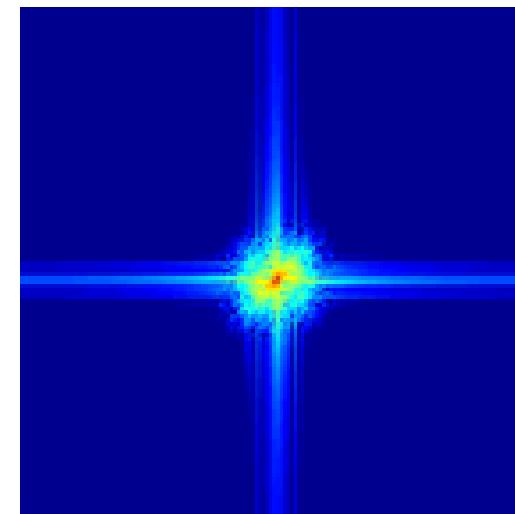
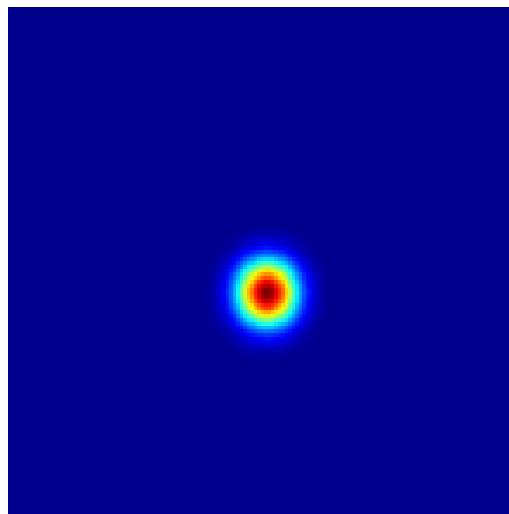
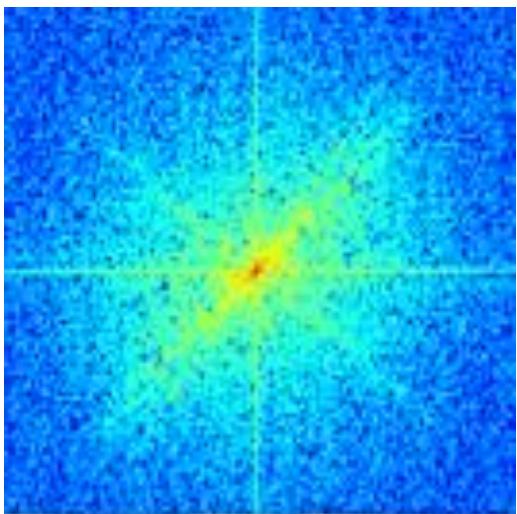
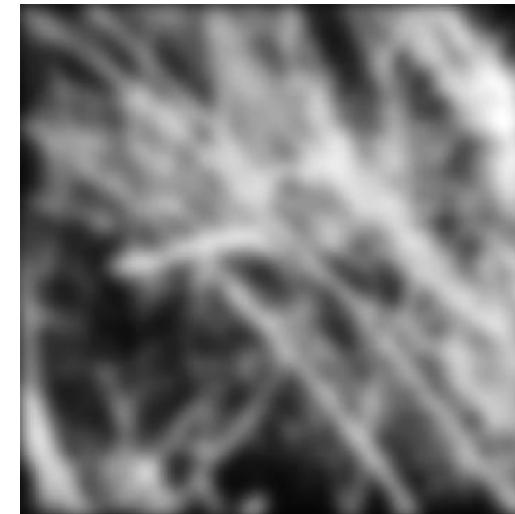
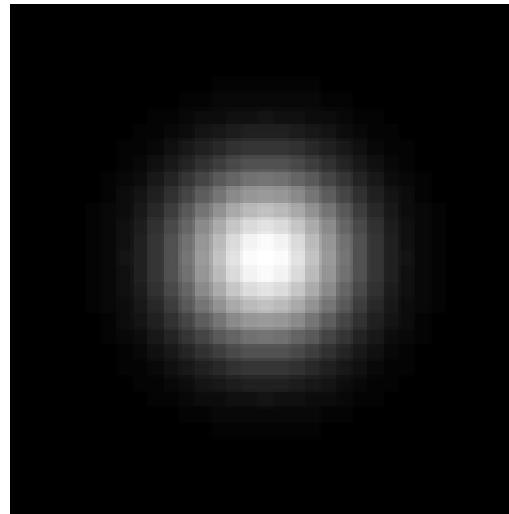


Fourier transform

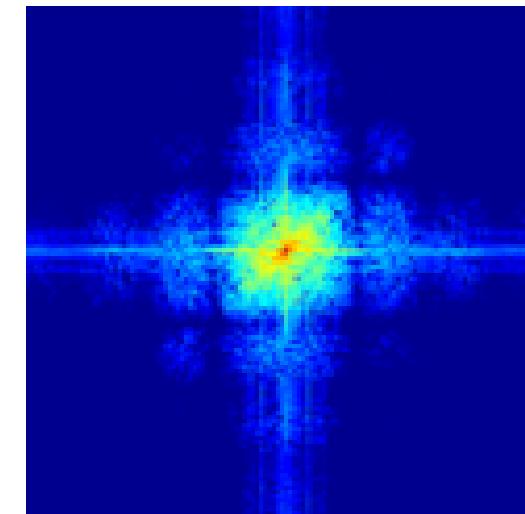
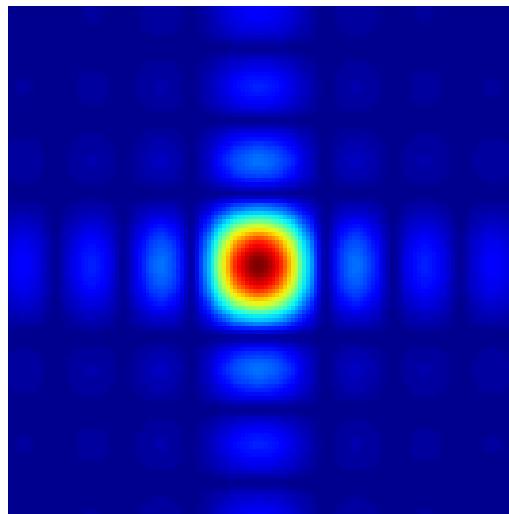
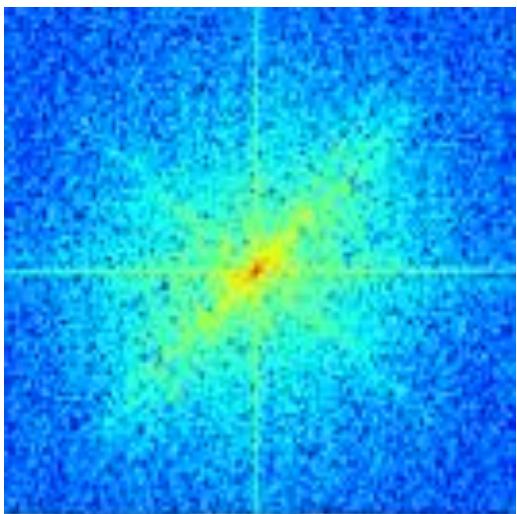
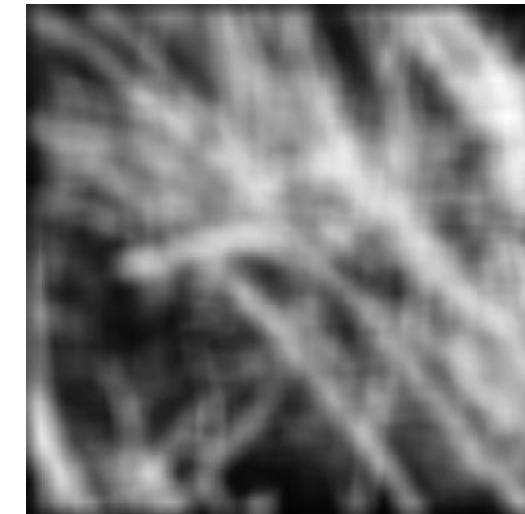
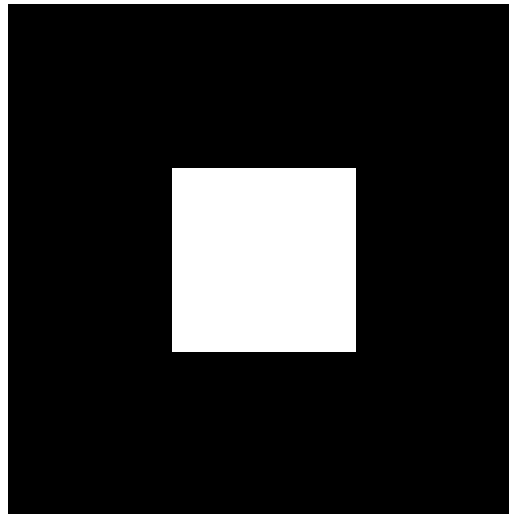


# Frequency domain filtering

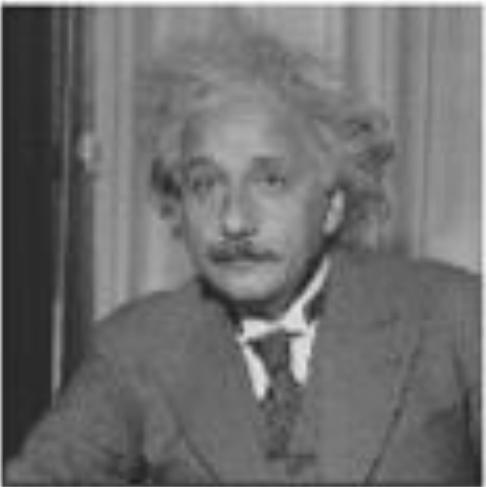
# Gaussian blur



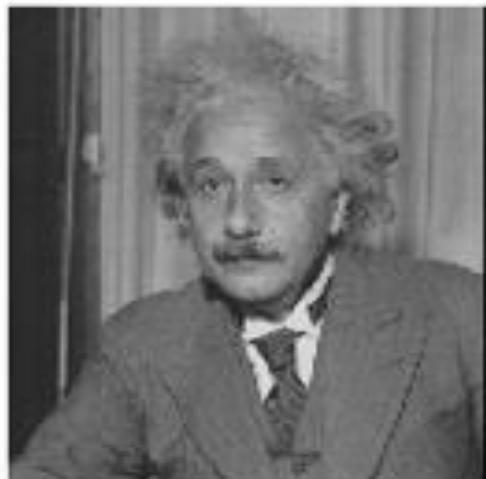
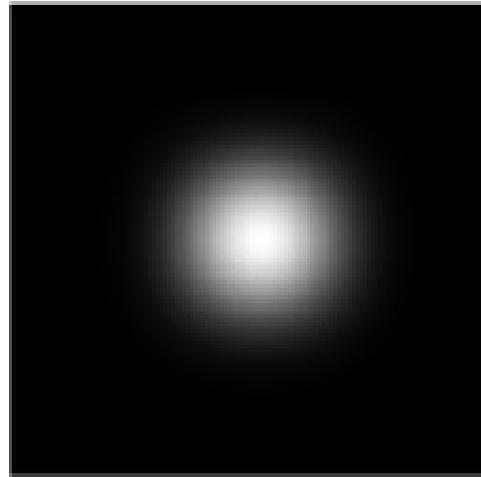
# Box blur



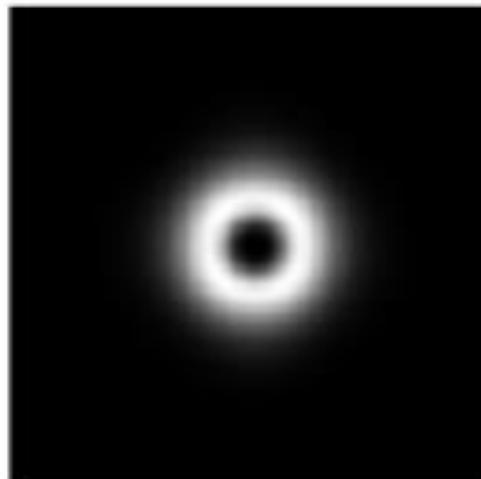
# More filtering examples



?

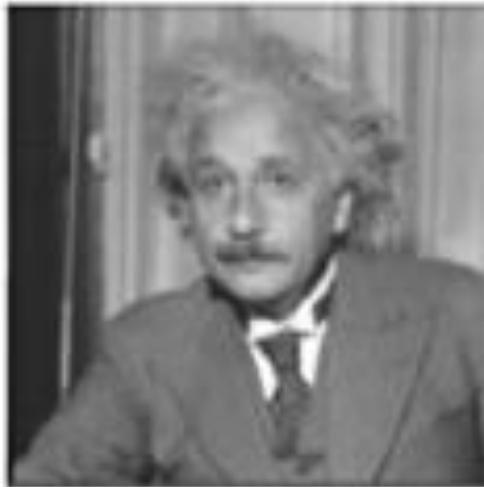
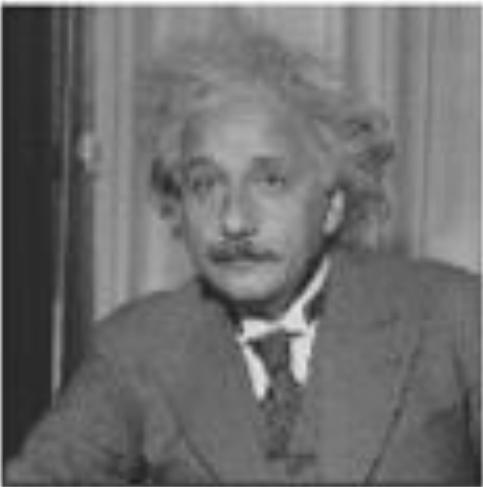


?

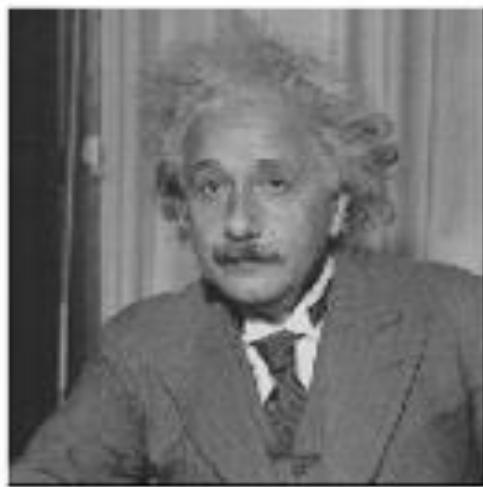
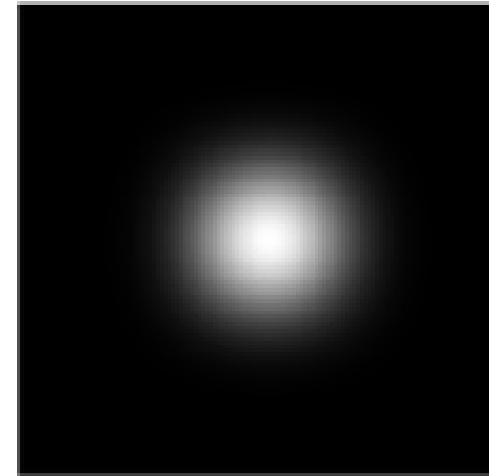


filters shown  
in frequency-  
domain

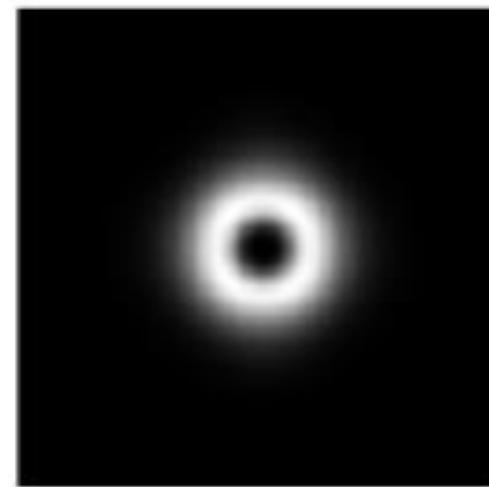
# More filtering examples



low-pass



band-pass

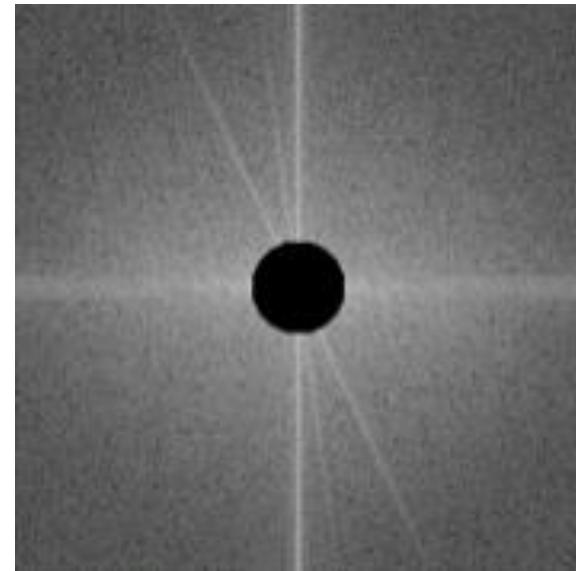


filters shown  
in frequency-  
domain

# More filtering examples

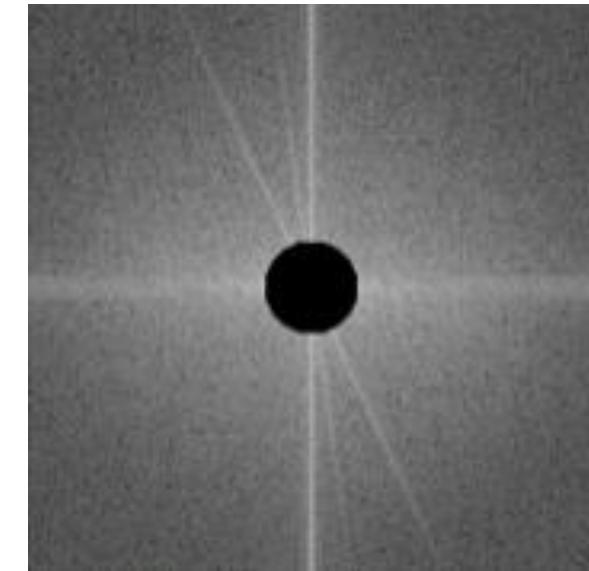
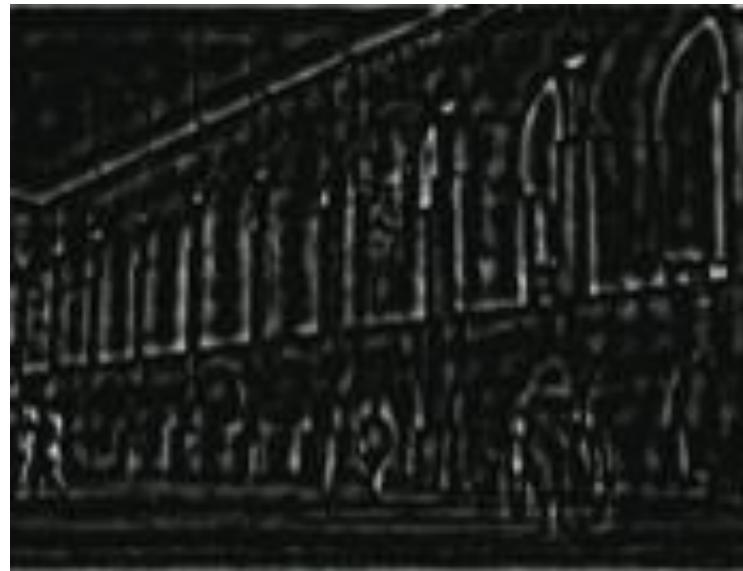


?



high-pass

# More filtering examples



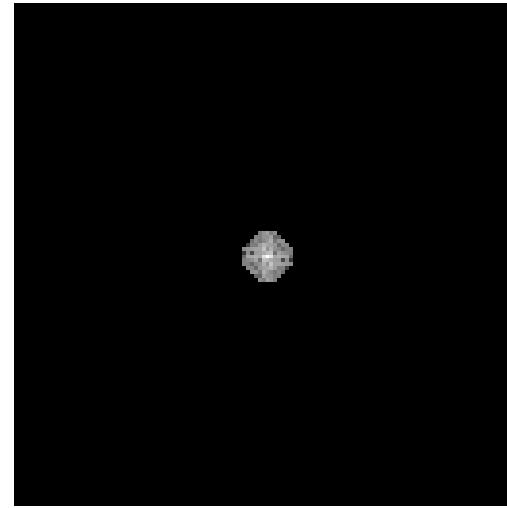
high-pass

# More filtering examples

original image

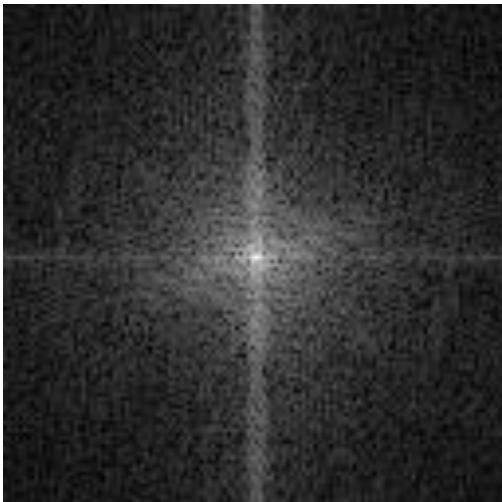


low-pass filter



?

frequency magnitude

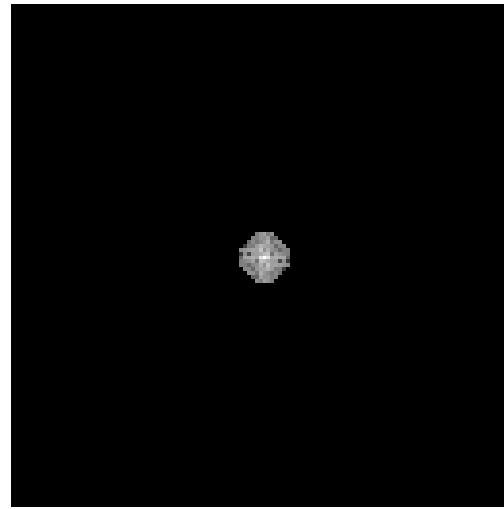


# More filtering examples

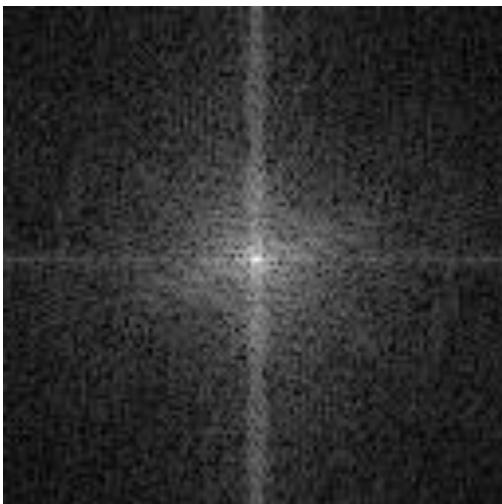
original image



low-pass filter



frequency magnitude

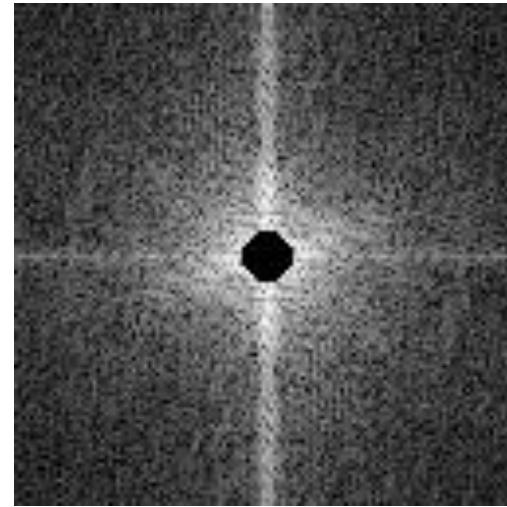


# More filtering examples

original image

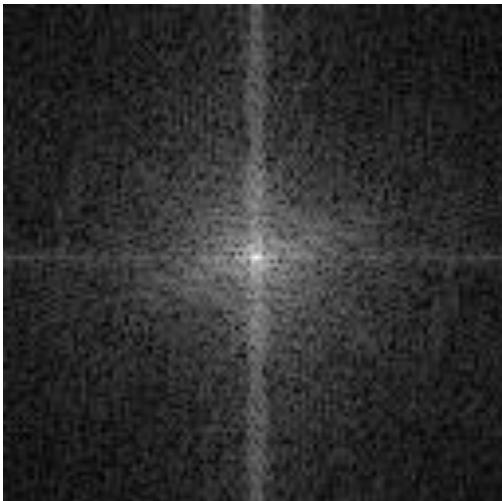


high-pass filter



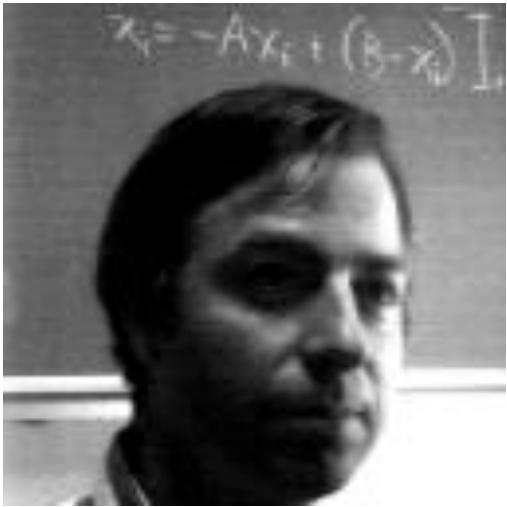
?

frequency magnitude

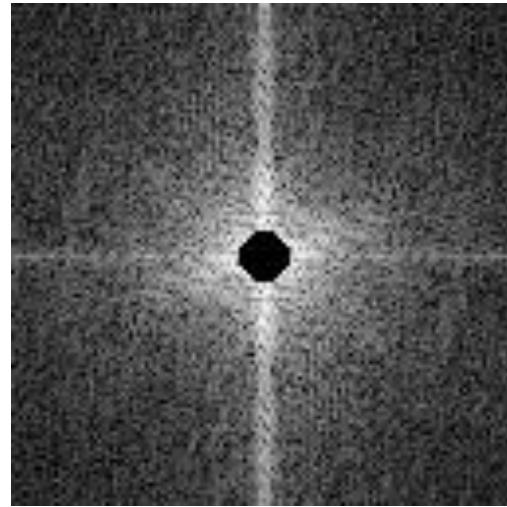


# More filtering examples

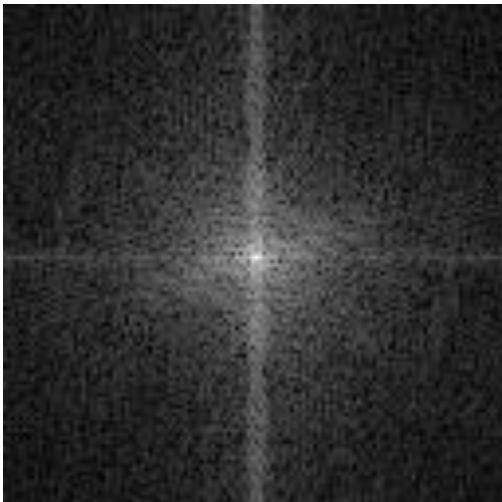
original image



high-pass filter



frequency magnitude

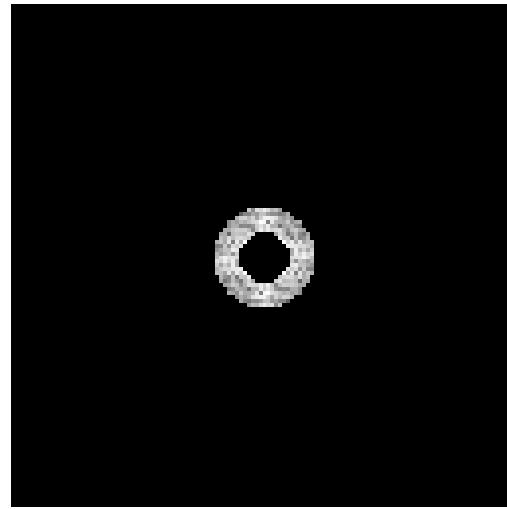


# More filtering examples

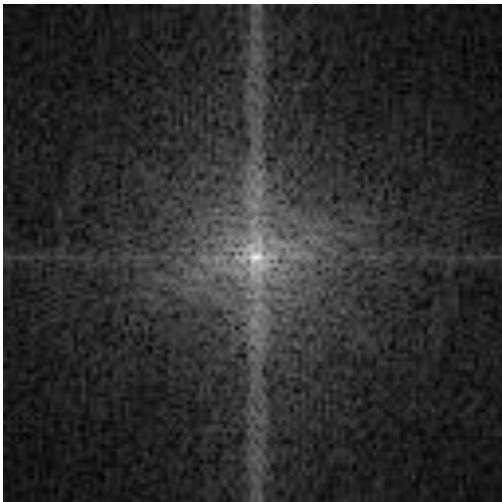
original image



band-pass filter



frequency magnitude

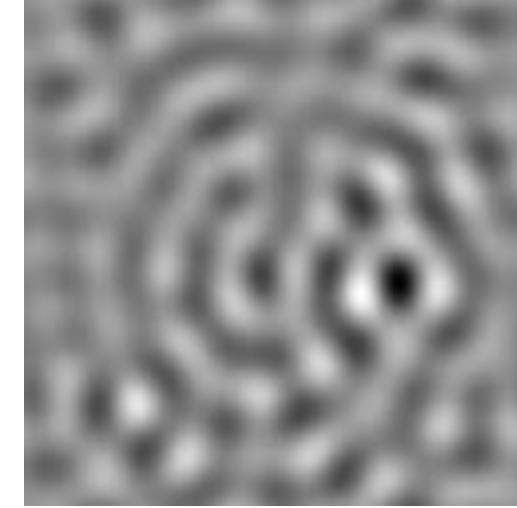
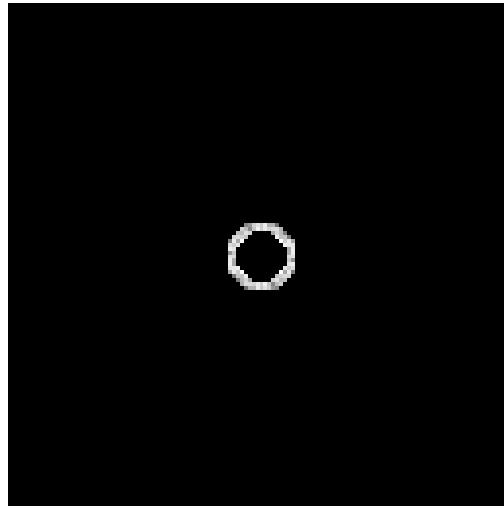


# More filtering examples

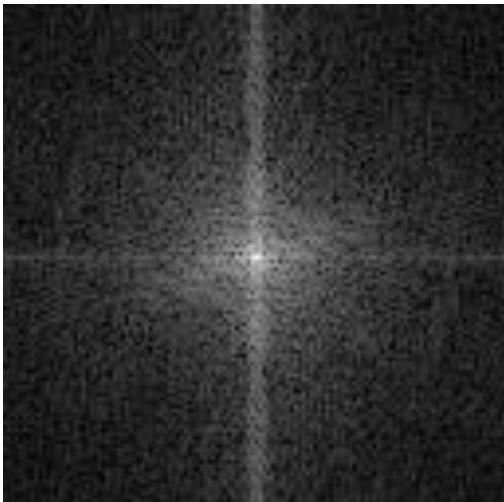
original image



band-pass filter



frequency magnitude

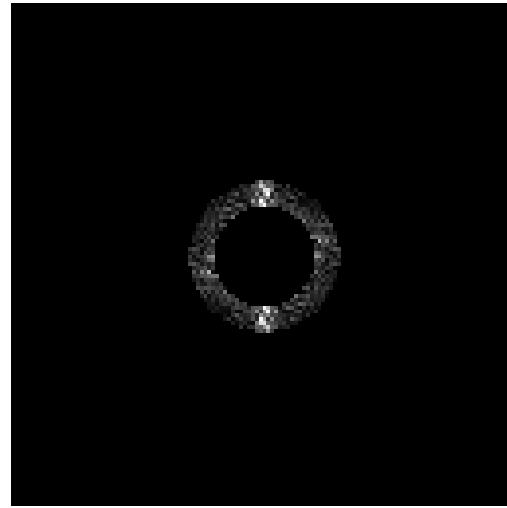


# More filtering examples

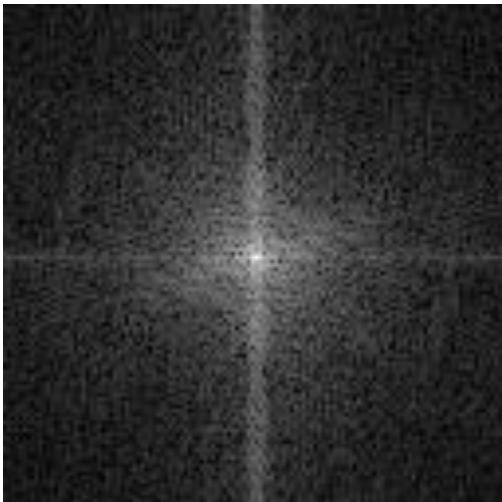
original image



band-pass filter



frequency magnitude

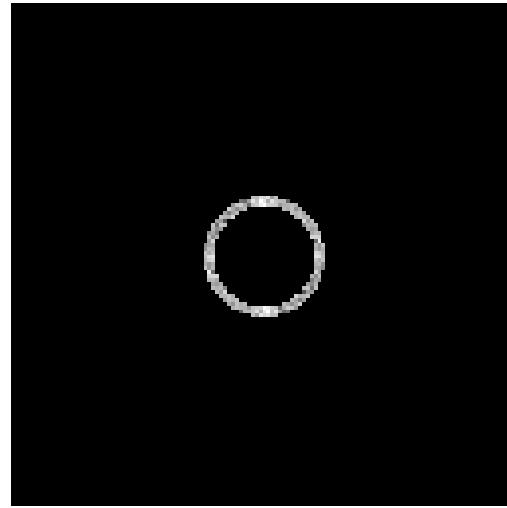


# More filtering examples

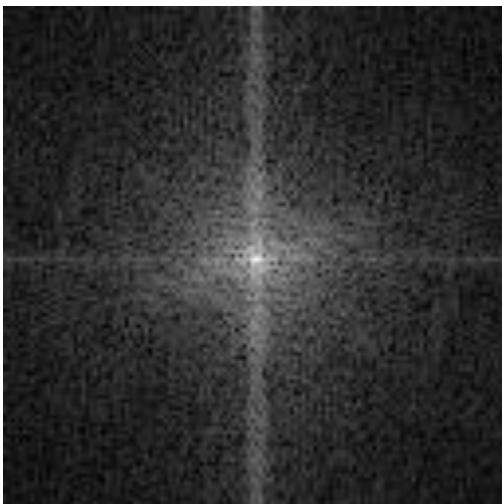
original image



band-pass filter



frequency magnitude



# Common types of noise

- **Salt and pepper noise:** random occurrences of black and white pixels
- **Impulse noise:** random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian normal distribution



Original



Salt and pepper noise



Impulse noise



Gaussian noise

# Filtering to reduce noise

- Noise is what we're not interested in.
  - We'll discuss simple, low-level noise today: Light fluctuations; Sensor noise; Quantization effects; Finite precision
  - Not complex: shadows; extraneous objects.
- A pixel's neighborhood contains information about its intensity.
- Averaging noise reduces its effect.

# Additive noise

- $I = S + N$ . Noise doesn't depend on signal.
- We'll consider:

$$I_i = s_i + n_i \text{ with } E(n_i) = 0$$

$s_i$  deterministic.

$n_i, n_j$  independent for  $n_i \neq n_j$

$n_i, n_j$  identically distributed

# Motivation: noise reduction



- We can measure **noise** in multiple images of the same static scene.
- How could we reduce the noise, i.e., give an estimate of the true intensities?
- **What if there's only one image?**

# Average Filter

- Mask with positive entries, that sum 1.
- Replaces each pixel with an average of its neighborhood.
- If all weights are equal, it is called a BOX filter.

$$F = \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

$1/9$

(Camps)

# Does it reduce noise?

- Intuitively, takes out small variations.

$$I(i, j) = \hat{I}(i, j) + N(i, j) \text{ with } N(i, j) \sim N(0, \sigma)$$

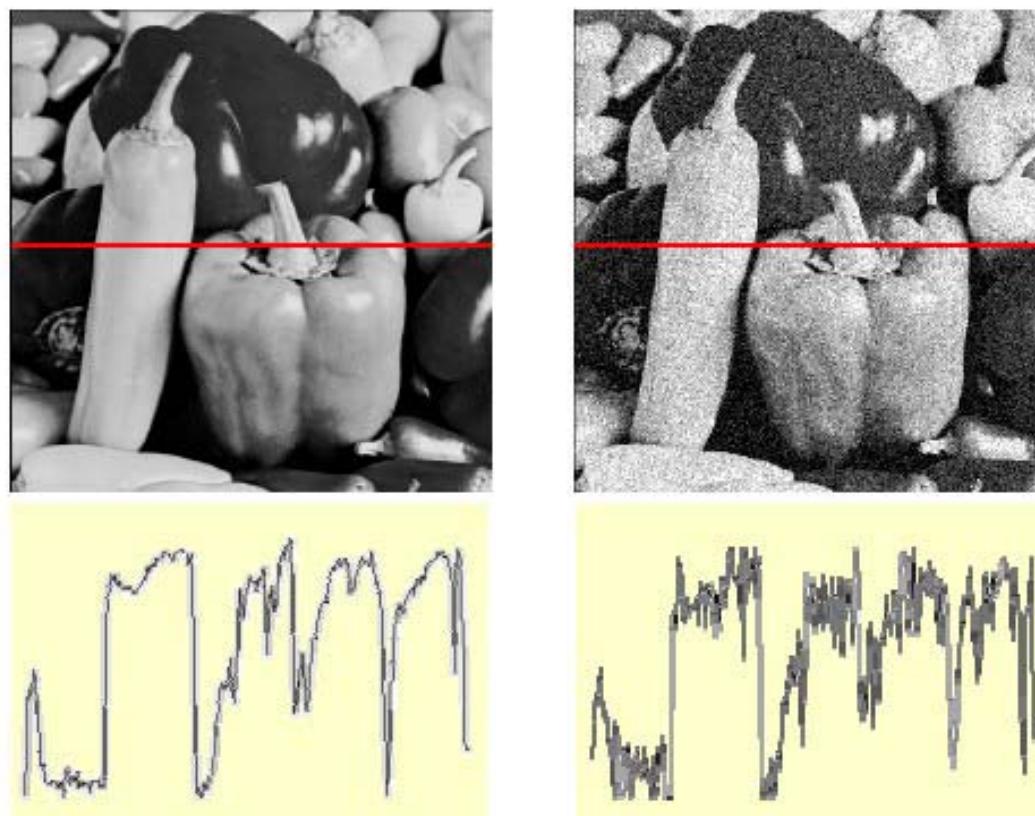
$$\begin{aligned} O(i, j) &= \frac{1}{m^2} \sum_{h=-m/2}^{m/2} \sum_{k=-m/2}^{m/2} \hat{I}(i-h, j-k) + N(i-h, j-k) = \\ &= \frac{1}{m^2} \sum_{h=-m/2}^{m/2} \sum_{k=-m/2}^{m/2} \hat{I}(i-h, j-k) + \underbrace{\frac{1}{m^2} \sum_{h=-m/2}^{m/2} \sum_{k=-m/2}^{m/2} N(i-h, j-k)}_{\hat{N}(i, j)} \end{aligned}$$

$$E(\hat{N}(i, j)) = 0$$

$$E(\hat{N}^2(i, j)) = \frac{1}{m^2} m \sigma^2 = \frac{\sigma^2}{m} \Rightarrow \hat{N}(i, j) \sim N(0, \frac{\sigma}{\sqrt{m}})$$

(Camps)

# Gaussian noise



$$f(x, y) = \overbrace{\hat{f}(x, y)}^{\text{Ideal Image}} + \overbrace{\eta(x, y)}^{\text{Noise process}}$$

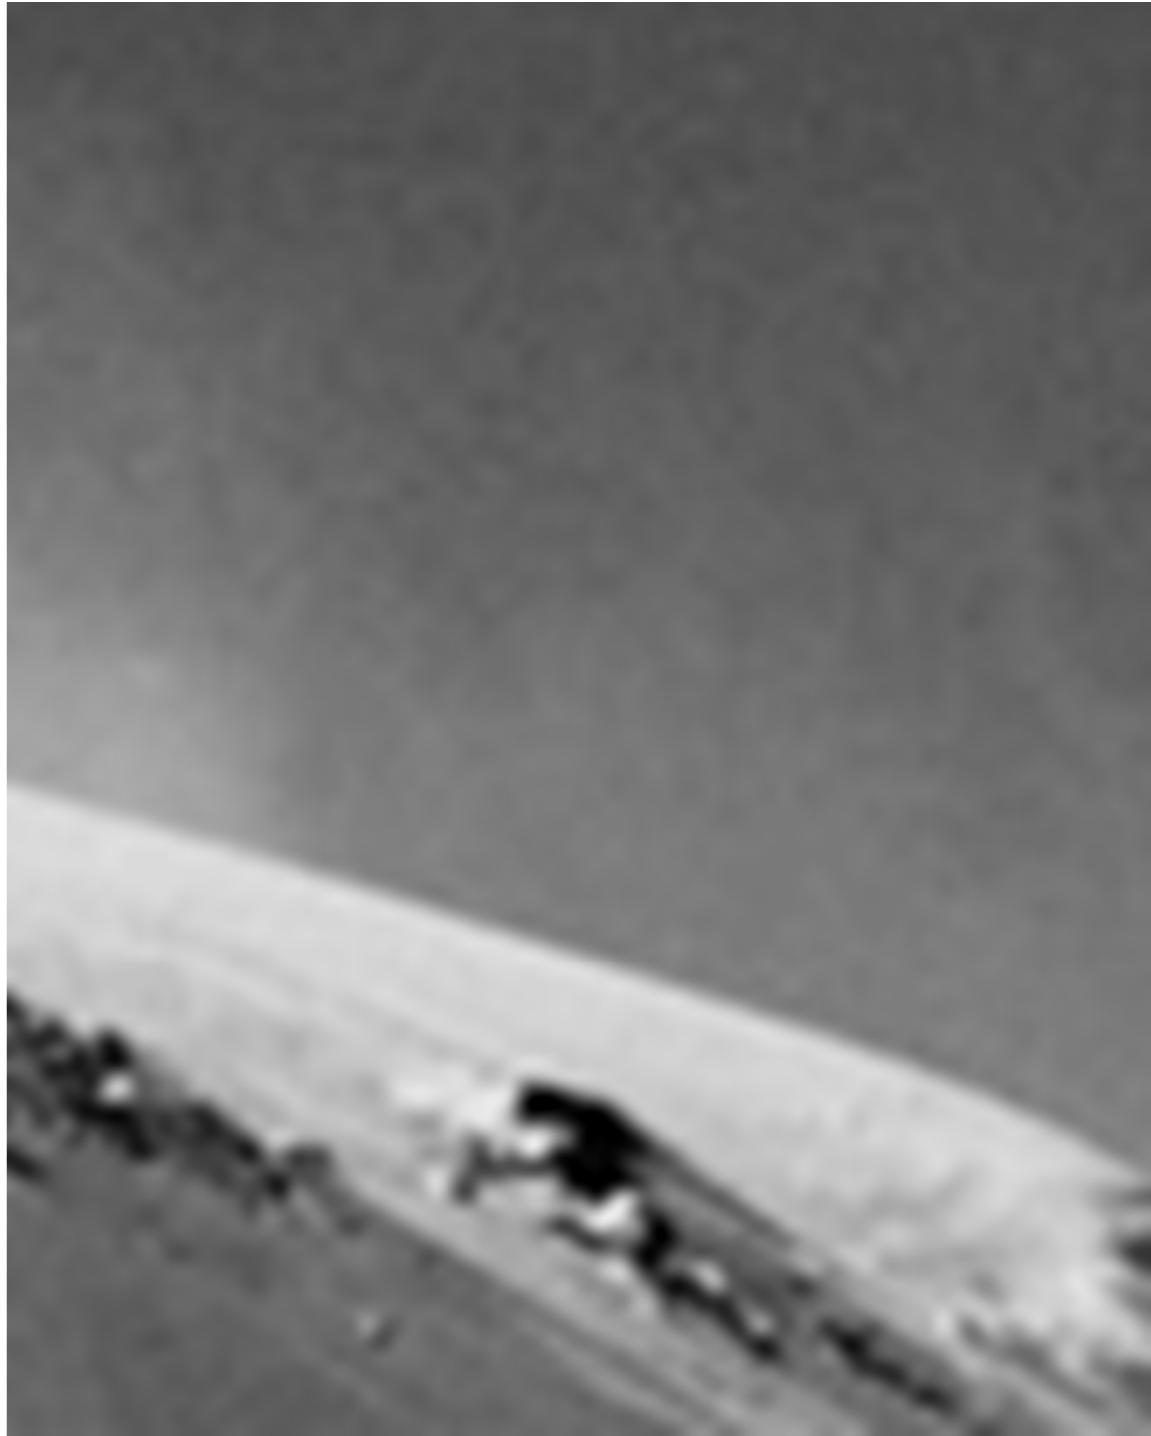
Gaussian i.i.d. ("white") noise:  
 $\eta(x, y) \sim \mathcal{N}(\mu, \sigma)$

```
>> noise = randn(size(im)).*sigma;  
  
>> output = im + noise;
```

Effect of  
sigma on  
Gaussian  
noise:

This shows  
the noise  
values added  
to the raw  
intensities of  
an image.

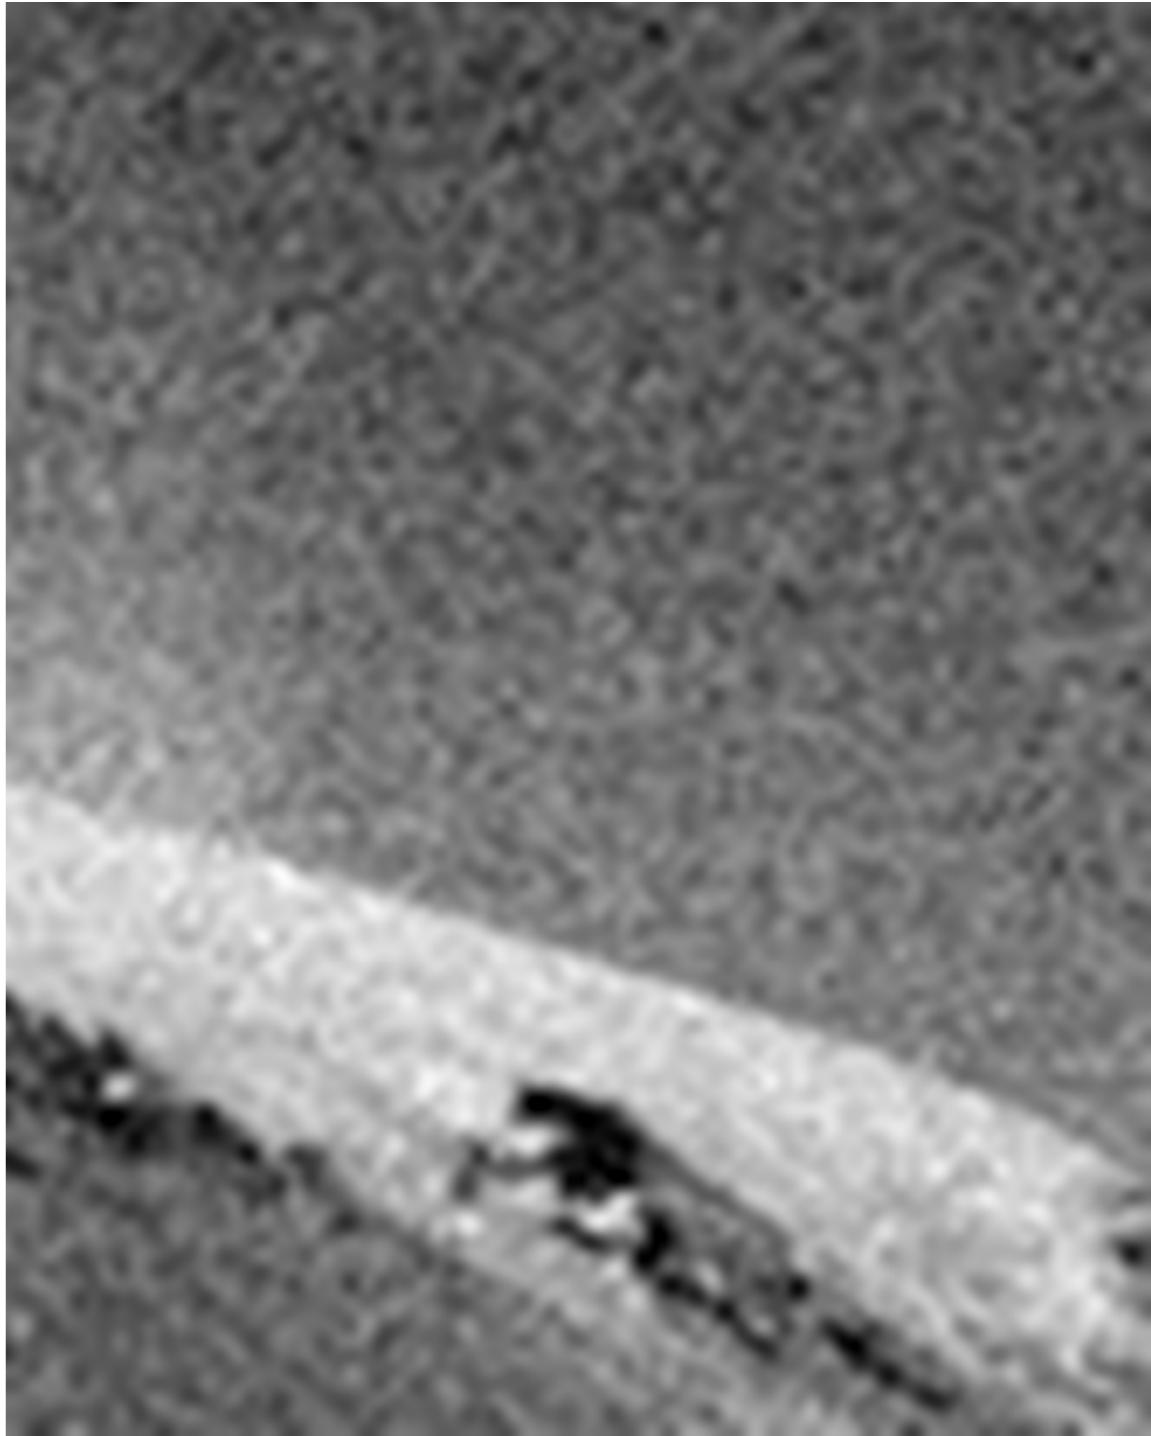
sigma=1



## Effect of sigma on Gaussian noise

This shows  
the noise  
values added  
to the raw  
intensities of  
an image.

sigma=16

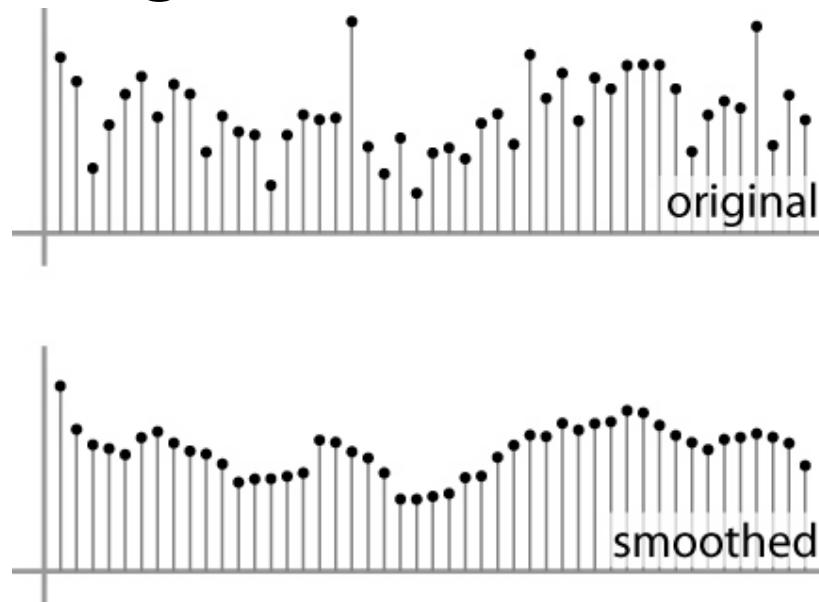


# First attempt at a solution

- Let's replace each pixel with an average of all the values in its neighborhood
- Assumptions:
  - Expect pixels to be like their neighbors
  - Expect noise processes to be independent from pixel to pixel

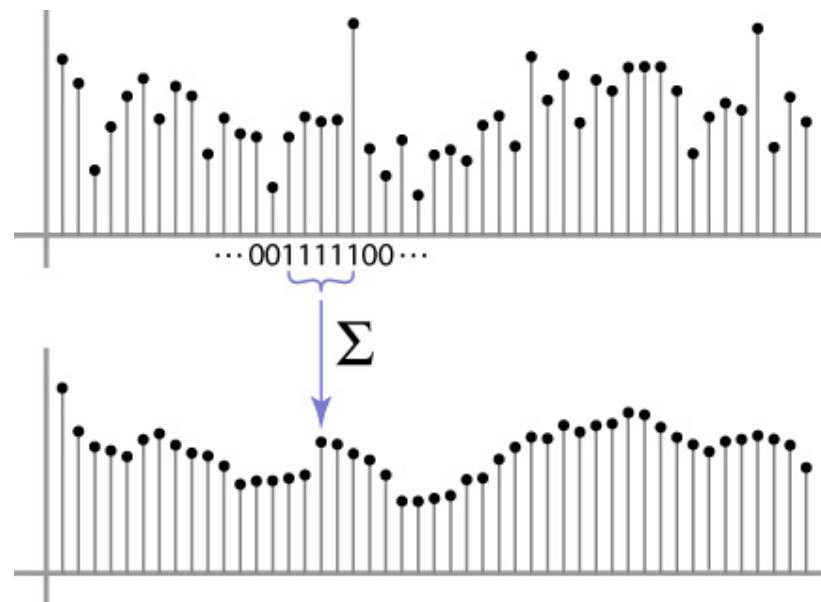
# First attempt at a solution

- Let's replace each pixel with an average of all the values in its neighborhood
- Moving average in 1D:



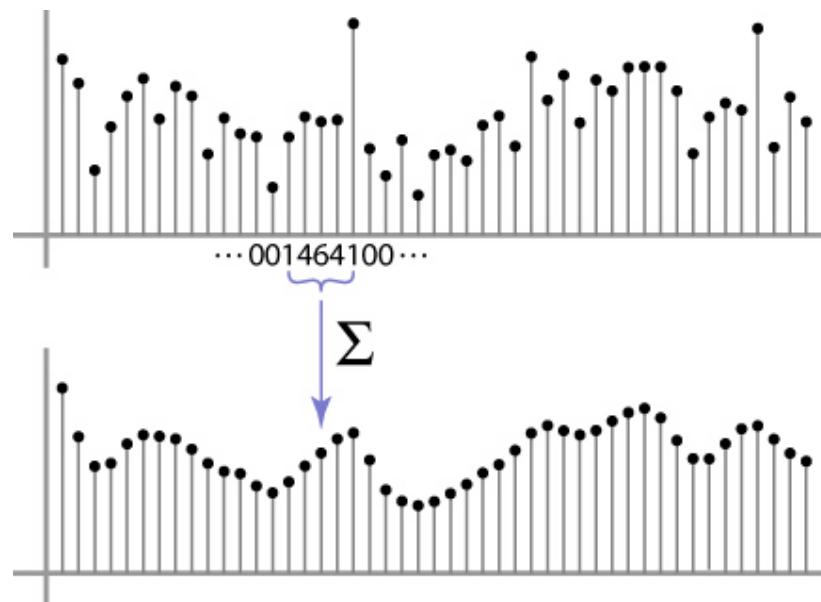
# Weighted Moving Average

- Can add weights to our moving average
- *Weights*  $[1, 1, 1, 1, 1] / 5$

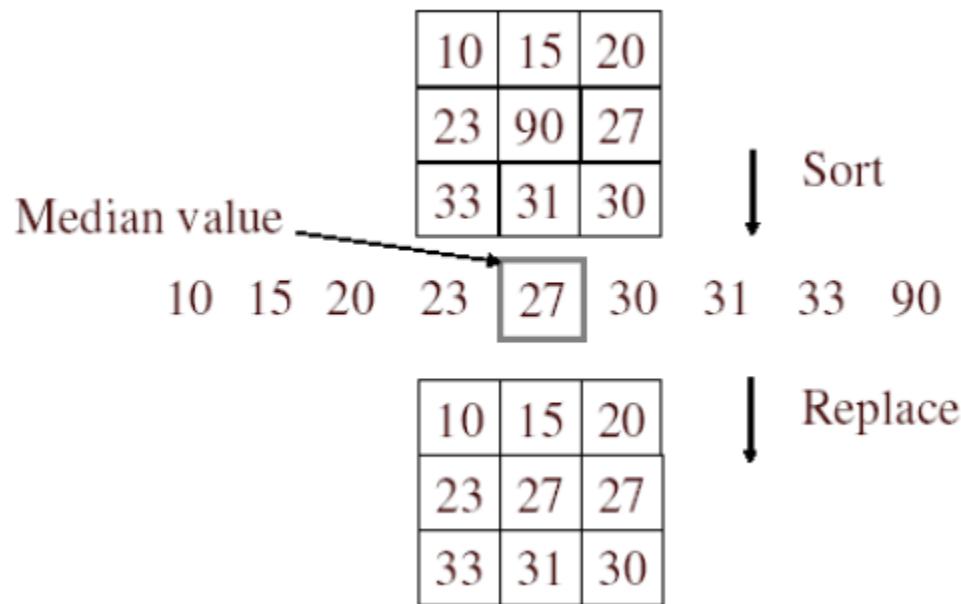


# Weighted Moving Average

- Non-uniform weights  $[1, 4, 6, 4, 1] / 16$



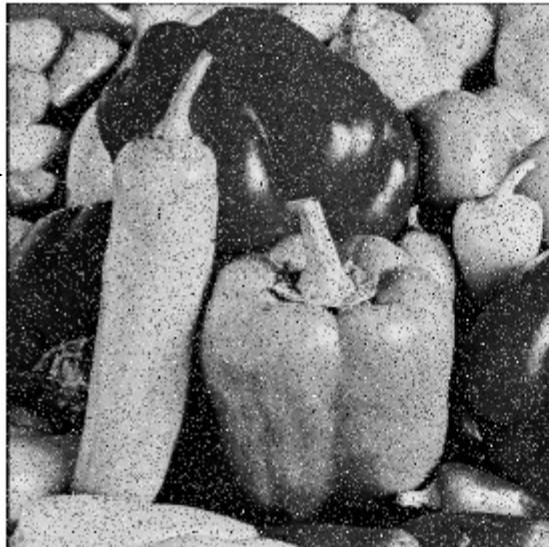
# Median filter



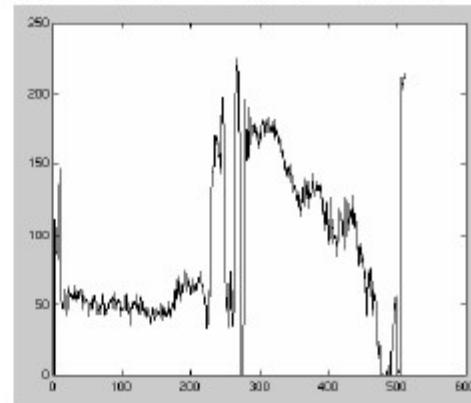
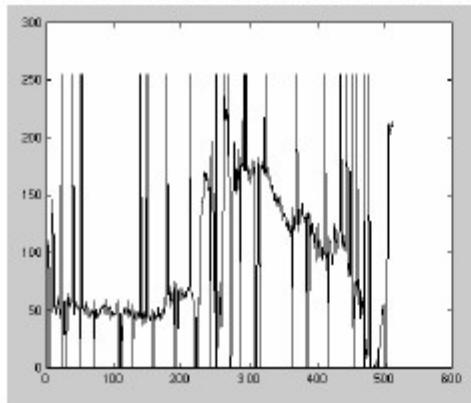
- No new pixel values introduced
- Removes spikes: good for impulse, salt & pepper noise

# Median filter

Salt and  
pepper noise



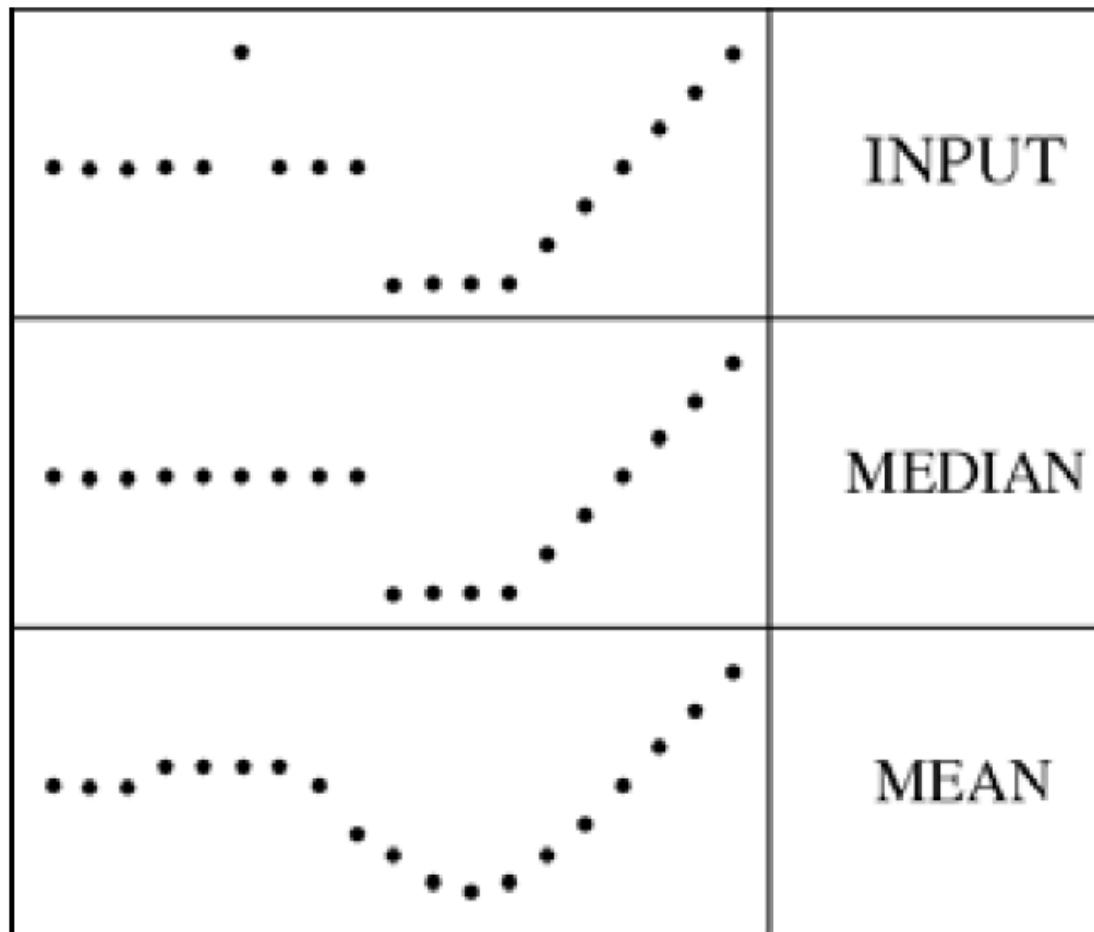
Median  
filtered



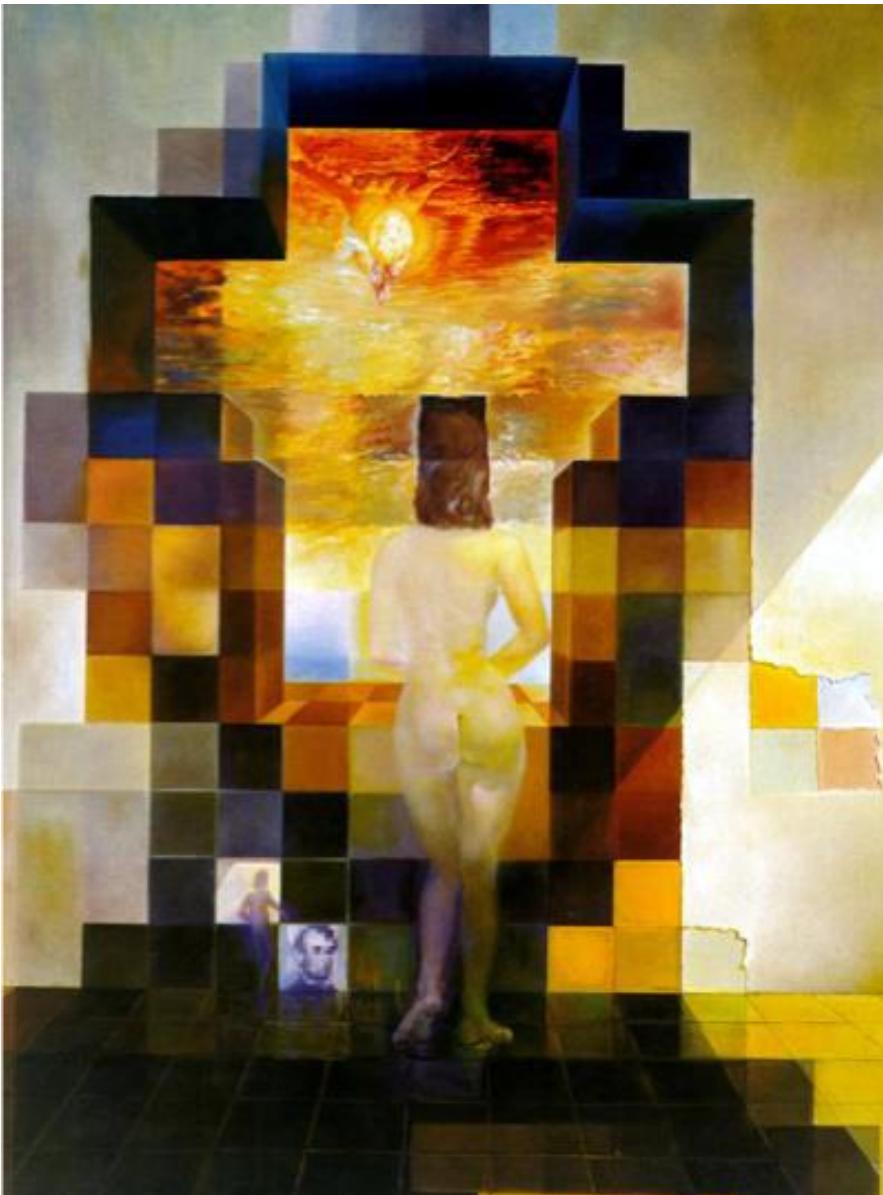
Plots of a row of the image

# Median filter

- Median filter is edge preserving



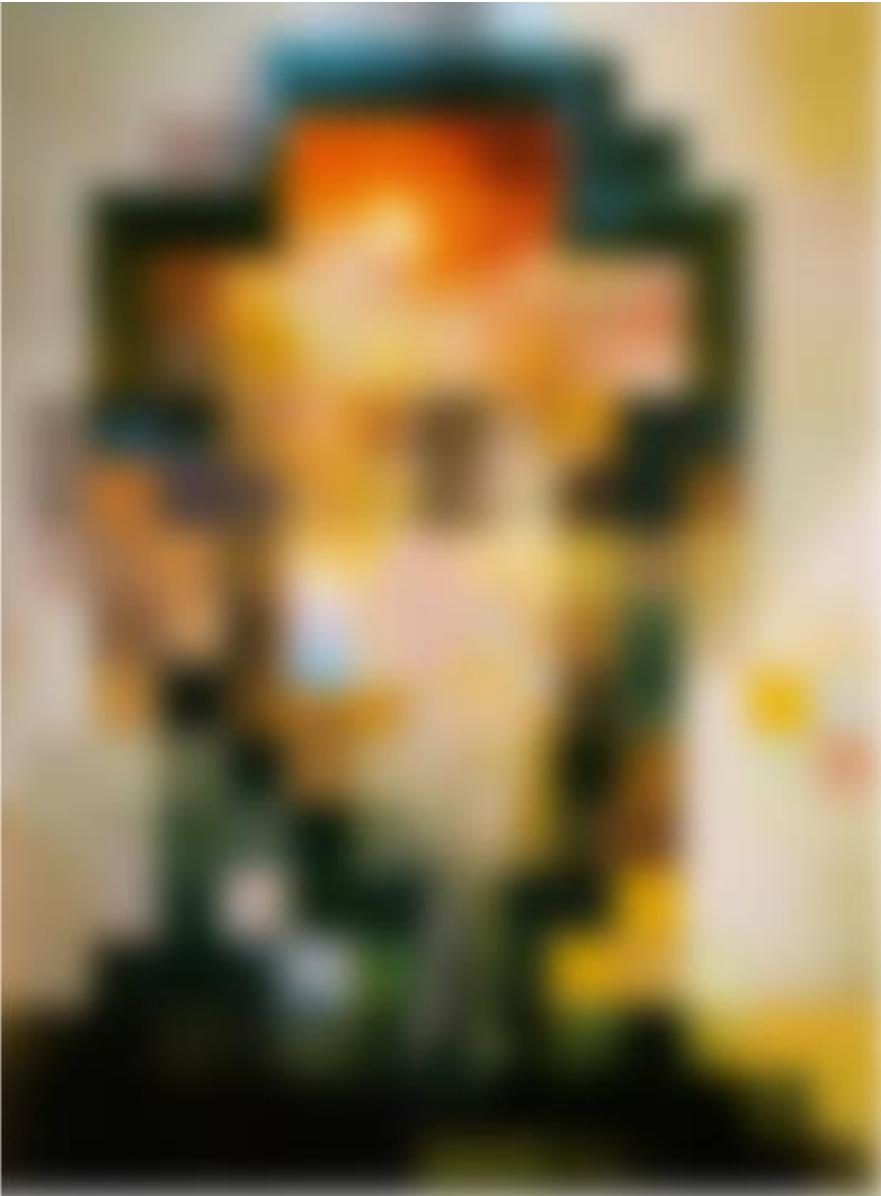
# Frequency-domain filtering in human vision



*Gala Contemplating the Mediterranean  
Sea Which at Twenty Meters Becomes  
the Portrait of Abraham Lincoln  
(Homage to Rothko)*

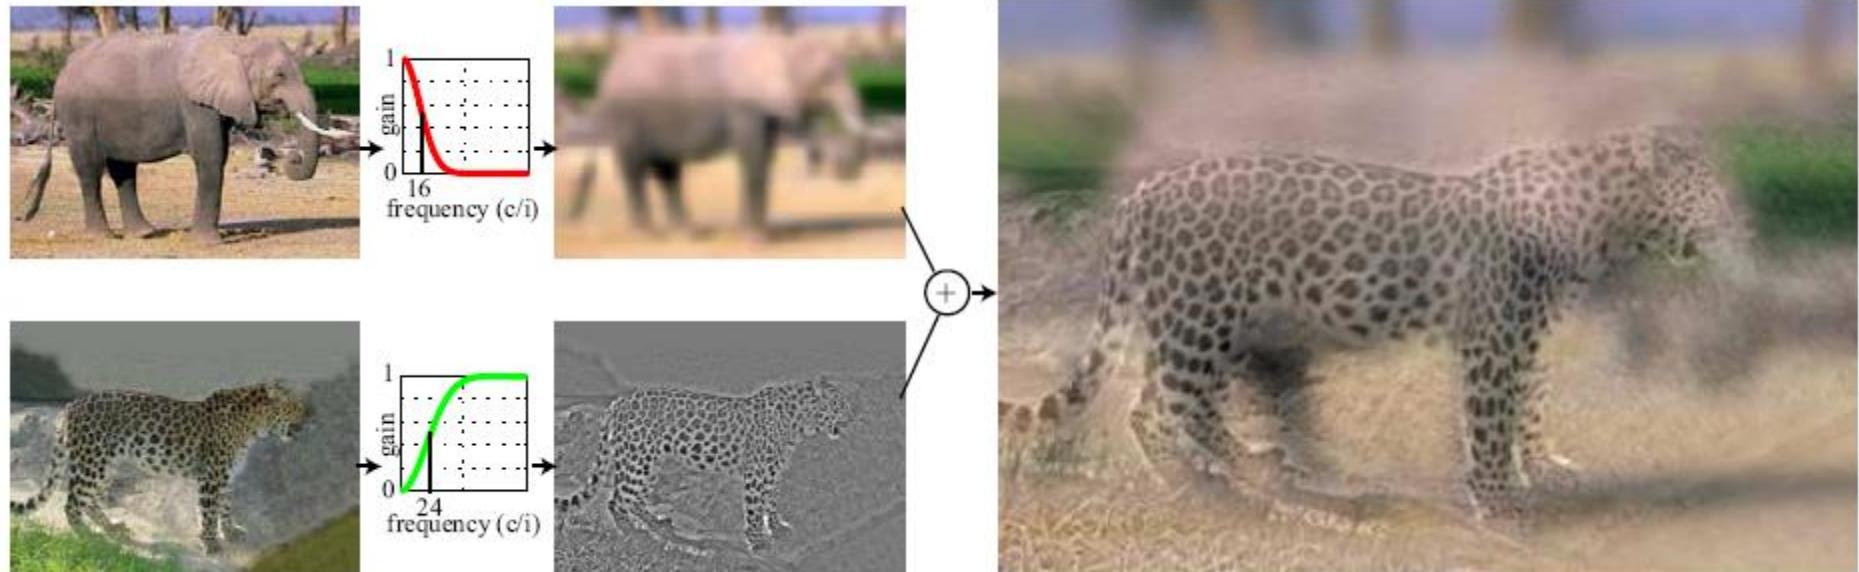
Salvador Dalí, 1976

# Frequency-domain filtering in human vision



Low-pass filtered version

# Hybrid Images



- A. Oliva, A. Torralba, P.G. Schyns,  
“Hybrid Images,” SIGGRAPH 2006

# Still used extensively



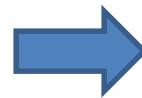
input image



foreground details enhanced, background details reduced

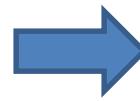
user-provided mask

# Why does a lower resolution image still make sense to us? What information do we lose?

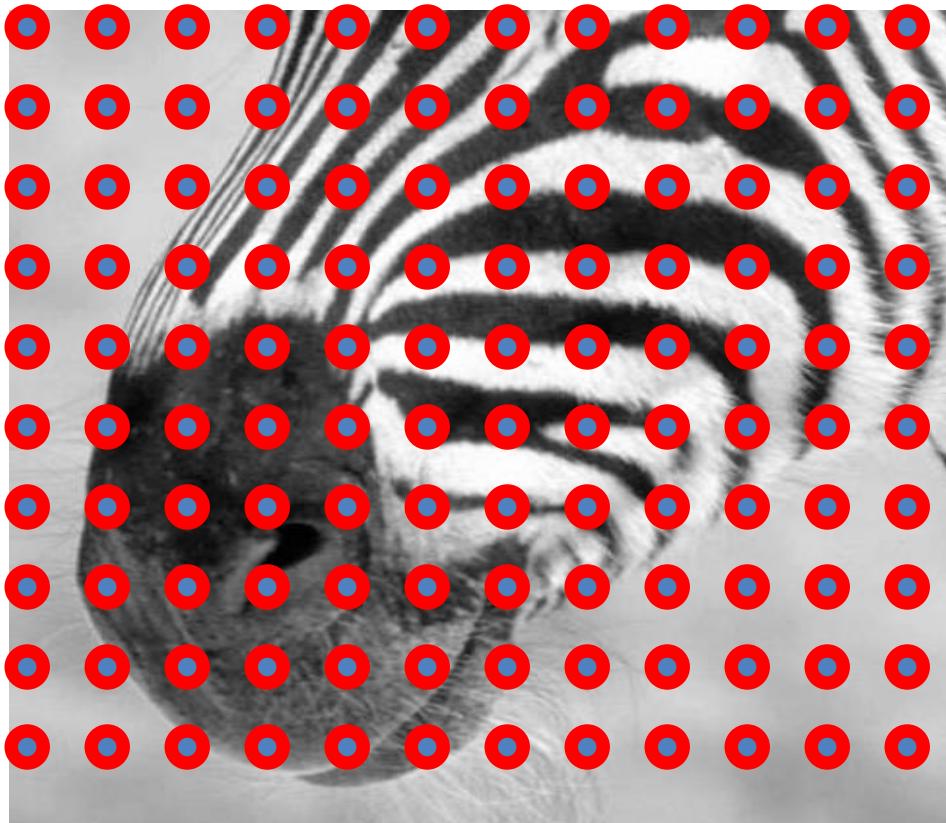


# Sampling

**Why does a lower resolution image still make sense to us? What do we lose?**



# Subsampling by a factor of 2



Throw away every other row and column to create a 1/2 size image

# Algorithm for downsampling by factor of 2

1. Start with image of  $w \times h$
2. Sample every other pixel
  - `im_small = image[ ::2:, ::2 ]`
3. Repeat until `im_small` is 1 pixel large.

*Numpy syntax:*  
`::2` -> start at 0,  
end at 'end',  
increase every 2,  
until the end.  
e.g.,  
`0,2,4,6,...,w`

(if  $w$  is not even,  
then this goes to  
 $w-1$ )

# Sampling and aliasing

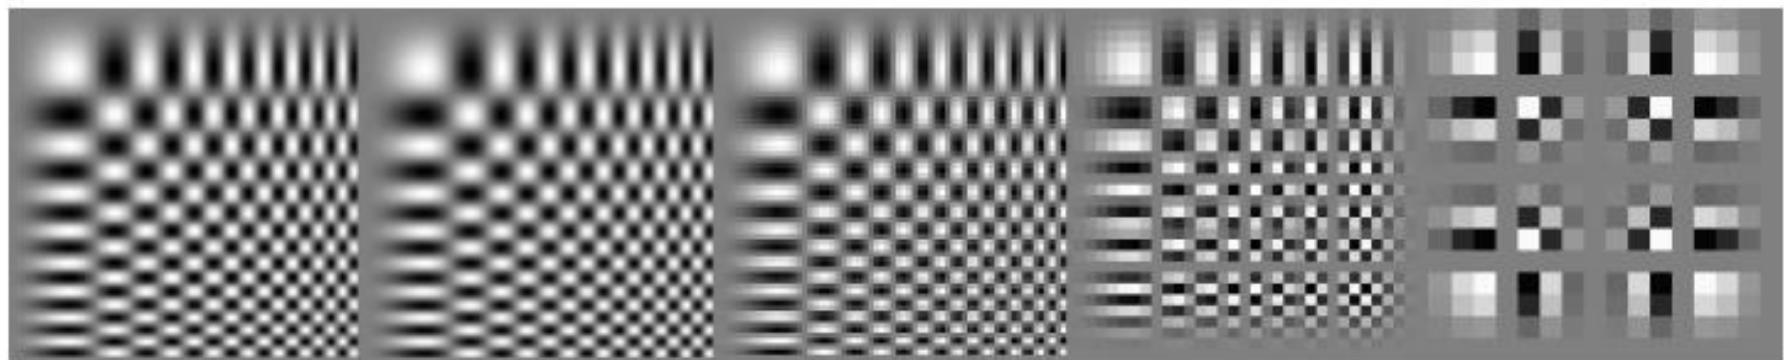
256x256

128x128

64x64

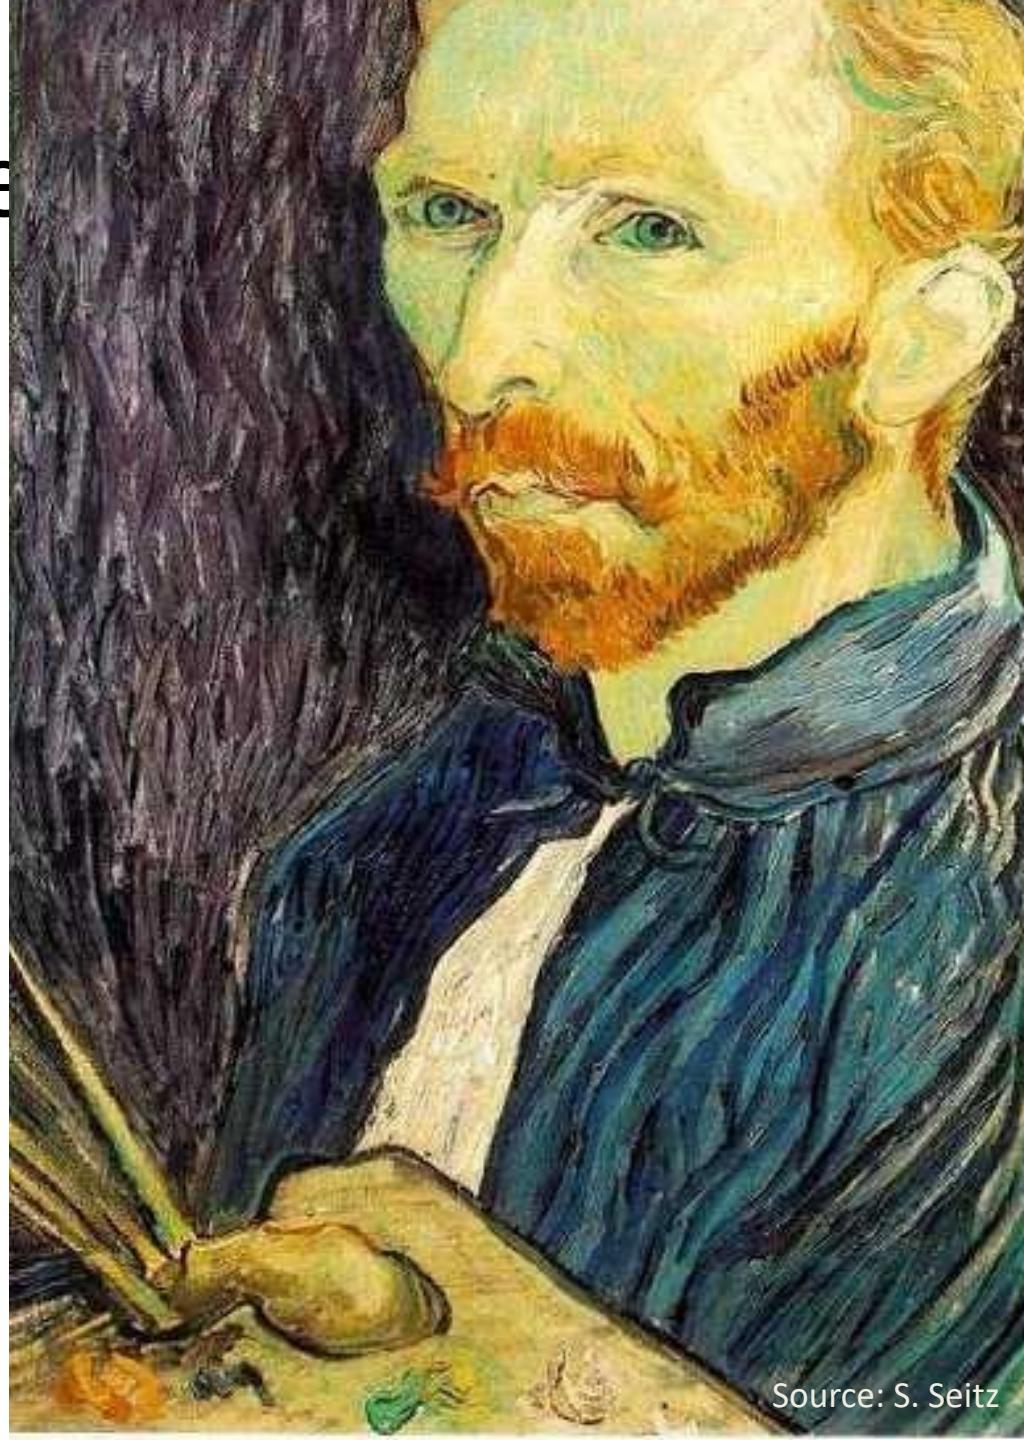
32x32

16x16



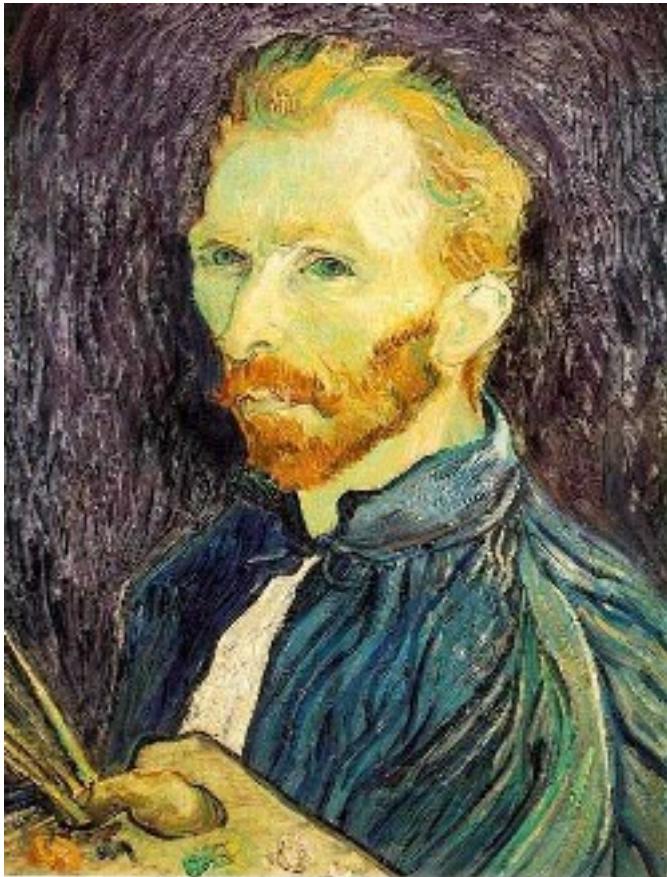
# Image

This image is too big to fit on the screen. How can we generate a half-sized version?



Source: S. Seitz

# Image sub-sampling



1/4

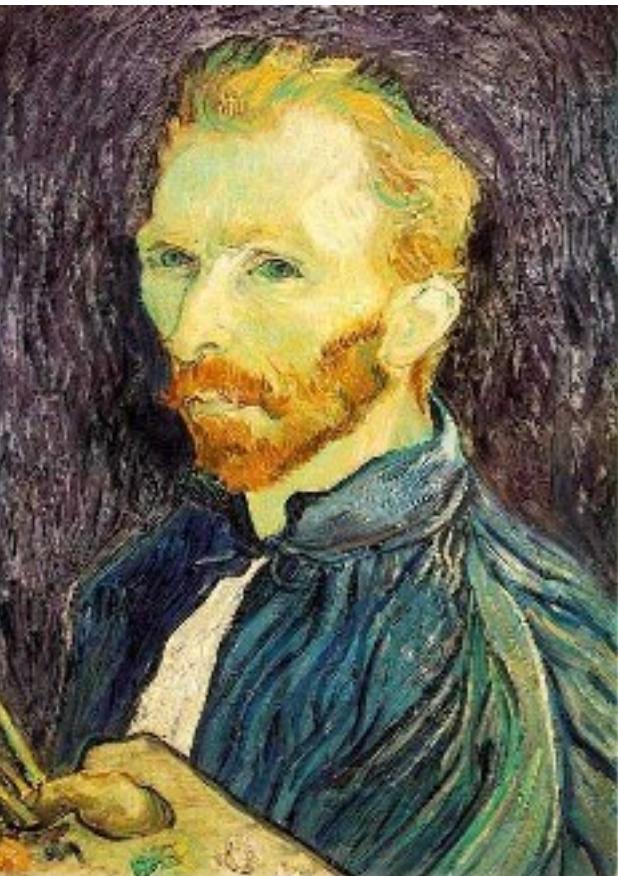


1/8

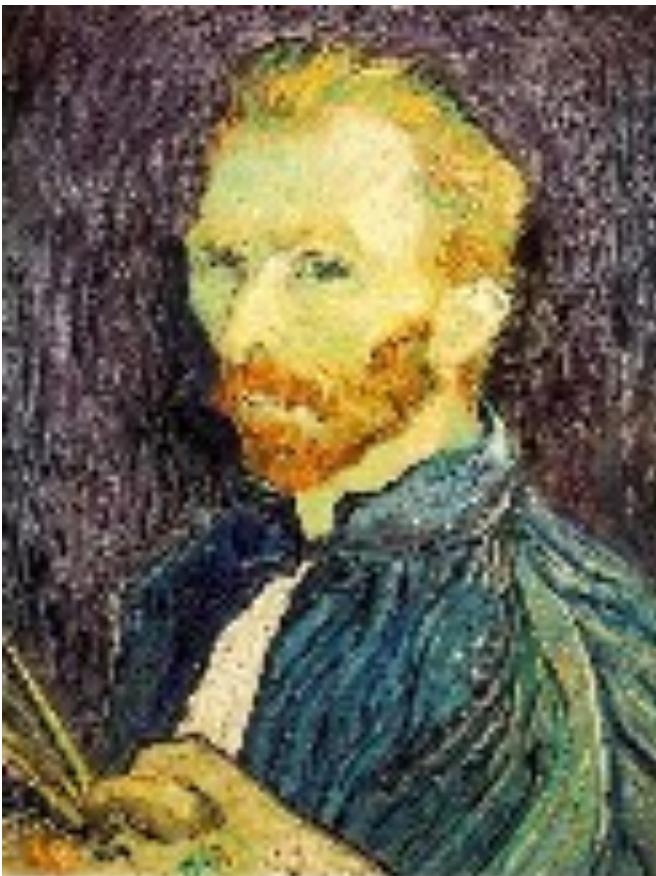
Throw away every other row and column to create a 1/2 size image  
- called *image sub-sampling*

Source: S. Seitz

# Image sub-sampling



1/2



1/4 (2x zoom)

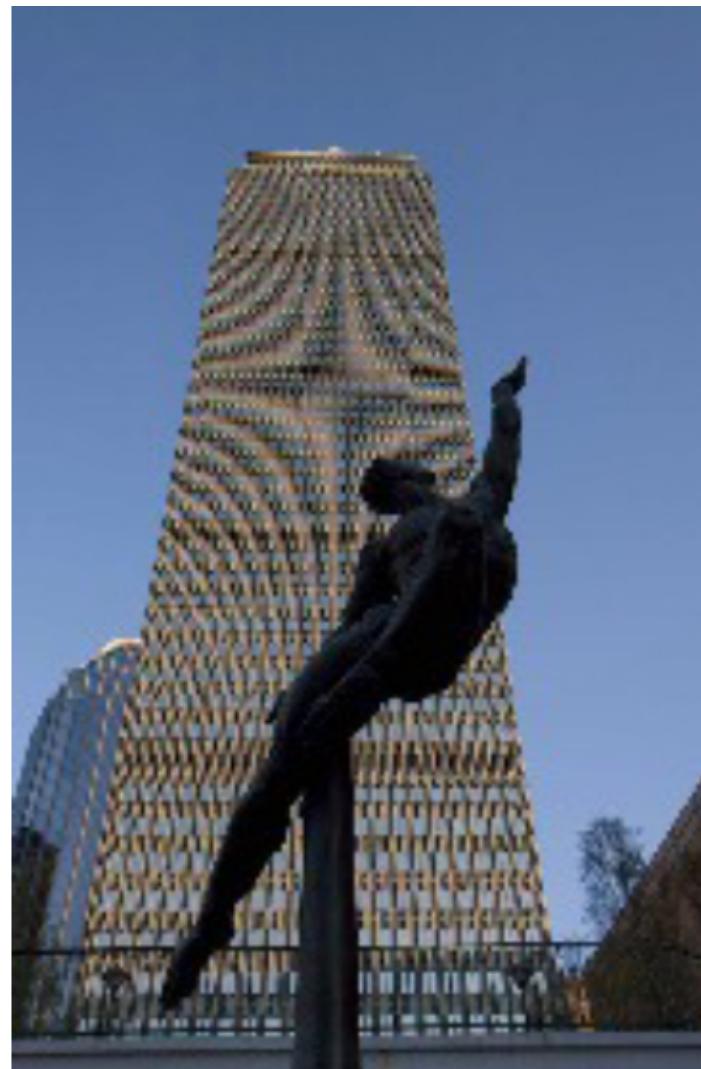


1/8 (4x zoom)

Why does this look so crusty?

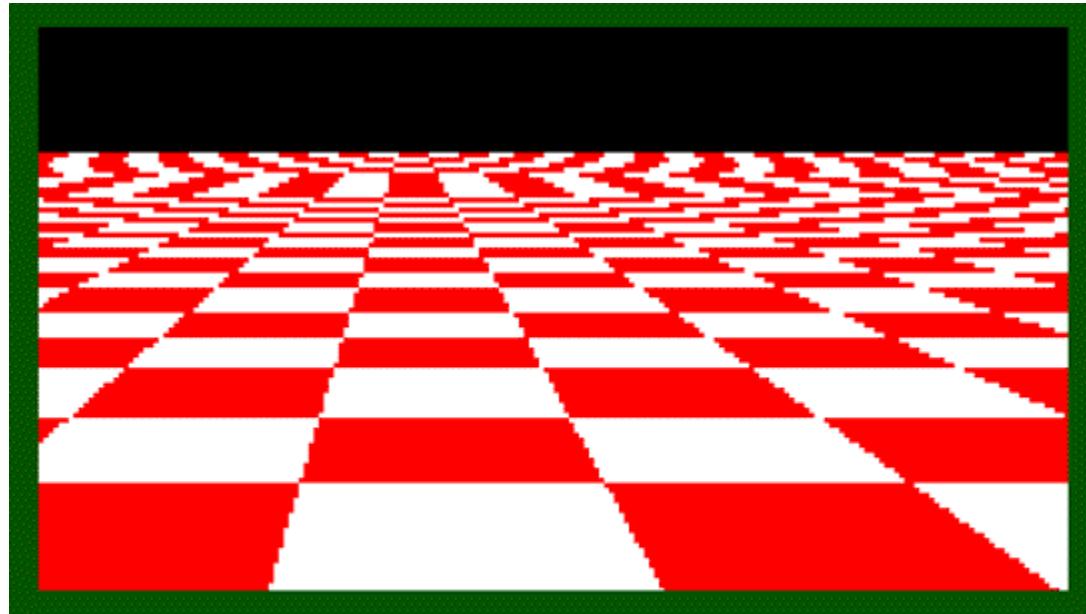
Source: S. Seitz

# Image sub-sampling – another example



Source: F. Durand

# Even worse for synthetic images



# Better sensors

Solutions:

- Sample more often

# Anti-aliasing

Solutions:

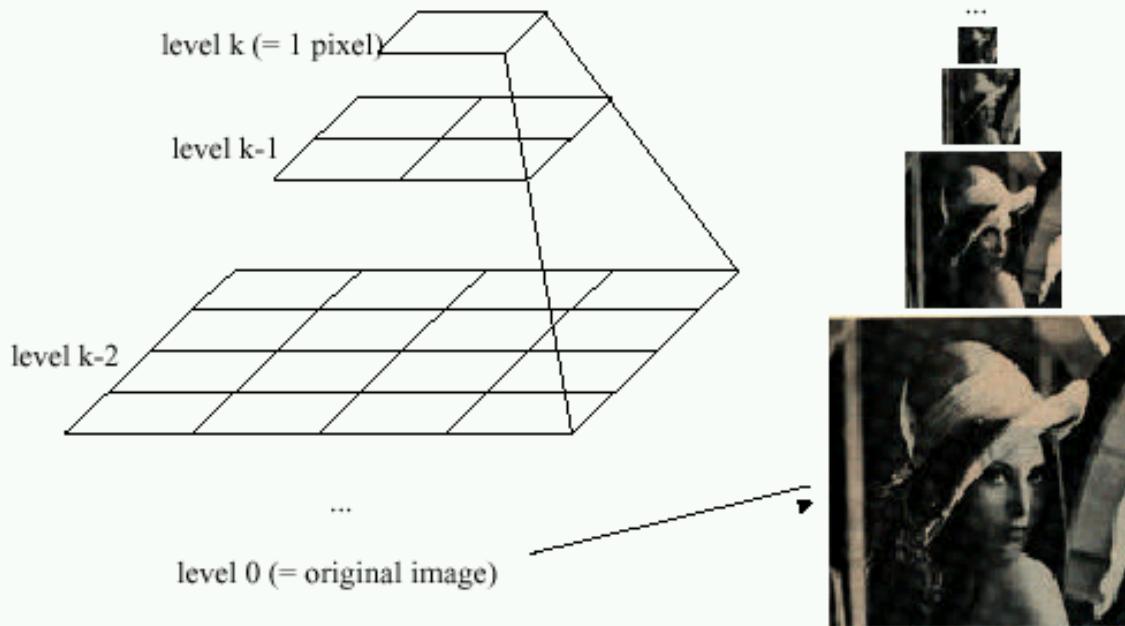
- Sample more often
- Get rid of all frequencies that are greater than half the new sampling frequency
  - Will lose information
  - But it's better than aliasing
  - Apply a smoothing (*low pass*) filter

# Gaussian image pyramid

# Gaussian pyramids

## [Burt and Adelson, 1983]

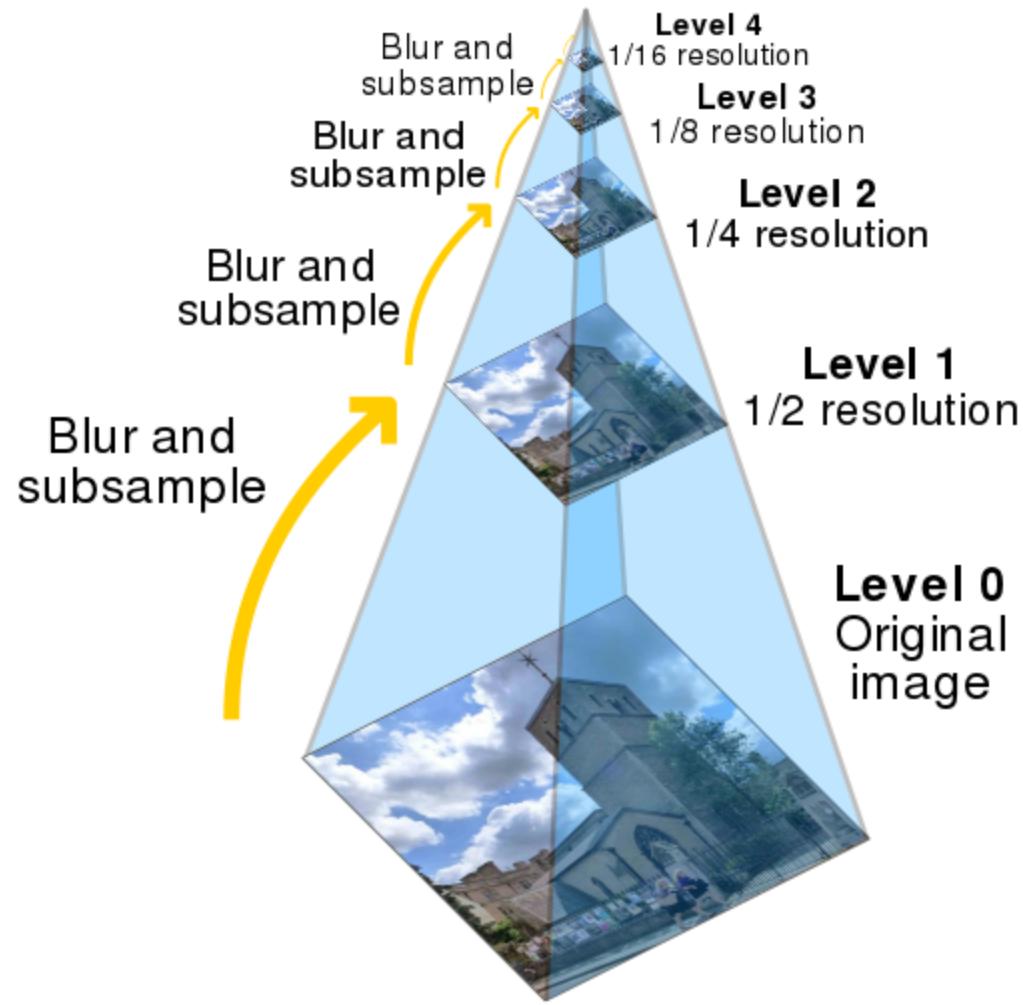
Idea: Represent  $N \times N$  image as a “pyramid” of  $1 \times 1, 2 \times 2, 4 \times 4, \dots, 2^k \times 2^k$  images (assuming  $N=2^k$ )

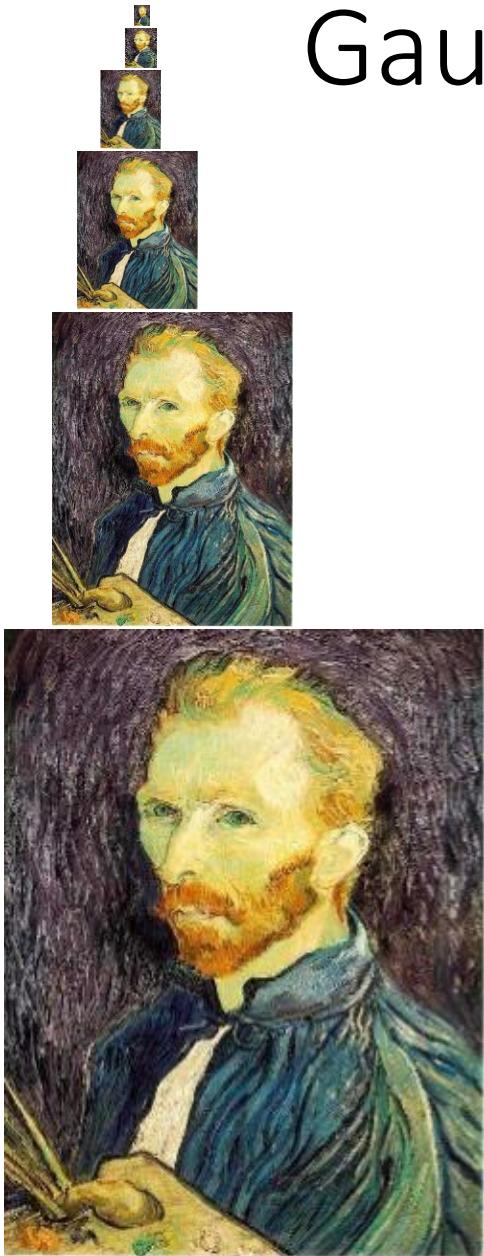


- In computer graphics, a *mip map* [Williams, 1983]
- A precursor to *wavelet transform*

Gaussian Pyramids have all sorts of applications in computer vision

Source: S. Seitz





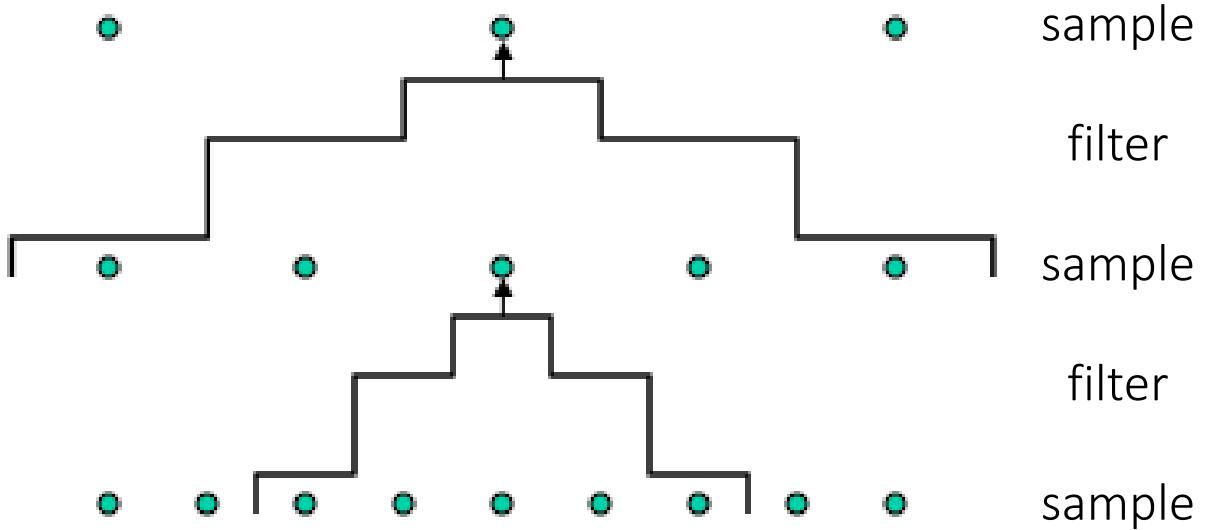
# Gaussian image pyramid

The name of this sequence of subsampled images

# Constructing a Gaussian pyramid

Algorithm

```
repeat:  
    filter  
    subsample  
until min resolution reached
```

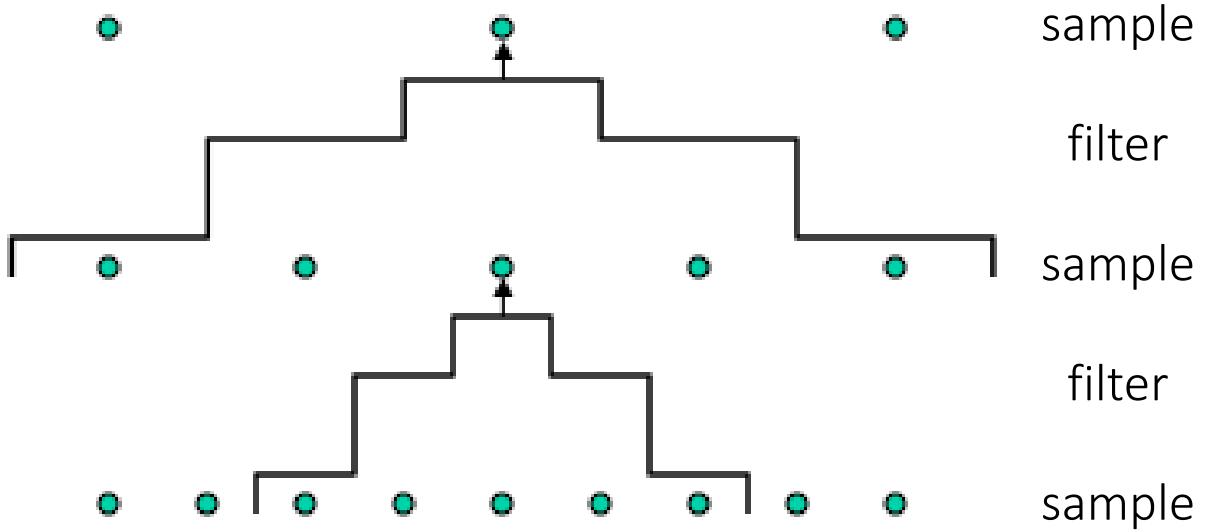


Question: How much bigger than the original image is the whole pyramid?

# Constructing a Gaussian pyramid

Algorithm

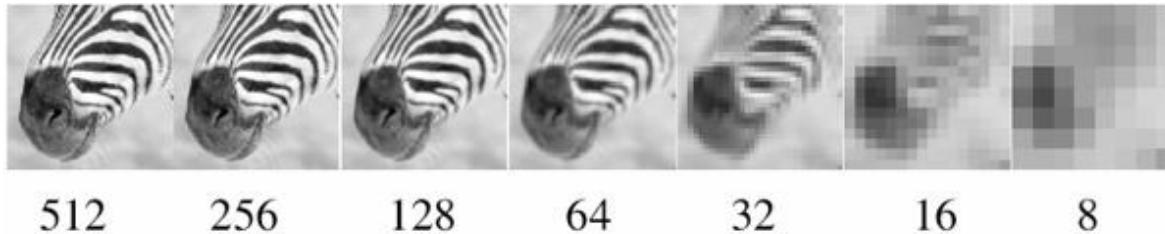
```
repeat:  
    filter  
    subsample  
until min resolution reached
```



Question: How much bigger than the original image is the whole pyramid?

Answer: Just  $4/3$  times the size of the original image! (How did I come up with this number?)

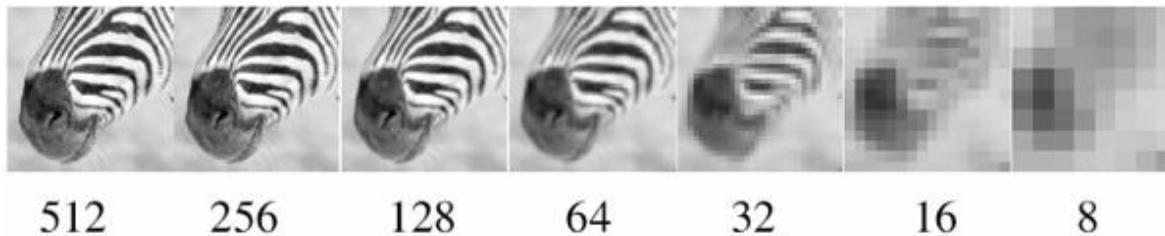
# Some properties of the Gaussian pyramid



What happens to the details of the image?



# Some properties of the Gaussian pyramid



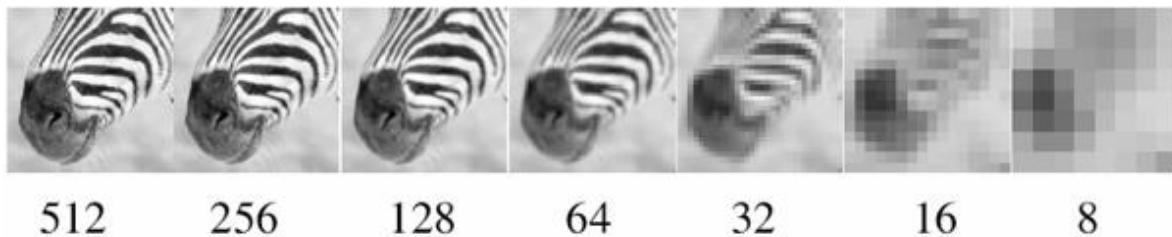
What happens to the details of the image?

- They get smoothed out as we move to higher levels.

What is preserved at the higher levels?



# Some properties of the Gaussian pyramid



What happens to the details of the image?

- They get smoothed out as we move to higher levels.

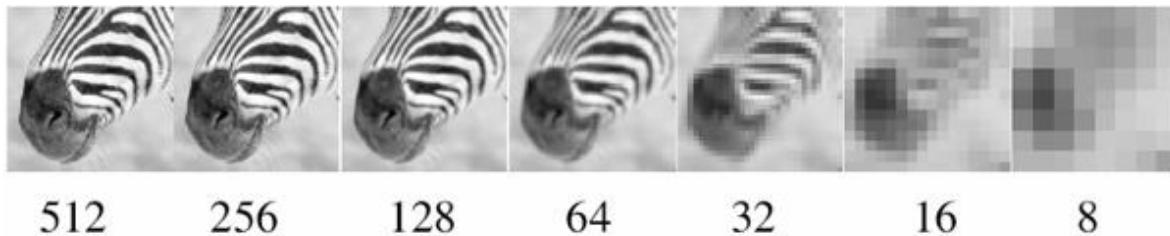


What is preserved at the higher levels?

- Mostly large uniform regions in the original image.

How would you reconstruct the original image from the image at the upper level?

# Some properties of the Gaussian pyramid



What happens to the details of the image?

- They get smoothed out as we move to higher levels.



What is preserved at the higher levels?

- Mostly large uniform regions in the original image.

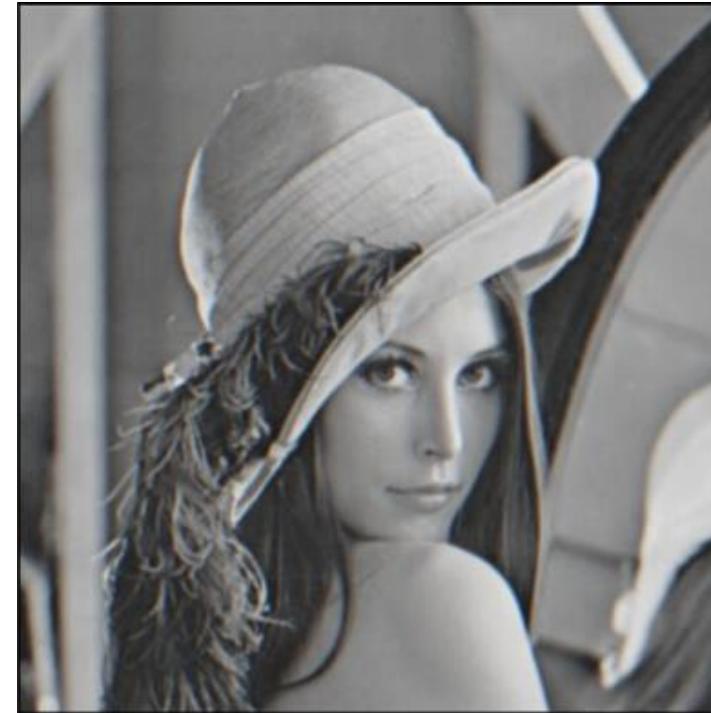
How would you reconstruct the original image from the image at the upper level?

- That's not possible.

# Blurring is lossy



level 0



level 1 (before downsampling)



residual

What does the residual look like?

# Blurring is lossy



level 0



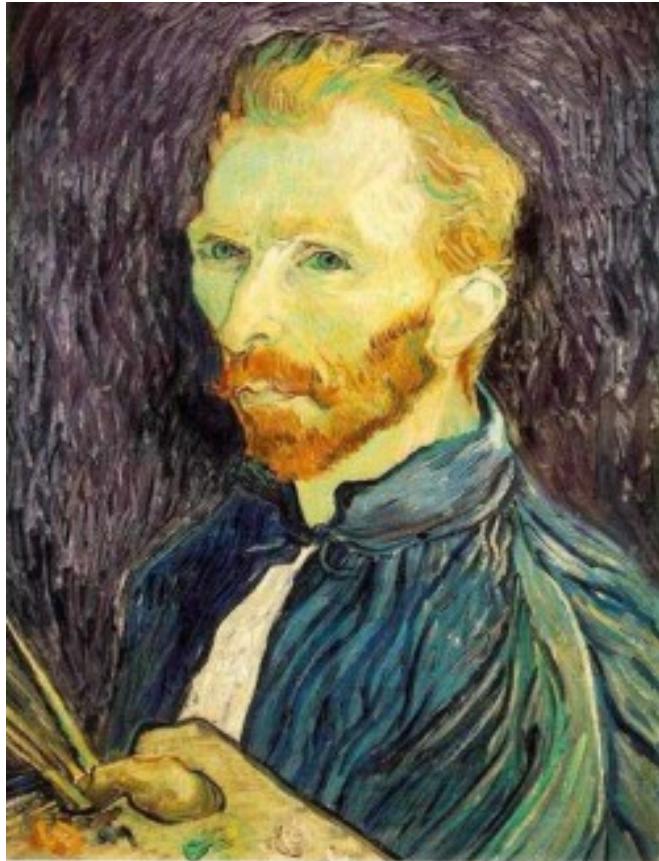
-  
level 1 (before downsampling)



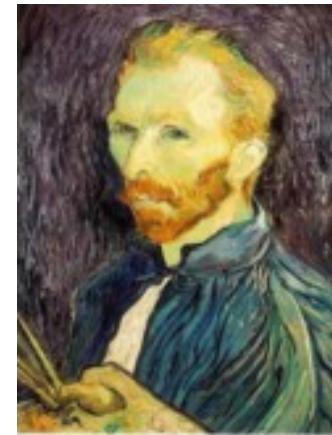
=  
residual

Can we make a pyramid that is lossless?

# Gaussian pre-filtering



Gaussian 1/2



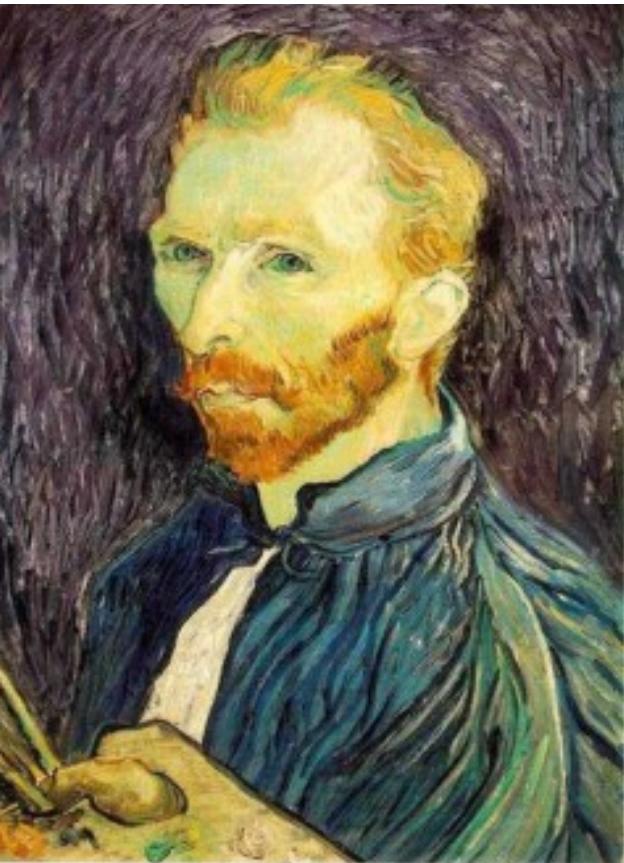
G 1/4



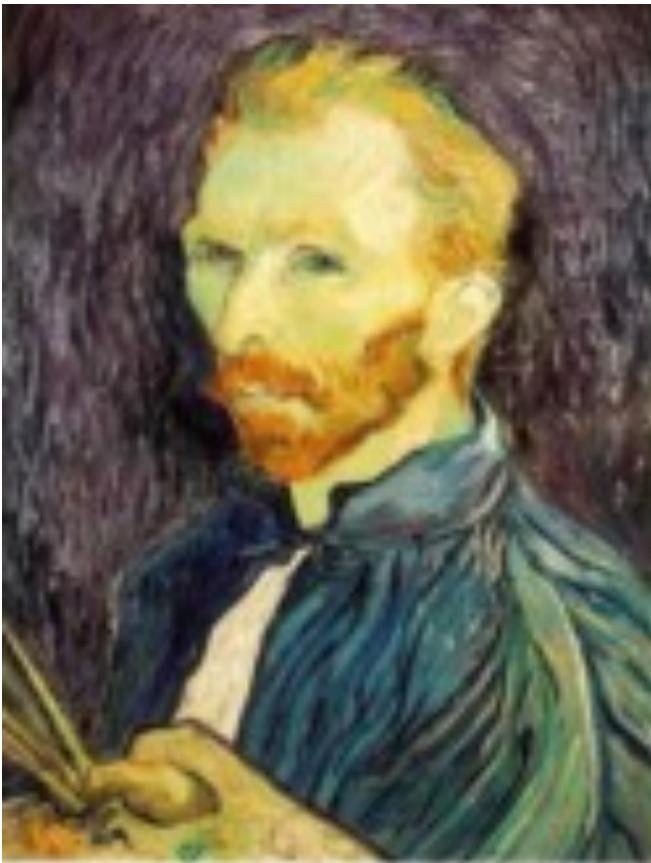
G 1/8

- Solution: filter the image, *then* subsample

# Subsampling with Gaussian pre-filtering



Gaussian 1/2



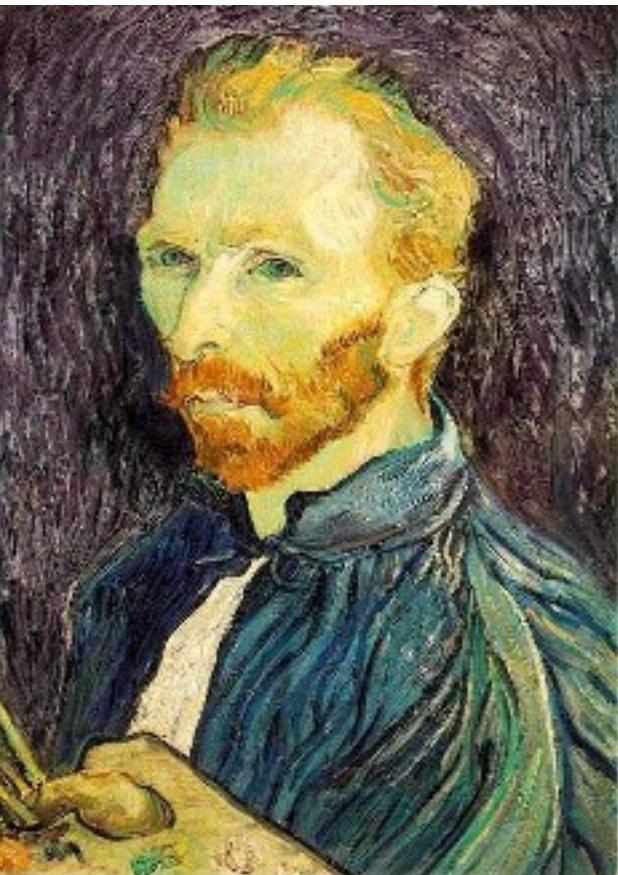
G 1/4



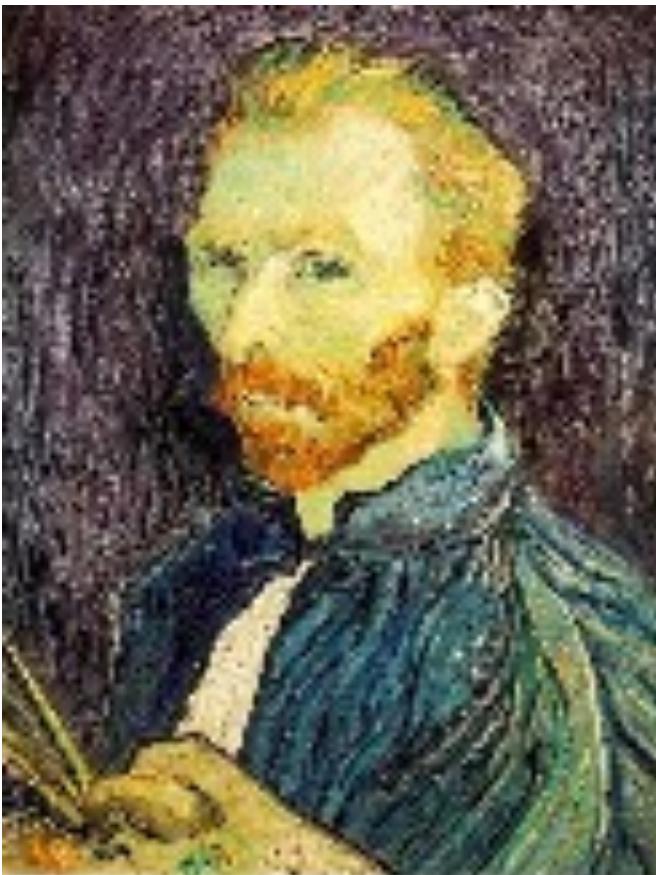
G 1/8

- Solution: filter the image, *then* subsample

# Compare with...



1/2



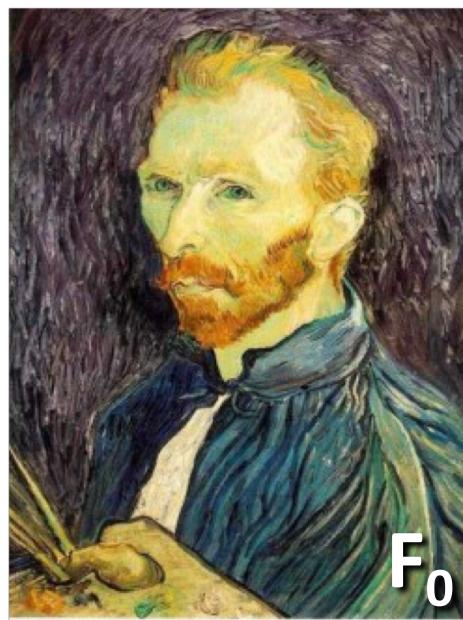
1/4 (2x zoom)



1/8 (4x zoom)

# Gaussian pre-filtering

- Solution: filter the image, *then* subsample



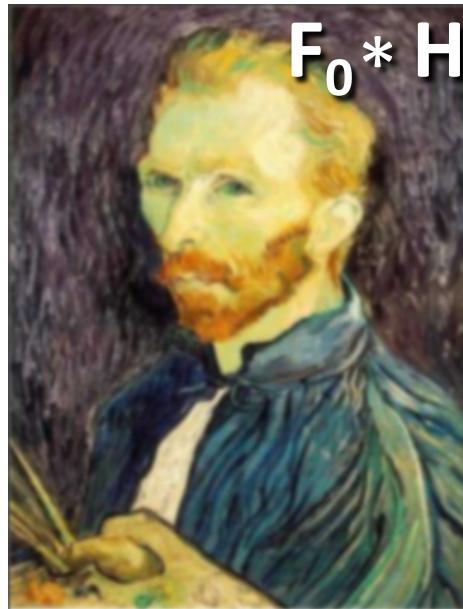
blur

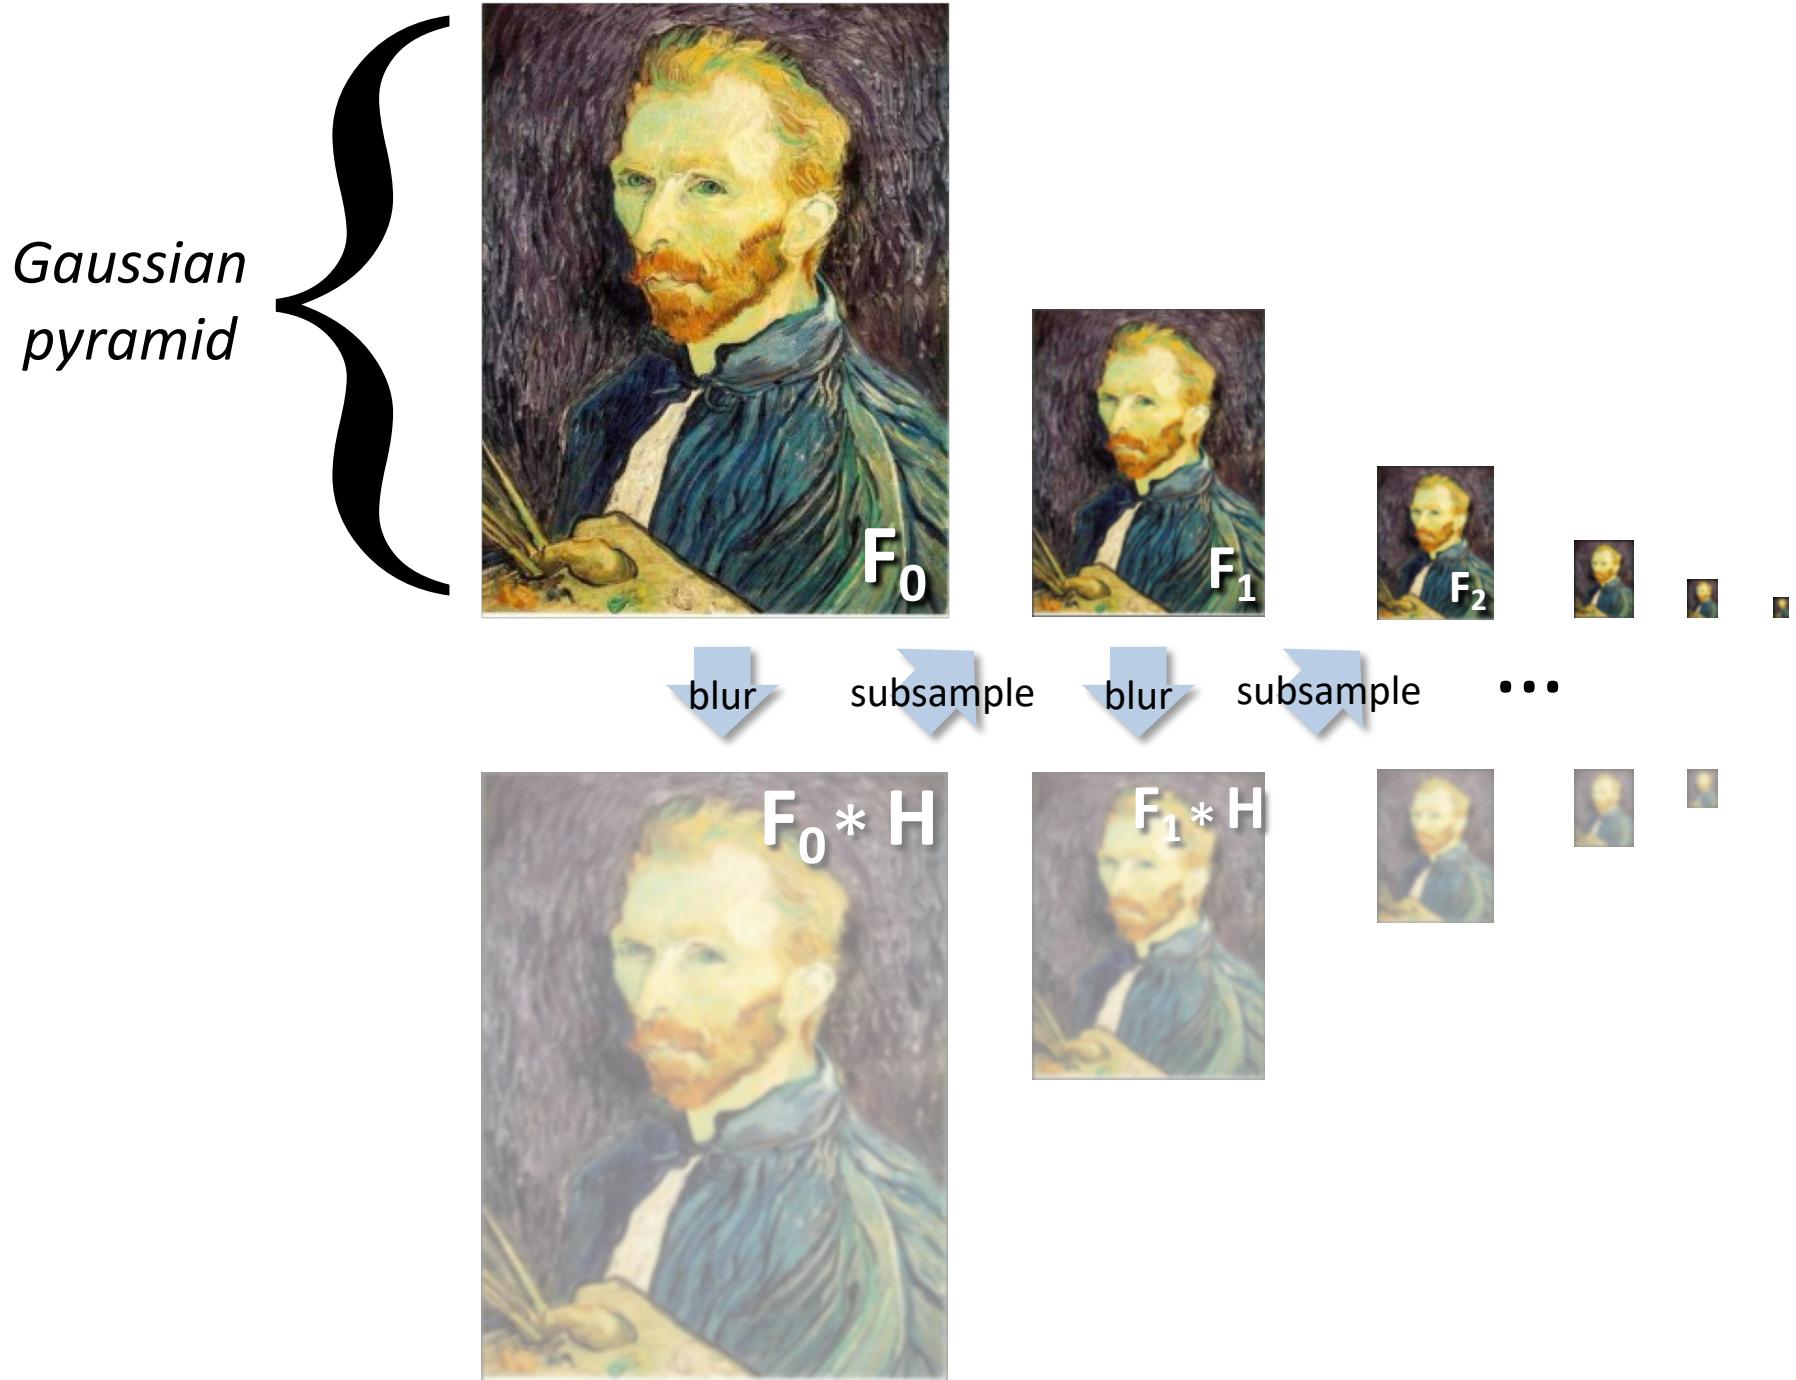
subsample

blur

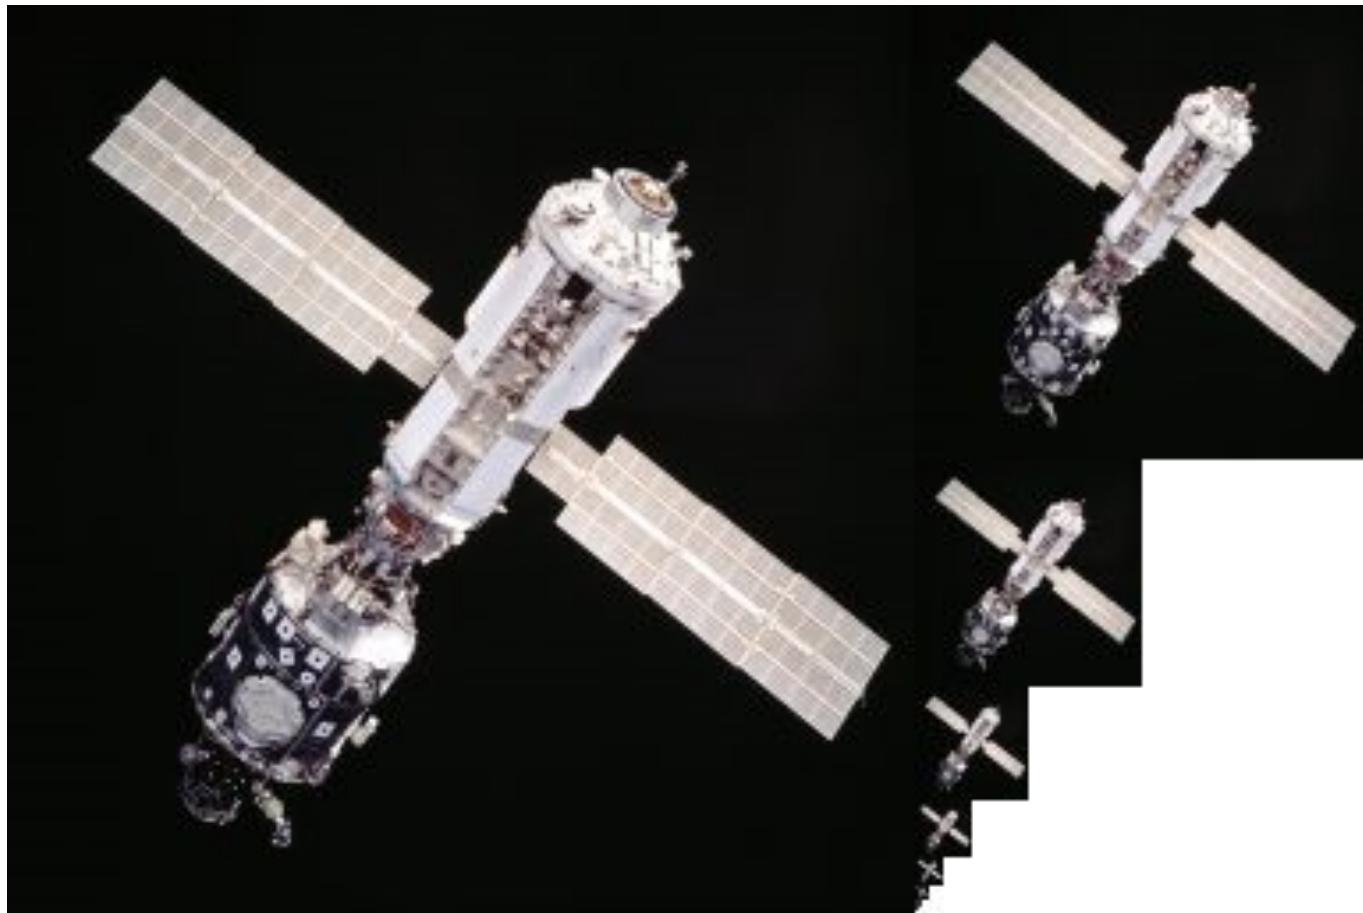
subsample

...





# Gaussian Pyramid



# What are image pyramids used for?

image compression



multi-scale  
texture mapping

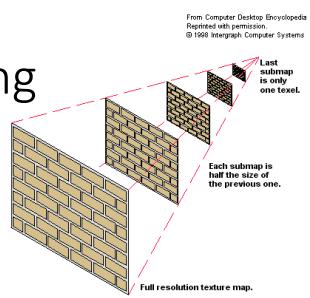
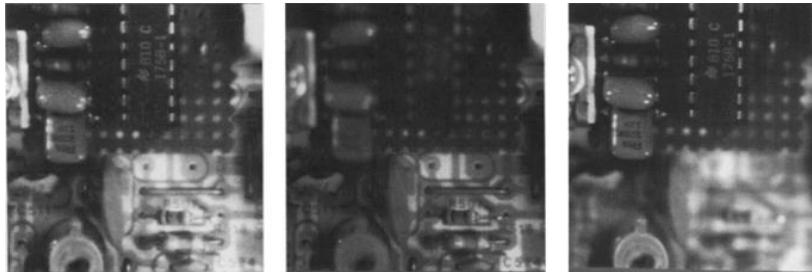


image blending



focal stack compositing



denoising



multi-scale detection



multi-scale registration

