

# Linear shift-invariant image filtering

- Replace each pixel by a *linear* combination of its neighbors (and possibly itself).
- The combination is determined by the filter's *kernel*.
- The same kernel is *shifted* to all pixel locations so that all pixels use the same linear combination of their neighbors.

# Motivation

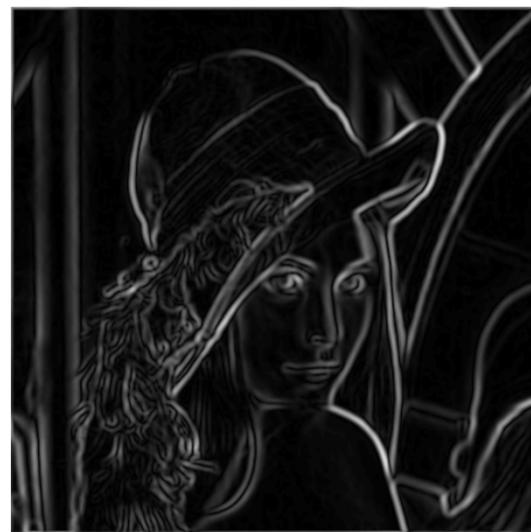
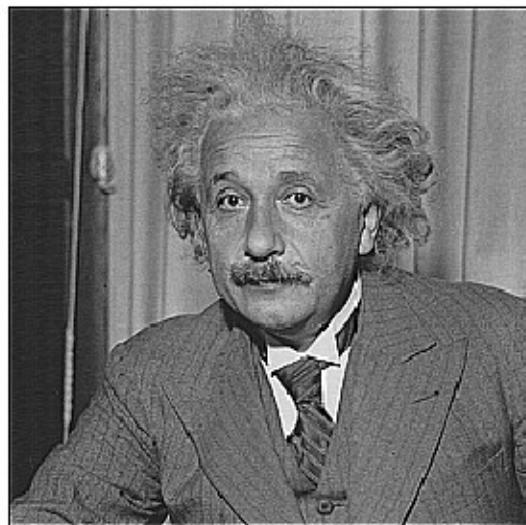
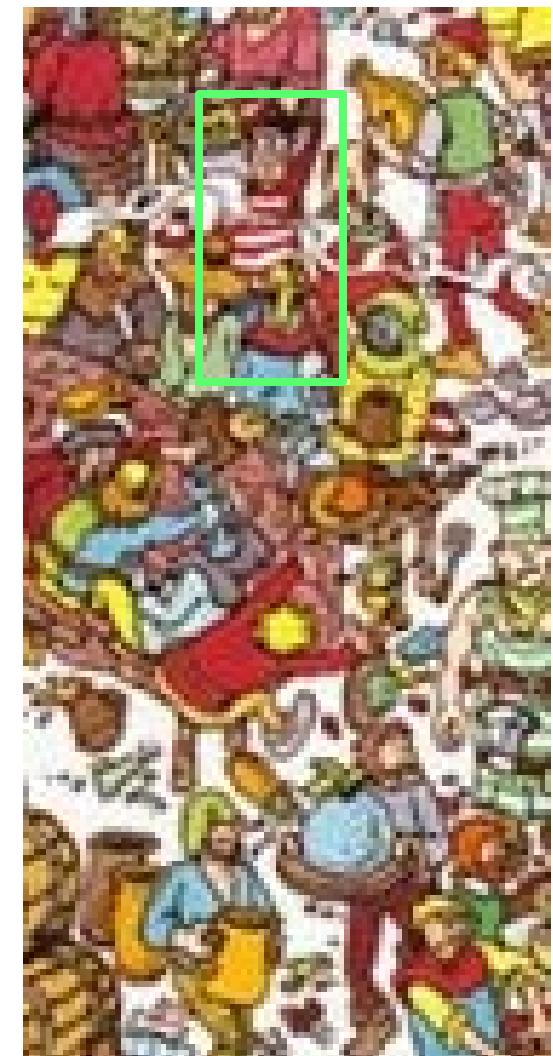
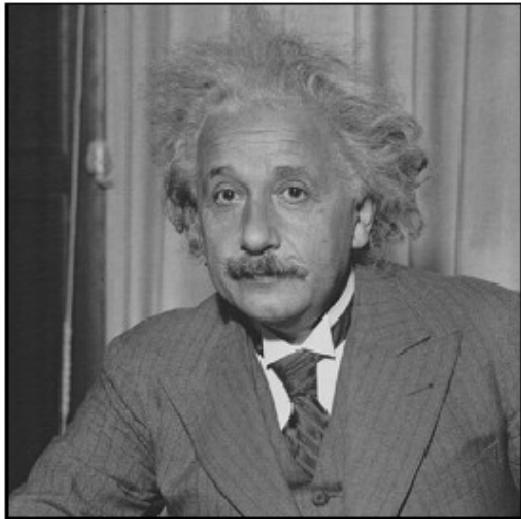
## Filtering:

- Form a new image whose pixels are a combination original pixel values

## Goals:

- Extract useful information from the images
  - Features (edges, corners, blobs...)
- Modify or enhance image properties:
  - super-resolution; in-painting; de-noising

# Image Filters



*Smooth/Sharpen Images...*

*Find edges...*

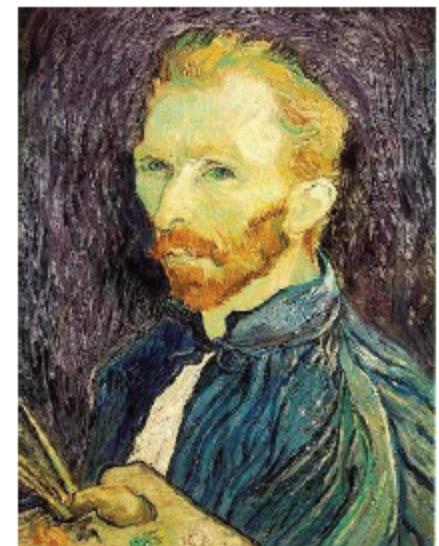
*Find waldo...*

De-noising



Salt and pepper noise

Super-resolution



In-painting



Bertamio et al

# Convolution for 2D discrete signals

Definition of filtering as convolution:

$$(f * g)(x, y) = \sum_{i,j=-\infty}^{\infty} f(i, j)I(x - i, y - j)$$

filtered image      filter      input image      notice the flip

# Convolution for 2D discrete signals

Definition of filtering as convolution:

$$(f * g)(x, y) = \sum_{i,j=-\infty}^{\infty} f(i, j)I(x - i, y - j)$$

filtered image      filter      input image      notice the flip

If the filter  $f(i, j)$  is non-zero only within  $-1 \leq i, j \leq 1$ , then

$$(f * g)(x, y) = \sum_{i,j=-1}^1 f(i, j)I(x - i, y - j)$$

The kernel we saw earlier is the 3x3 matrix representation of  $f(i, j)$ .

# Image filtering

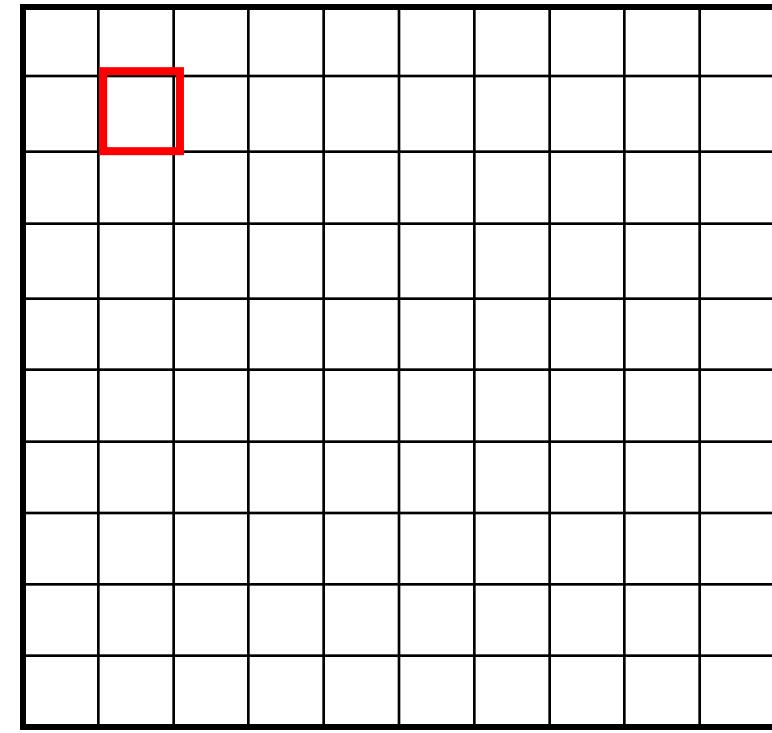
 $I[.,.]$ 

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

 $h[.,.]$ 

$$f[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1



$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l] \quad \begin{matrix} m = 1, n = 1 \\ k, l = [-1, 0, 1] \end{matrix}$$

Credit: S. Seitz

# Image filtering

 $I[.,.]$ 

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

 $h[.,.]$ 

0	10									

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

$$\begin{aligned}m &= 2, n &= 1 \\k, l &= [-1, 0, 1]\end{aligned}$$

Credit: S. Seitz

1	1	1
1	1	1
1	1	1

# Image filtering

 $I[.,.]$ 

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	0	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

 $h[.,.]$ 


$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

$$\begin{aligned}m &= 3, n &= 1 \\k, l &= [-1, 0, 1]\end{aligned}$$

$$f[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

Credit: S. Seitz

# Image filtering

 $I[.,.]$ 

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

 $h[.,.]$ 


$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

$$\begin{aligned}m &= 4, n &= 1 \\k, l &= [-1, 0, 1]\end{aligned}$$

1	1	1
1	1	1
1	1	1

Credit: S. Seitz

# Image filtering

 $I[.,.]$ 

0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0
0	0	0	90	90	90	90	90	0
0	0	0	90	0	90	90	90	0
0	0	0	90	90	90	90	90	0
0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

 $h[.,.]$ 


0    10    20    30    30

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

$$\begin{aligned} m &= 5, n &= 1 \\ k, l &= [-1, 0, 1] \end{aligned}$$

1	1	1
1	1	1
1	1	1

Credit: S. Seitz

# Image filtering

 $I[.,.]$ 

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

 $h[.,.]$ 

	0	10	20	30	30					

$$h[m, n] = \sum_{k,l} f[k, l] I[m+k, n+l]$$

$$\begin{aligned} m &= 4, n &= 6 \\ k, l &= [-1, 0, 1] \end{aligned}$$

1	1	1
1	1	1
1	1	1

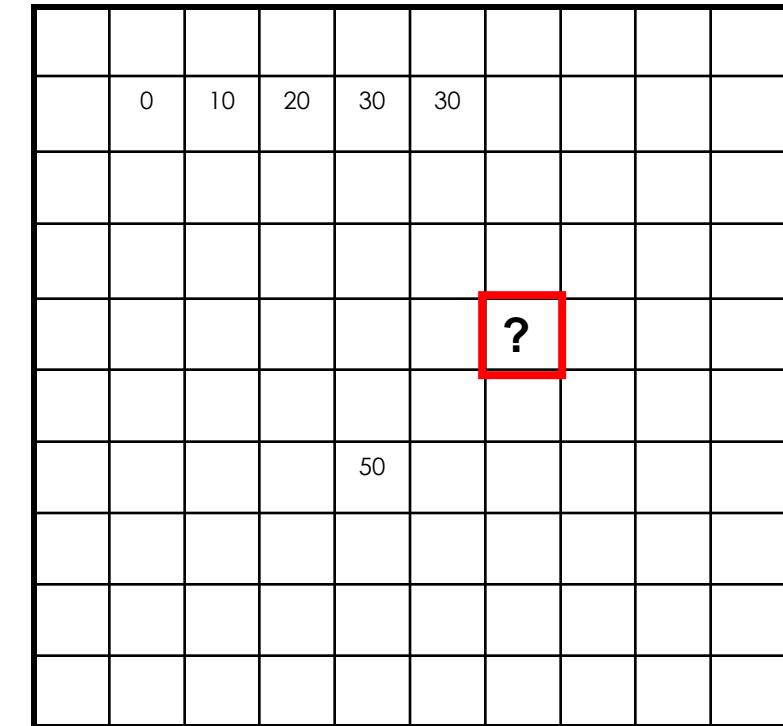
Credit: S. Seitz

# Image filtering

$$I[.,.]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$h[.,.]$$



$$h[m, n] = \sum_{k,l} f[k, l] I[m+k, n+l]$$

$$\begin{aligned} m &= 6, n &= 4 \\ k, l &= [-1, 0, 1] \end{aligned}$$

$$f[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

Credit: S. Seitz

# Image filtering

$$f[\cdot, \cdot] \quad \frac{1}{9} \begin{array}{|ccc|} \hline 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \\ \hline \end{array}$$

$$I[., .]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$h[., .]$$

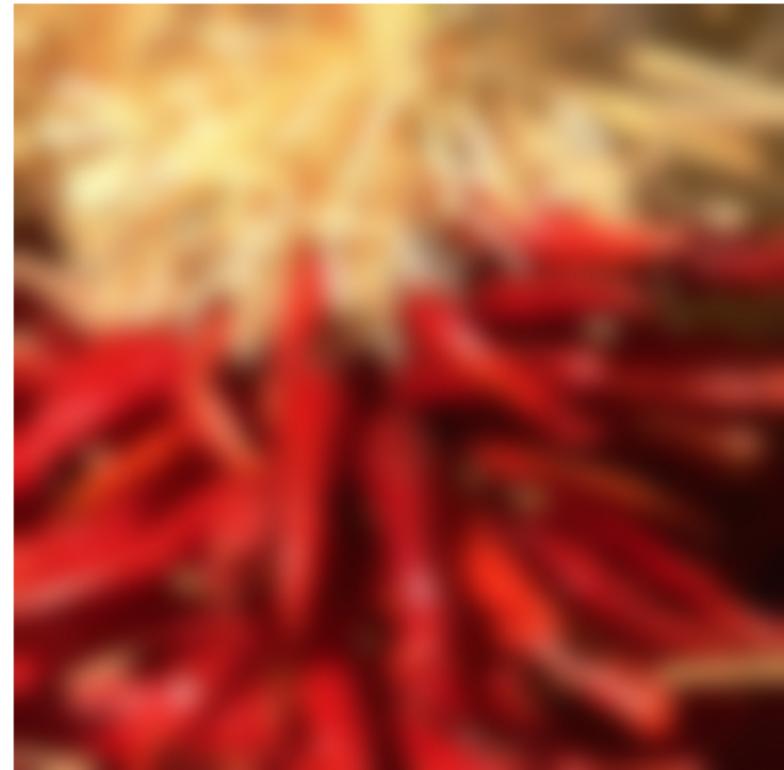
	0	10	20	30	30	30	20	10		
	0	20	40	60	60	60	40	20		
	0	30	60	90	90	90	60	30		
	0	30	50	80	80	90	60	30		
	0	30	50	80	80	90	60	30		
	0	20	30	50	50	60	40	20		
	10	20	30	30	30	30	20	10		
	10	10	10	0	0	0	0	0		

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

Credit: S. Seitz

# Boundary issues

- What about near the edge?
  - the filter window falls off the edge of the image
  - need to extrapolate
  - methods:
    - clip filter (black)
    - wrap around
    - copy edge
    - reflect across edge



# Properties of convolution

- Linear & shift invariant

- Commutative:

$$f * g = g * f$$

- Associative

$$(f * g) * h = f * (g * h)$$

- Identity:

unit impulse  $e = [..., 0, 0, 1, 0, 0, ...]$ .  $f * e = f$

- Differentiation:

$$\frac{\partial}{\partial x} (f * g) = \frac{\partial f}{\partial x} * g$$

# Separable filters

A 2D filter is separable if it can be written as the product of a “column” and a “row”.

example:  
box filter

1	1	1
1	1	1
1	1	1

=

1
1
1

\*

1	1	1
---	---	---

row

column

What is the rank of this filter matrix?

# Separable filters

A 2D filter is separable if it can be written as the product of a “column” and a “row”.

example:  
box filter

1	1	1
1	1	1
1	1	1

=

1
1
1

\*

1	1	1
---	---	---

row

column

Why is this important?

# Separable filters

A 2D filter is separable if it can be written as the product of a “column” and a “row”.

example:  
box filter

1	1	1
1	1	1
1	1	1

=

1
1
1

\*

1	1	1
---	---	---

row

column

2D convolution with a separable filter is equivalent to two 1D convolutions (with the “column” and “row” filters).

# Separable filters

A 2D filter is separable if it can be written as the product of a “column” and a “row”.

example:  
box filter

1	1	1
1	1	1
1	1	1

=

1
1
1

\*

1	1	1
---	---	---

row

column

2D convolution with a separable filter is equivalent to two 1D convolutions (with the “column” and “row” filters).

If the image has  $M \times M$  pixels and the filter kernel has size  $N \times N$ :

- What is the cost of convolution with a non-separable filter?

# Separable filters

A 2D filter is separable if it can be written as the product of a “column” and a “row”.

example:  
box filter

1	1	1
1	1	1
1	1	1

=

1
1
1

\*

1	1	1
---	---	---

row

column

2D convolution with a separable filter is equivalent to two 1D convolutions (with the “column” and “row” filters).

If the image has  $M \times M$  pixels and the filter kernel has size  $N \times N$ :

- What is the cost of convolution with a non-separable filter?  $\longrightarrow M^2 \times N^2$
- What is the cost of convolution with a separable filter?

# Separable filters

A 2D filter is separable if it can be written as the product of a “column” and a “row”.

example:  
box filter

1	1	1
1	1	1
1	1	1

=

1
1
1

\*

1	1	1
---	---	---

row

column

2D convolution with a separable filter is equivalent to two 1D convolutions (with the “column” and “row” filters).

If the image has  $M \times M$  pixels and the filter kernel has size  $N \times N$ :

- What is the cost of convolution with a non-separable filter?  $\longrightarrow M^2 \times N^2$
- What is the cost of convolution with a separable filter?  $\longrightarrow 2 \times N \times M^2$

# Think-Pair-Share time



1.

0	0	0
0	1	0
0	0	0

2.

0	0	0
0	0	1
0	0	0

3.

1	0	-1
2	0	-2
1	0	-1

4.

0	0	0
0	2	0
0	0	0

-

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

# 1. Practice with linear filters



Original

0	0	0
0	1	0
0	0	0

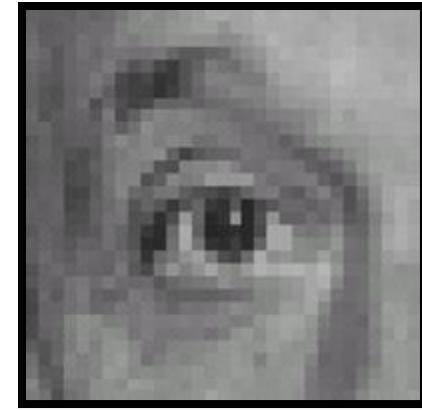
?

# 1. Practice with linear filters



Original

0	0	0
0	1	0
0	0	0



Filtered  
(no change)

## 2. Practice with linear filters



0	0	0
0	0	1
0	0	0

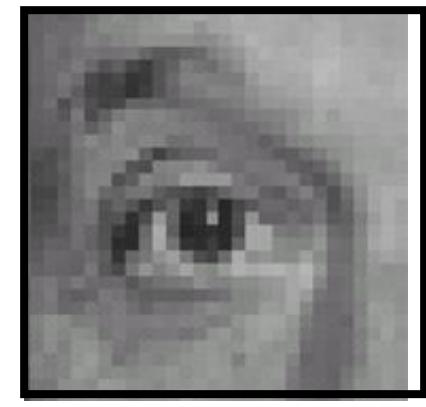
?

## 2. Practice with linear filters



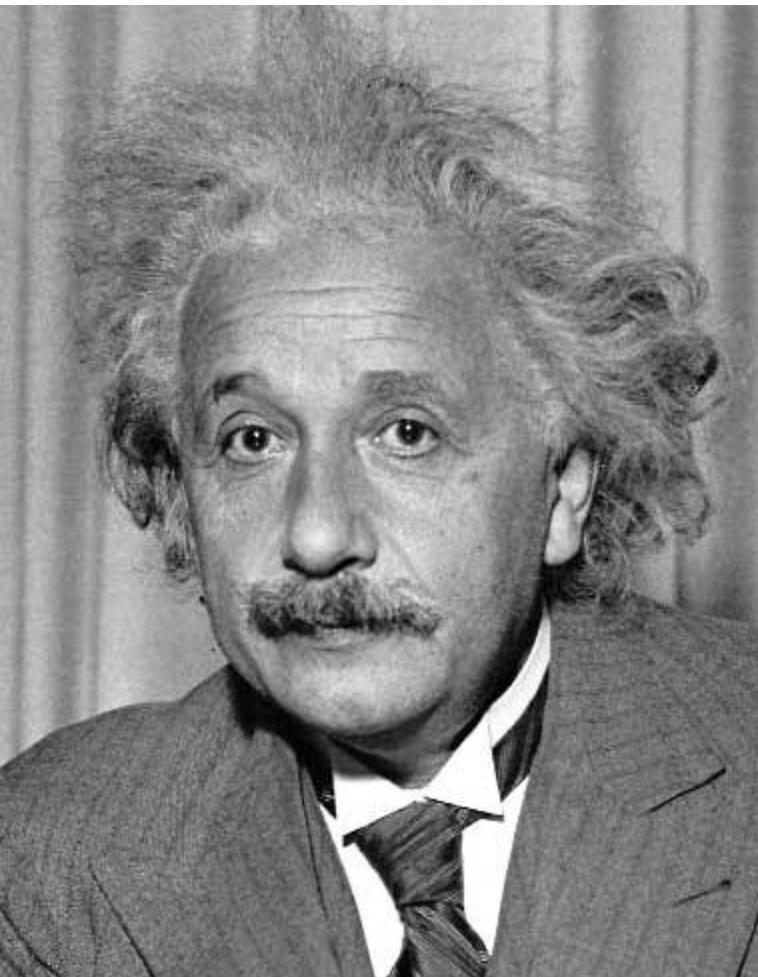
Original

0	0	0
0	0	1
0	0	0



Shifted left  
By 1 pixel

### 3. Practice with linear filters



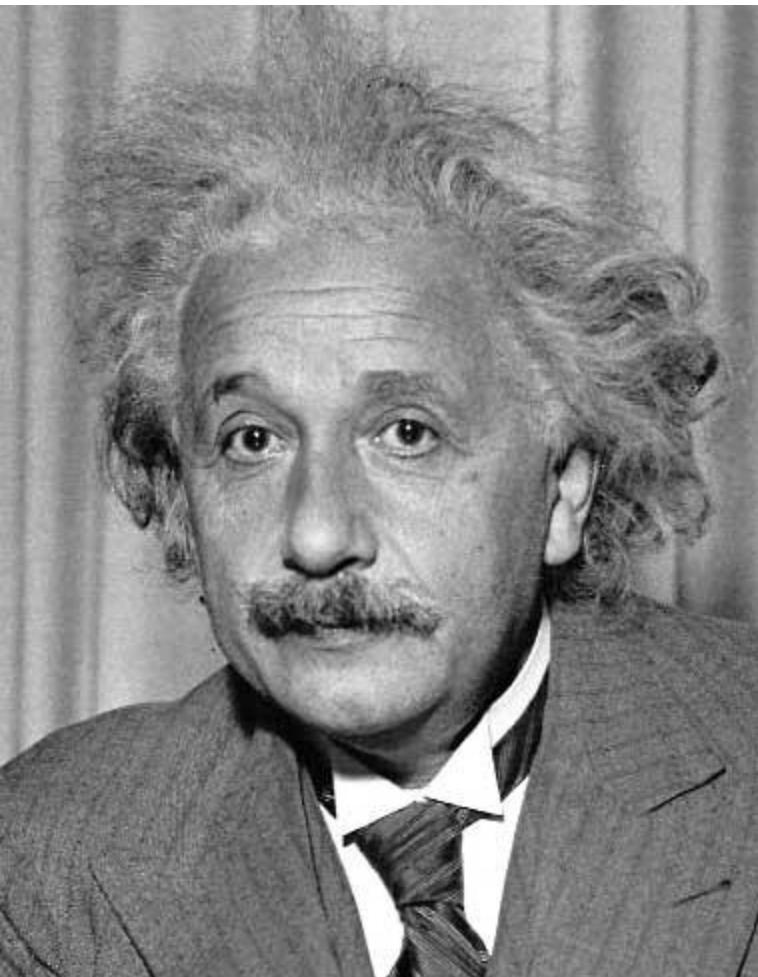
1	0	-1
2	0	-2
1	0	-1

Sobel



Vertical Edge  
(absolute value)

### 3. Practice with linear filters



1	2	1
0	0	0
-1	-2	-1

Sobel



Horizontal Edge  
(absolute value)

# 4. Practice with linear filters



0	0	0
0	2	0
0	0	0

-

$\frac{1}{9}$	1	1	1
1	1	1	1
1	1	1	1

?

(Note that filter sums to 1)

Original

# 4. Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

-

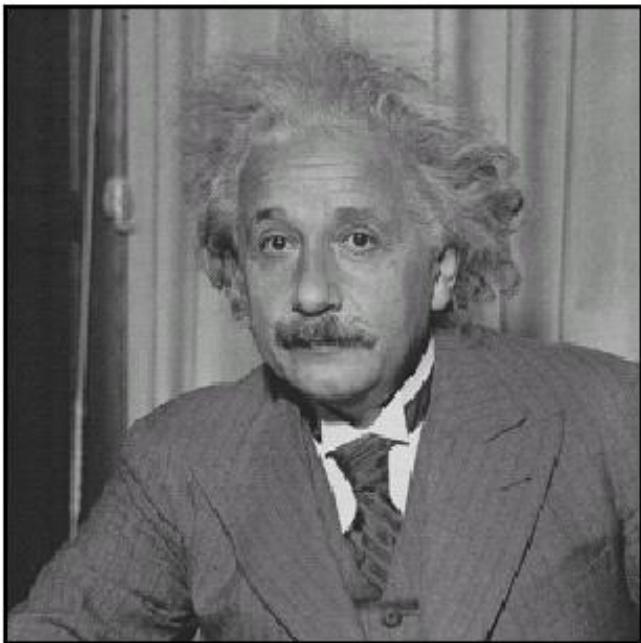
$\frac{1}{9}$	1	1	1
1	1	1	1
1	1	1	1



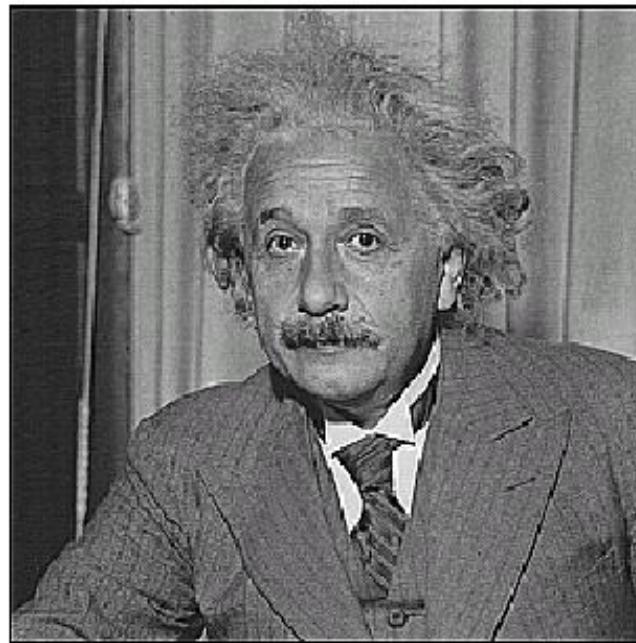
## Sharpening filter

- Accentuates differences with local average

# 4. Practice with linear filters



**before**



**after**

# Box Filter

What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove sharp features)

$$f[\cdot, \cdot]$$

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

# Box Filter

What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove sharp features)
- Why does it sum to one?

$$f[\cdot, \cdot]$$

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

# What does blurring remove?



Now let's add it back.....

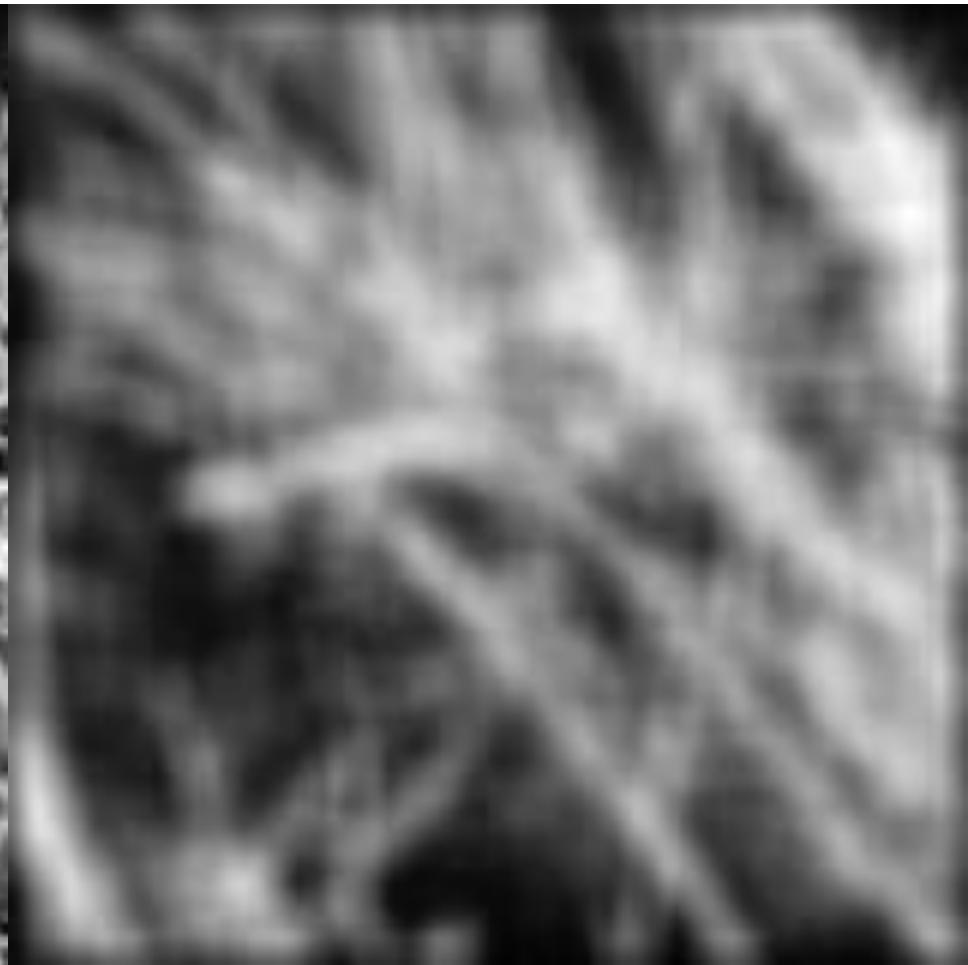


# Smoothing with box filter

$$f[\cdot, \cdot]$$

$$\frac{1}{9}$$

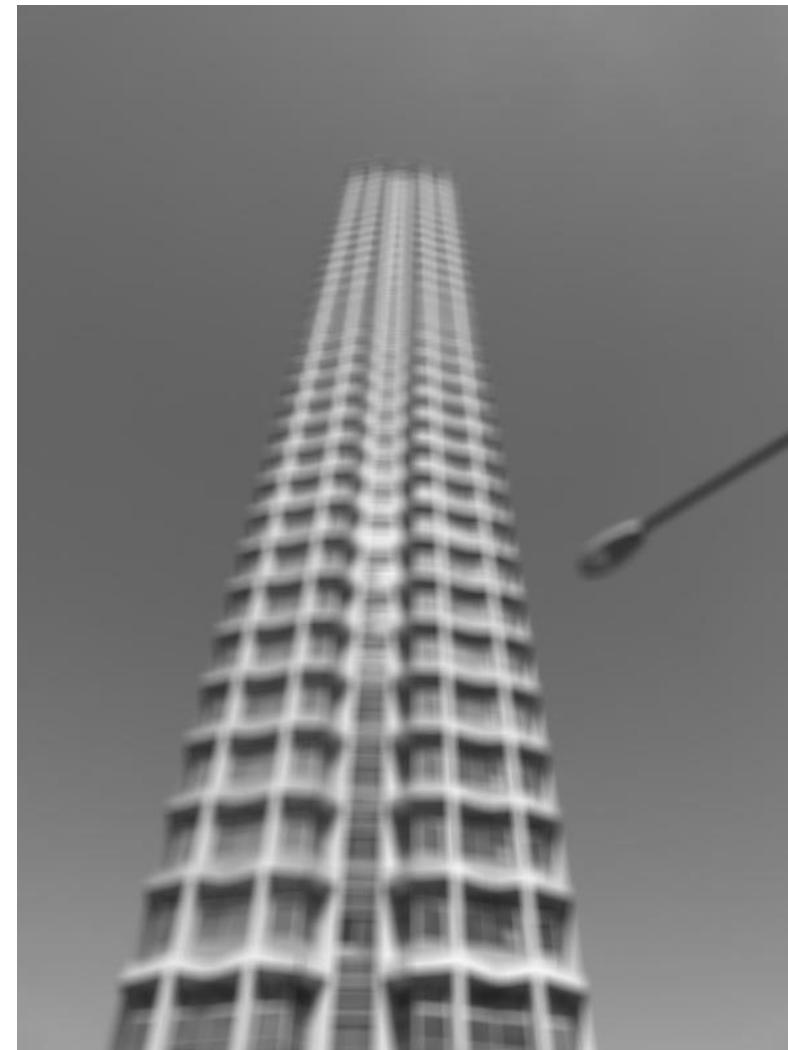
1	1	1
1	1	1
1	1	1



# A few more filters



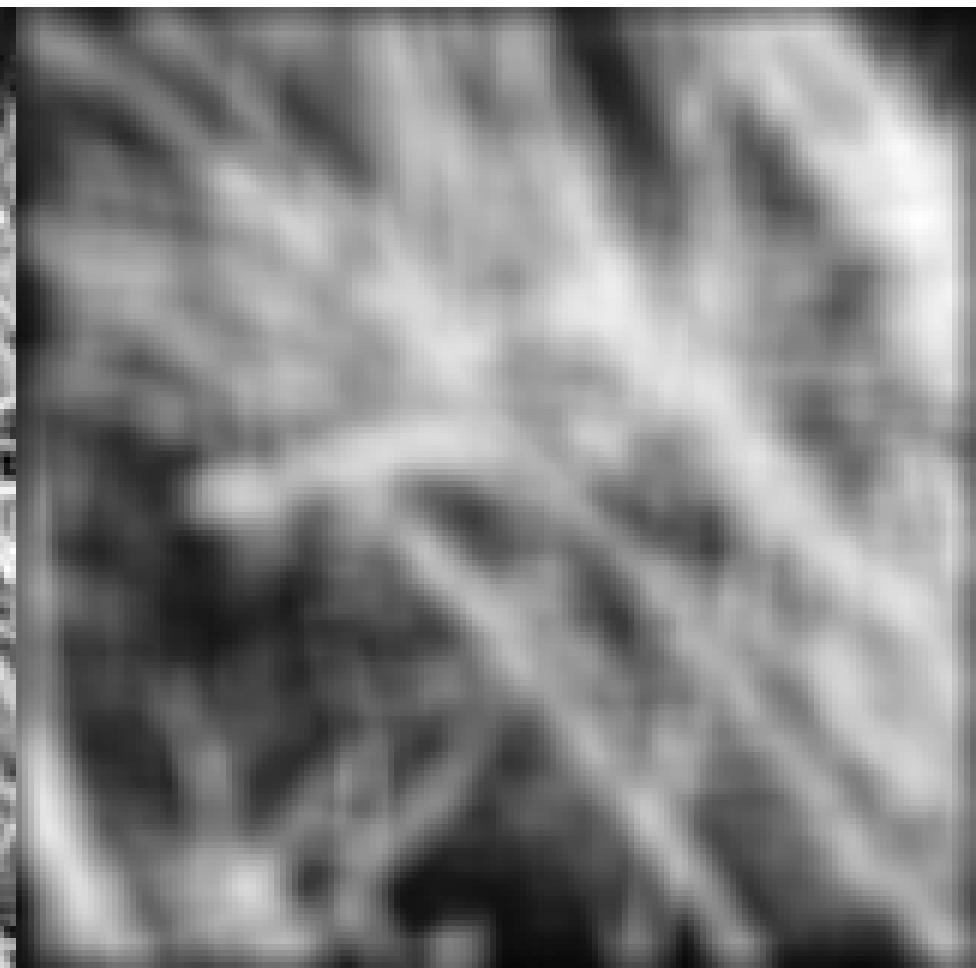
original



3x3 box filter

do you see  
any problems  
in this image?

# Example: Smoothing by Averaging

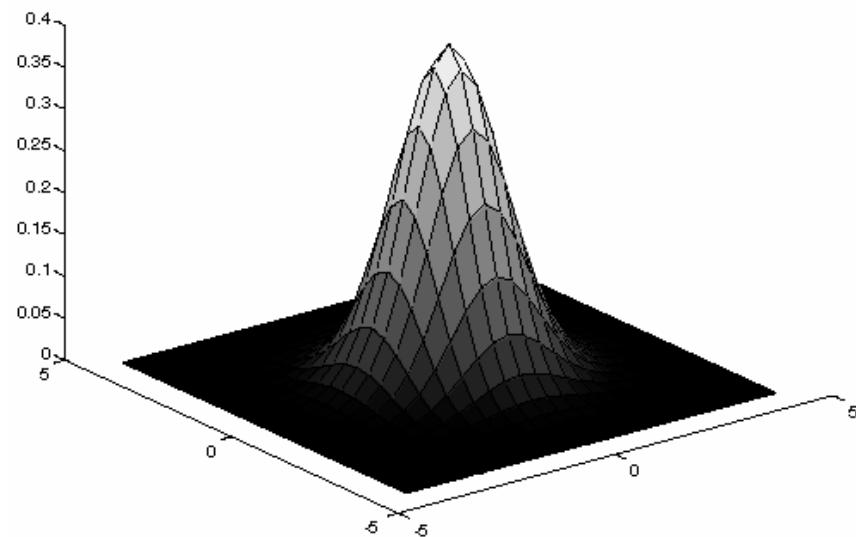


# Gaussian filters

- Remove “high-frequency” components from the image (a low-pass filter)
  - Images become more smooth
- Gaussian convolved with Gaussian...  
...is another Gaussian
  - So can smooth with small-width kernel, repeat, and get same result as larger-width kernel would have
  - Convoluting twice with Gaussian kernel of width  $\sigma$  is same as convoluting once with kernel of width  $\sigma\sqrt{2}$
- *Separable* kernel
  - Factors into product of two 1D Gaussians

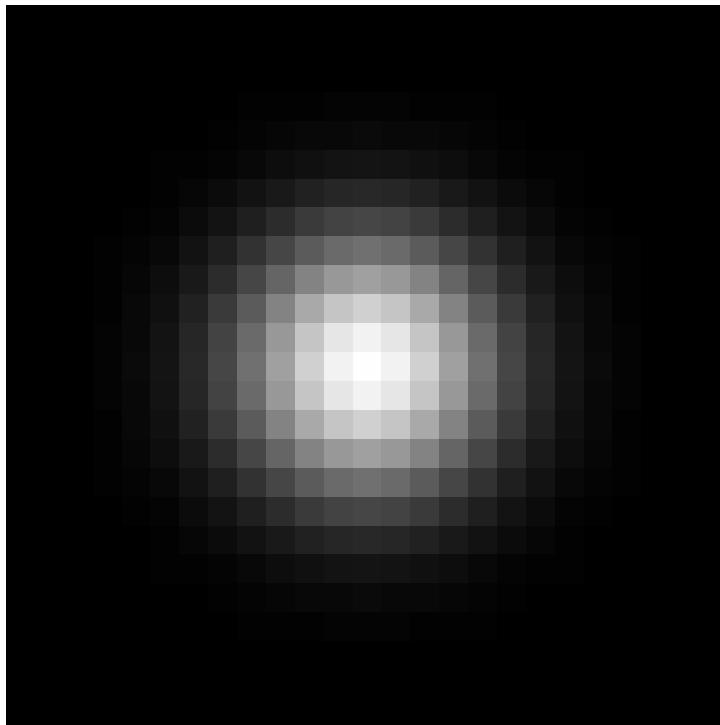
# Gaussian Averaging

- Rotationally symmetric.
- Weights nearby pixels more than distant ones.
  - This makes sense as probabilistic inference.



- A Gaussian gives a good model of a fuzzy blob

# An Isotropic Gaussian



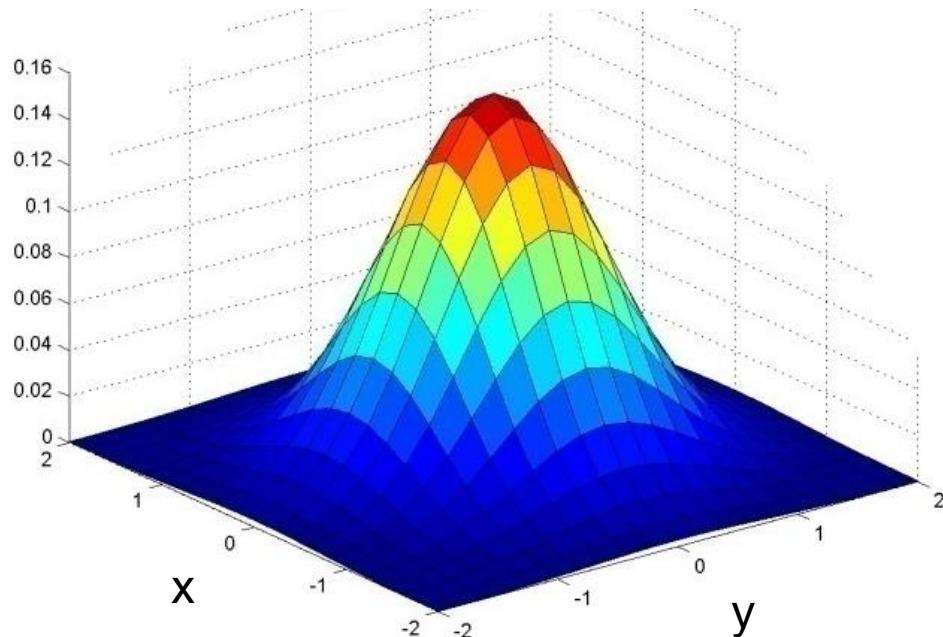
- The picture shows a smoothing kernel proportional to

$$\exp\left(-\left(\frac{x^2 + y^2}{2\sigma^2}\right)\right)$$

(which is a reasonable model of a circularly symmetric fuzzy blob)

# Important filter: Gaussian

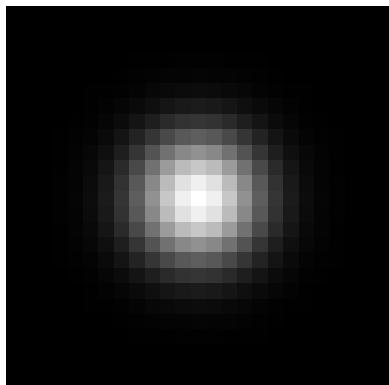
Weight contributions of neighboring pixels by nearness



x	0.003	0.013	0.022	0.013	0.003
y	0.013	0.059	0.097	0.059	0.013
x	0.022	0.097	0.159	0.097	0.022
y	0.013	0.059	0.097	0.059	0.013
x	0.003	0.013	0.022	0.013	0.003

Kernel size  $5 \times 5$ ,  
Standard deviation  $\sigma = 1$

Viewed  
from top



$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

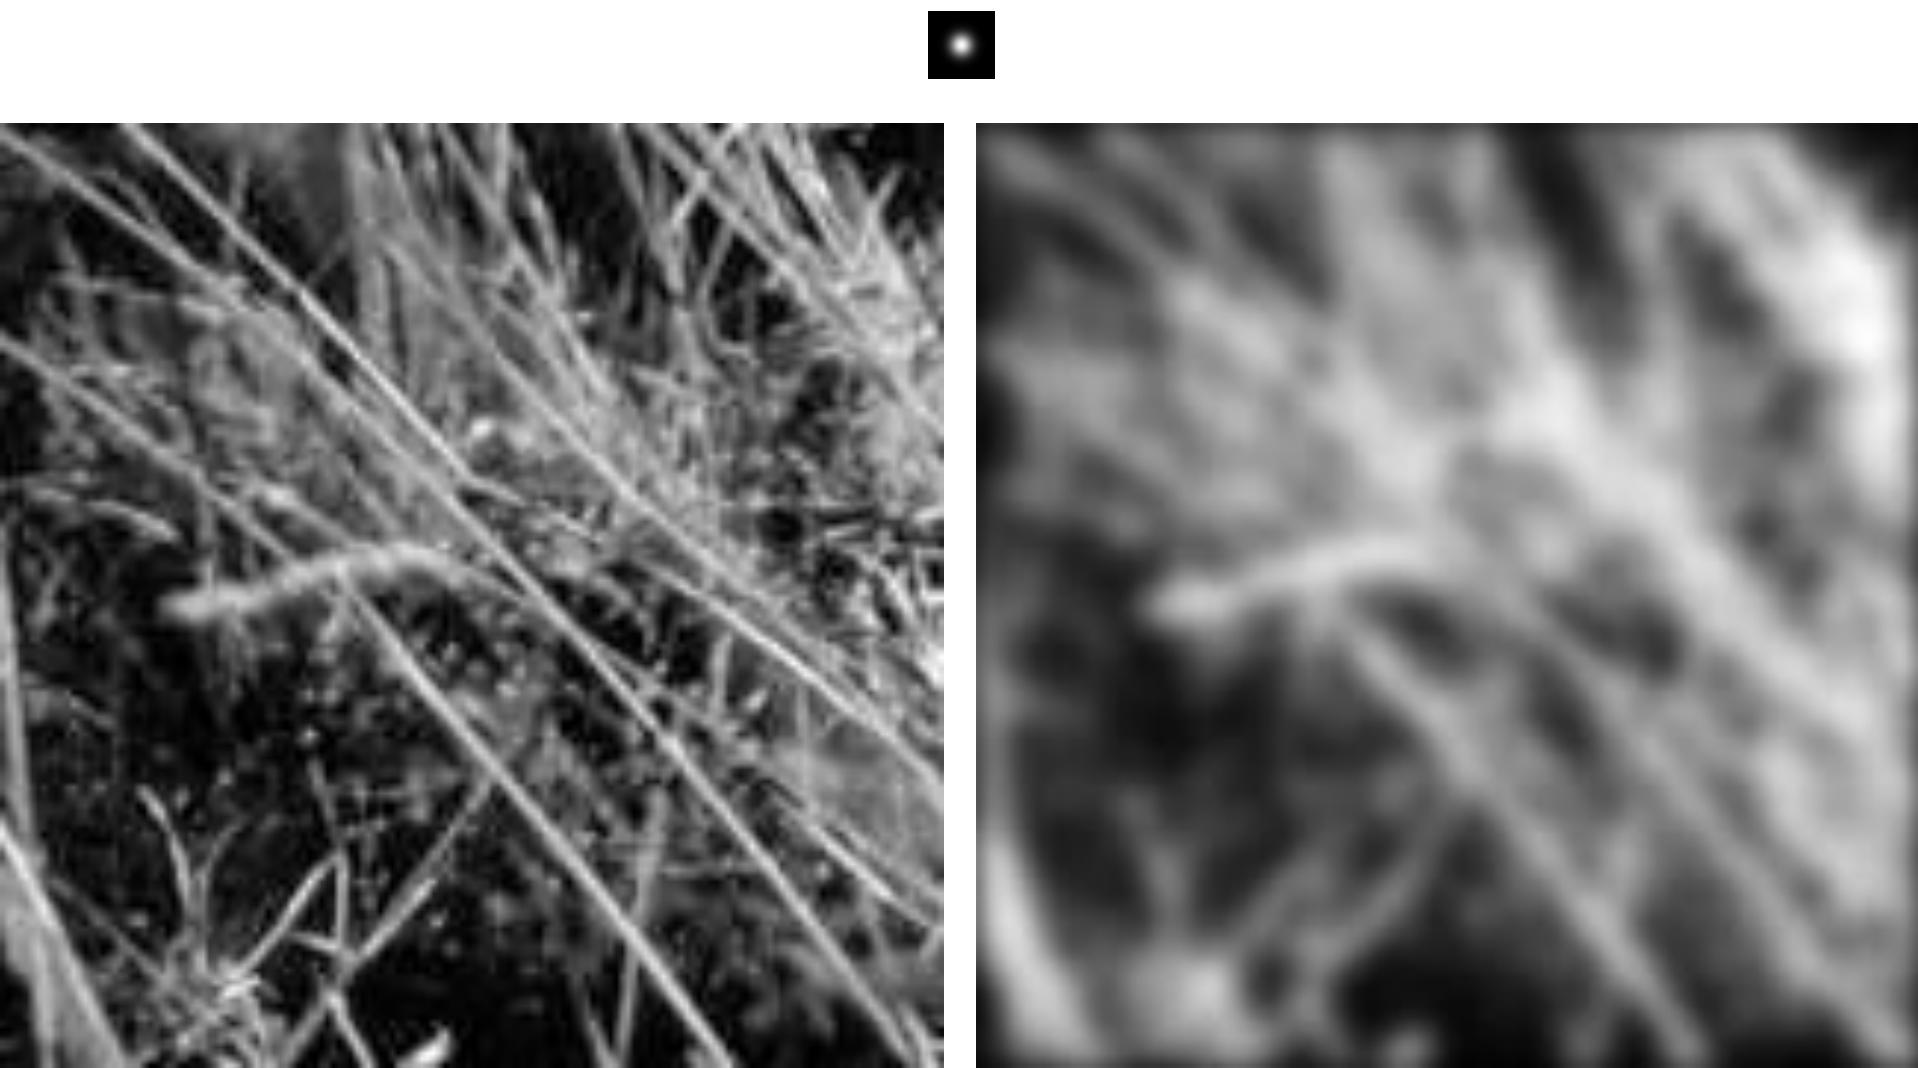
# Separability of the Gaussian filter

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

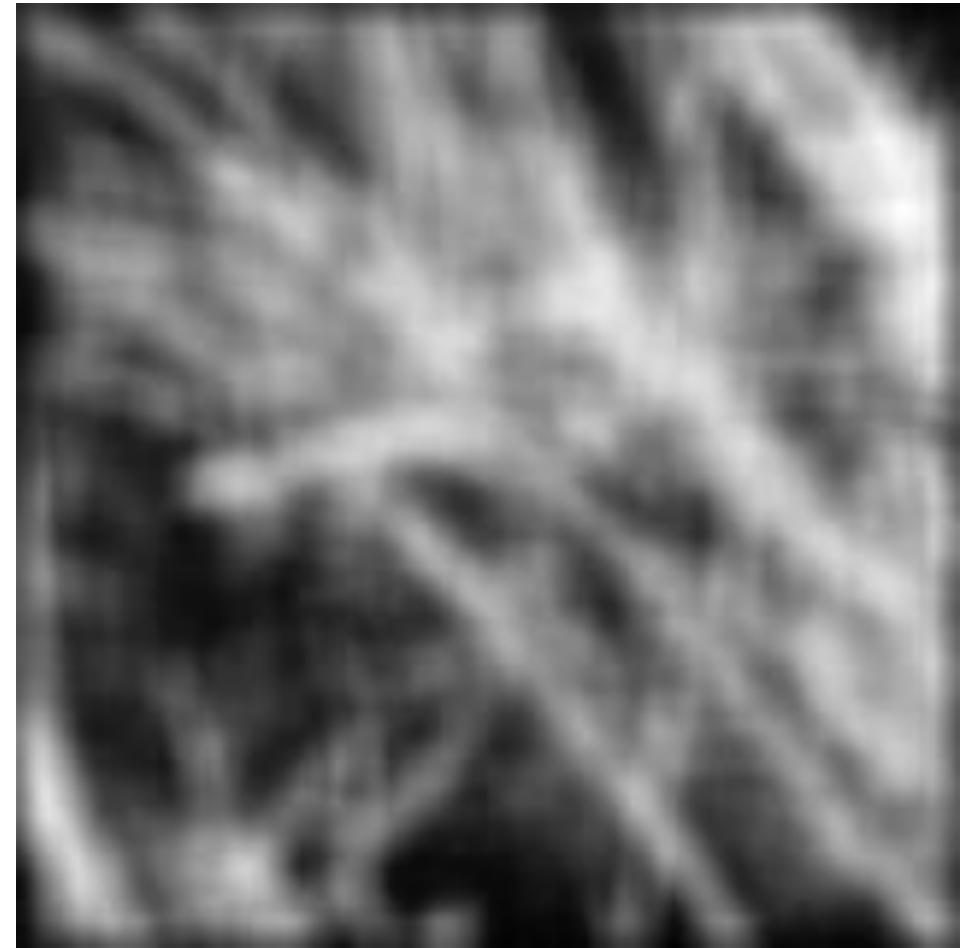
The 2D Gaussian can be expressed as the product of two functions, one a function of  $x$  and the other a function of  $y$

In this case, the two functions are the (identical) 1D Gaussian

# Smoothing with Gaussian filter



# Smoothing with box filter

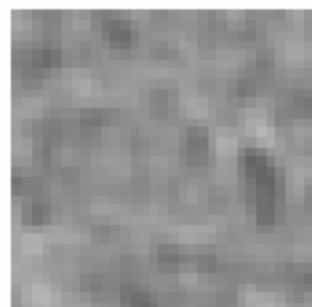
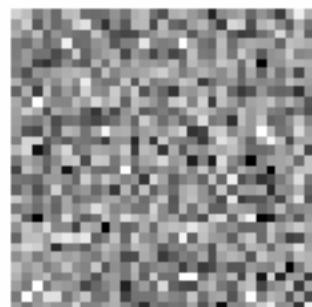
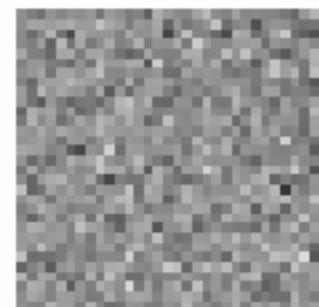


$\sigma=0.05$

$\sigma=0.1$

$\sigma=0.2$

no  
smoothing



$\sigma=1$  pixel

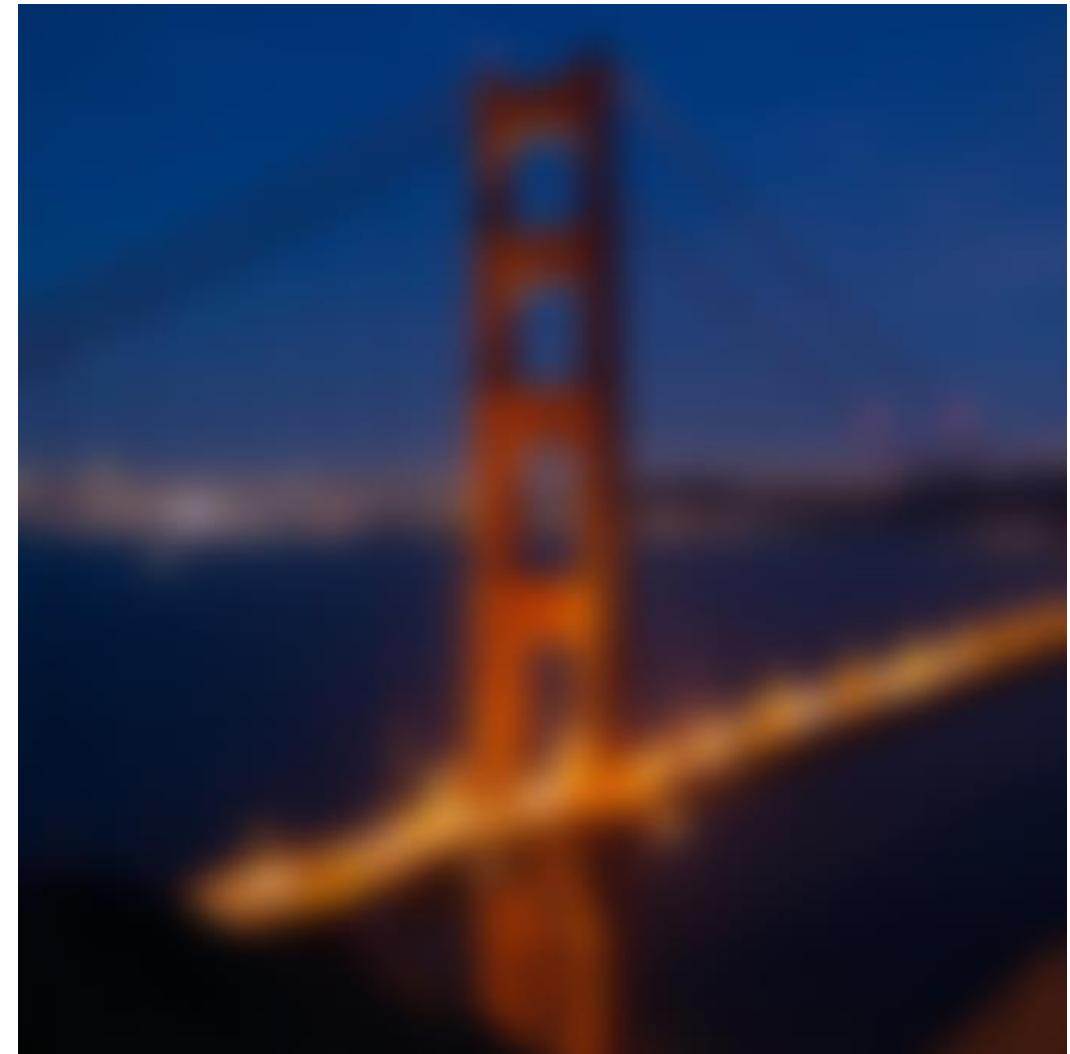


$\sigma=2$  pixels

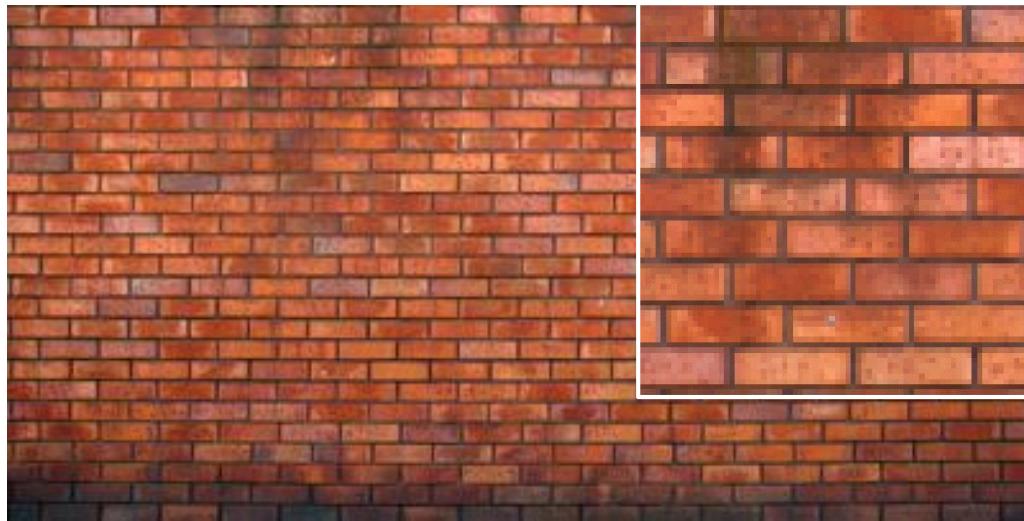
## The effects of smoothing

Each row shows smoothing with gaussians of different width; each column shows different realizations of an image of gaussian noise.

# Gaussian filtering example

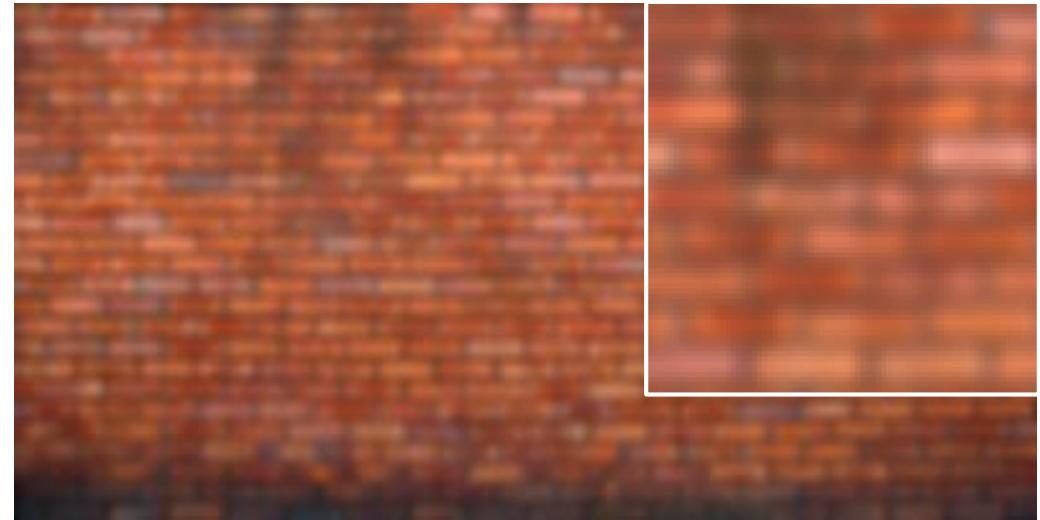


# Gaussian vs box filtering

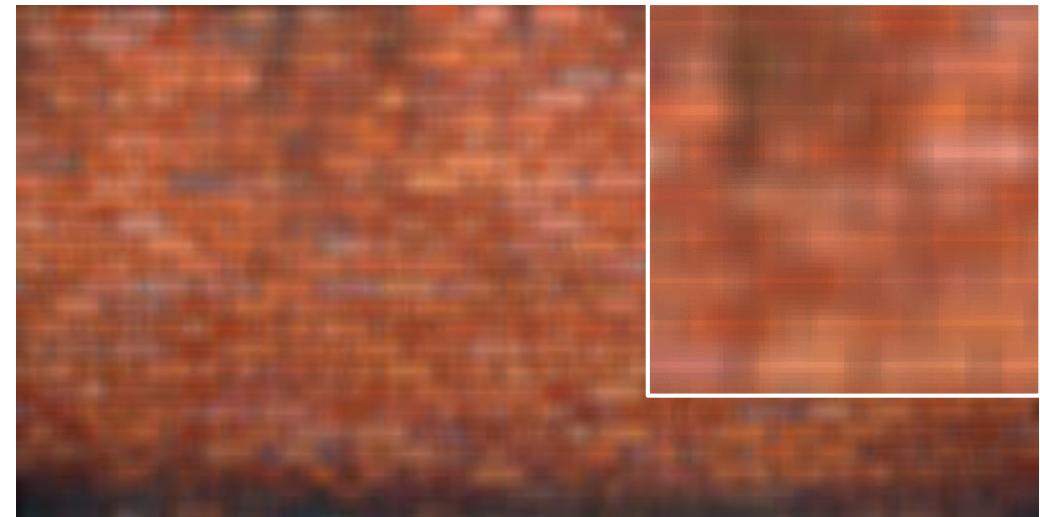


original

Which blur do you like better?



7x7 Gaussian



7x7 box

# Edge Filters

- Sobel

$$\frac{1}{8} \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix} \quad \frac{1}{8} \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

- Prewitt

$$\frac{1}{6} \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix} \quad \frac{1}{6} \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$

- Roberts

$$\begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix} \quad \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$$

- Laplacian

$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix} \quad \frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \\ 1 & -8 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

# The Sobel filter

1	0	-1
2	0	-2
1	0	-1

Sobel filter

=

1
2
1

\*

1	0	-1
---	---	----

1D derivative  
filter

What filter  
is this?

# The Sobel filter

1	0	-1
2	0	-2
1	0	-1

Sobel filter

=

1
2
1

Blurring

\*

1	0	-1
---	---	----

1D derivative  
filter

In a 2D image, does this filter responses along horizontal or vertical lines?

# The Sobel filter

1	0	-1
2	0	-2
1	0	-1

Sobel filter

=

1
2
1

Blurring

\*

1	0	-1
---	---	----

1D derivative  
filter

Does this filter return large responses on vertical or horizontal lines?

# The Sobel filter

Horizontal Sober filter:

$$\begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline \end{array}$$

What does the vertical Sobel filter look like?

# The Sobel filter

Horizontal Sober filter:

$$\begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline \end{array}$$

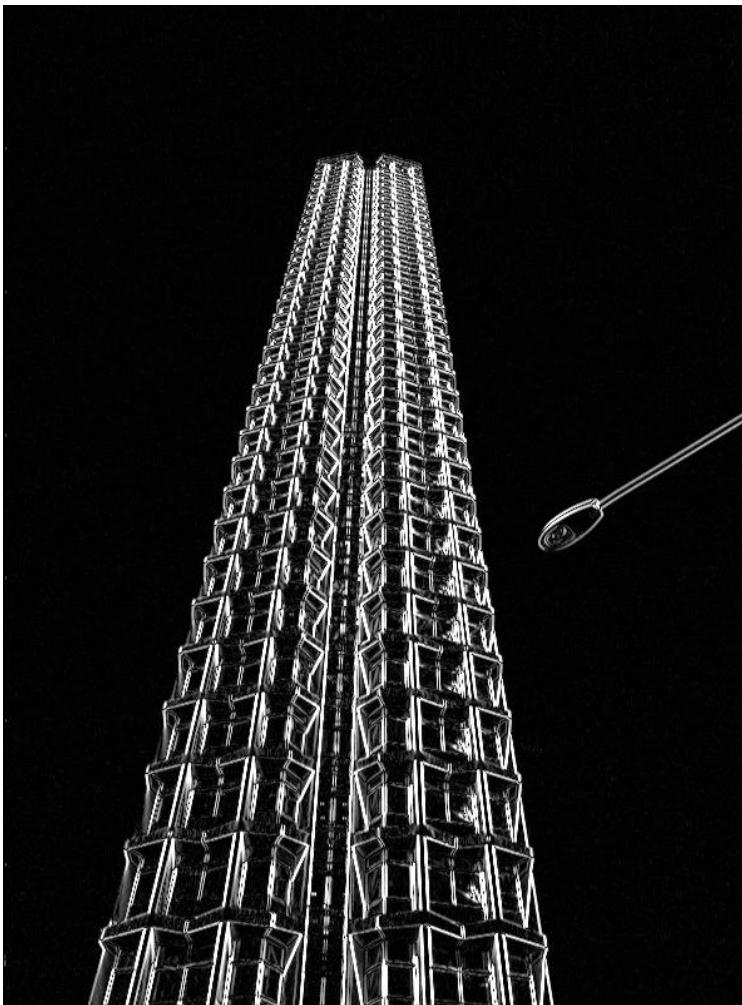
Vertical Sobel filter:

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 0 \\ \hline -1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array}$$

# Sobel filter example



original



which Sobel filter?

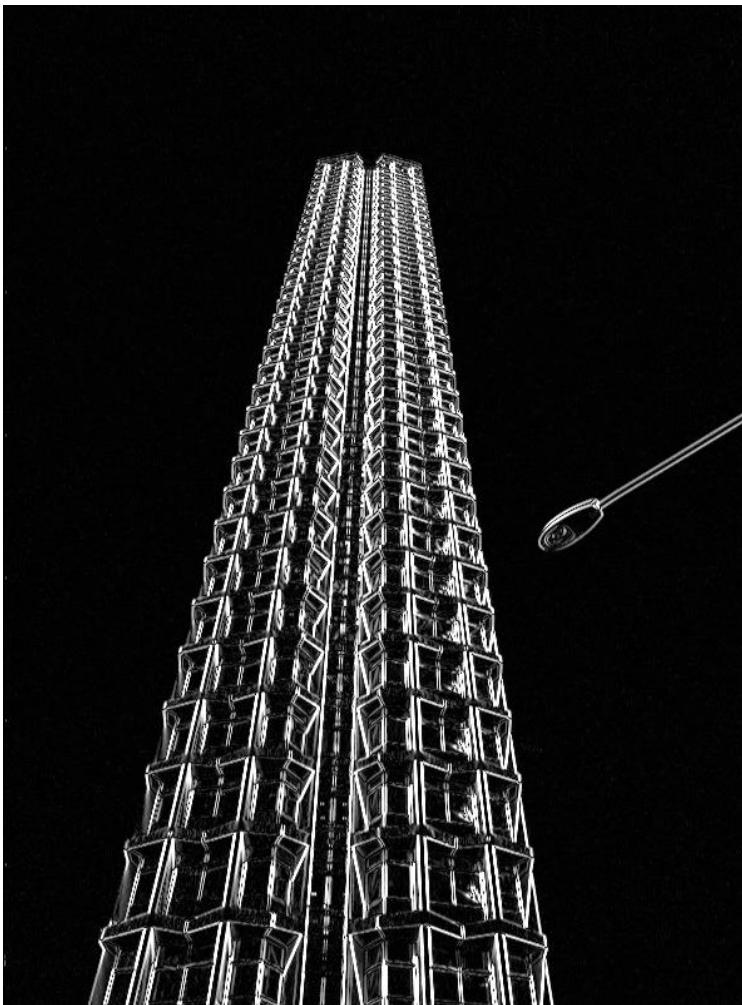


which Sobel filter?

# Sobel filter example



original

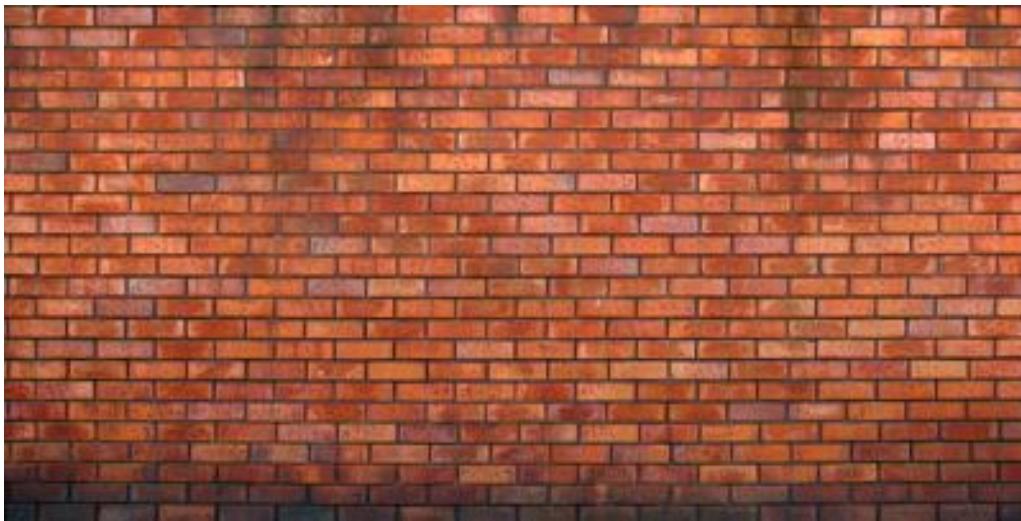


horizontal Sobel filter

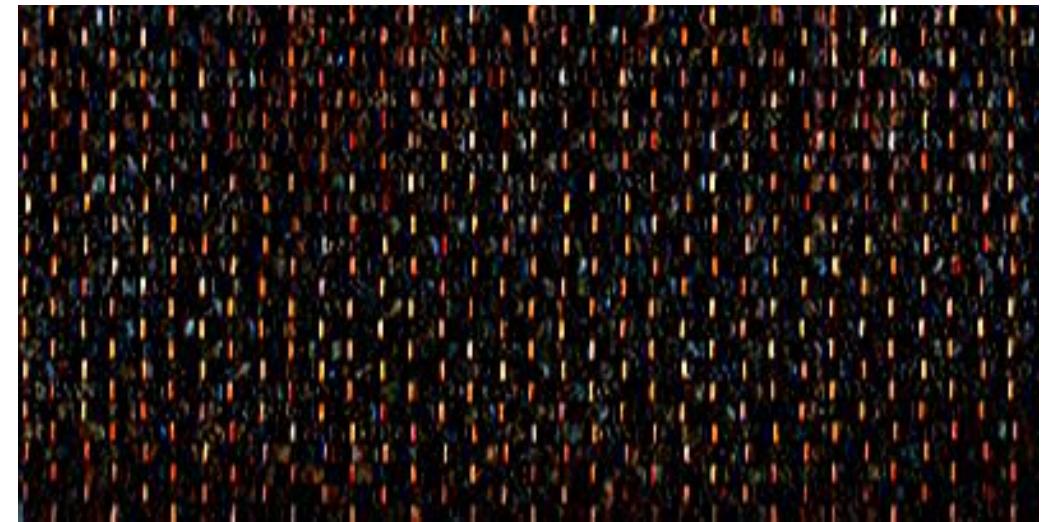


vertical Sobel filter

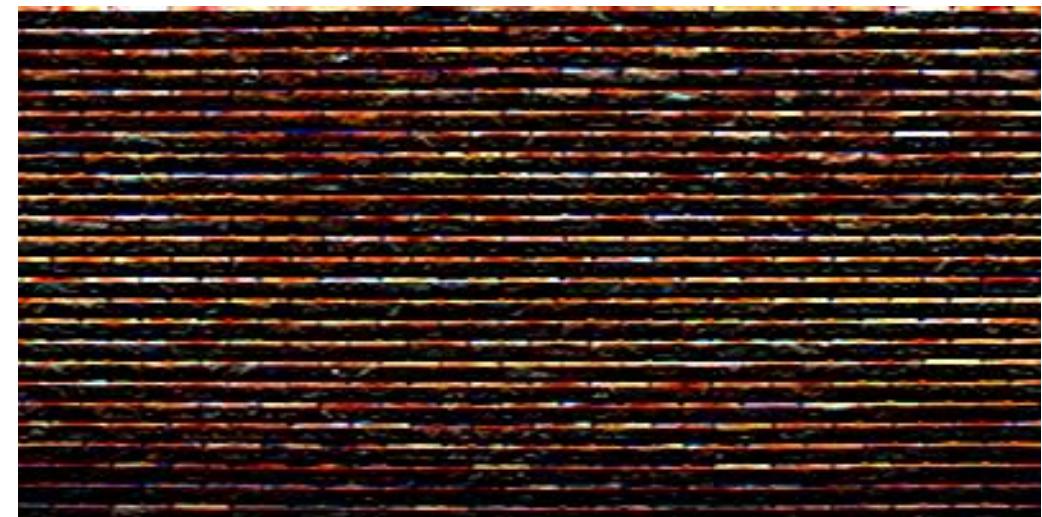
# Sobel filter example



original



horizontal Sobel filter



vertical Sobel filter

# Computing image gradients

1. Select your favorite derivative filters.

$$\mathbf{S}_x = \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array}$$

$$\mathbf{S}_y = \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

# Computing image gradients

1. Select your favorite derivative filters.

$$\mathbf{S}_x = \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array} \quad \mathbf{S}_y = \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

2. Convolve with the image to compute derivatives.

$$\frac{\partial \mathbf{f}}{\partial x} = \mathbf{S}_x \otimes \mathbf{f}$$

$$\frac{\partial \mathbf{f}}{\partial y} = \mathbf{S}_y \otimes \mathbf{f}$$

# Computing image gradients

1. Select your favorite derivative filters.

$$\mathbf{S}_x = \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array}$$

$$\mathbf{S}_y = \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

2. Convolve with the image to compute derivatives.

$$\frac{\partial \mathbf{f}}{\partial x} = \mathbf{S}_x \otimes \mathbf{f}$$

$$\frac{\partial \mathbf{f}}{\partial y} = \mathbf{S}_y \otimes \mathbf{f}$$

3. Form the image gradient, and compute its direction and amplitude.

$$\nabla \mathbf{f} = \left[ \frac{\partial \mathbf{f}}{\partial x}, \frac{\partial \mathbf{f}}{\partial y} \right]$$

gradient

$$\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

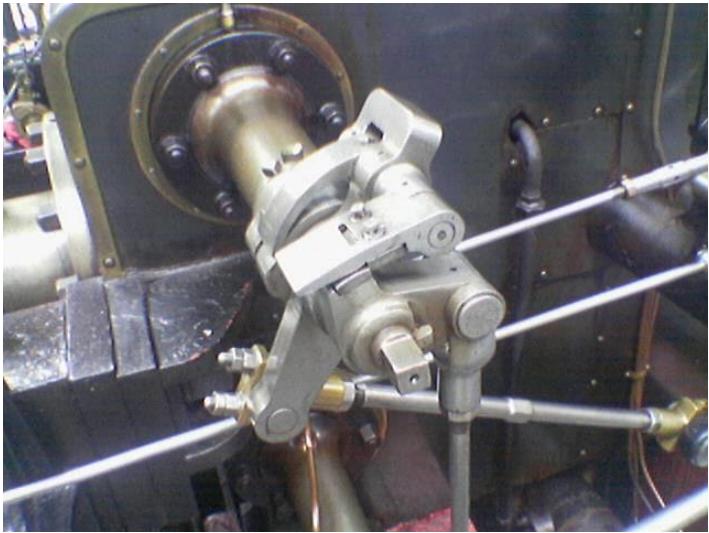
direction

$$\|\nabla f\| = \sqrt{\left( \frac{\partial f}{\partial x} \right)^2 + \left( \frac{\partial f}{\partial y} \right)^2}$$

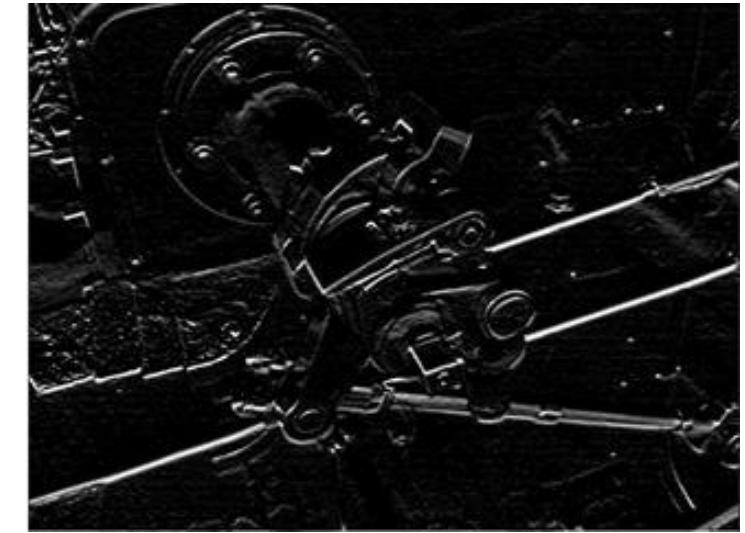
amplitude

# Image gradient example

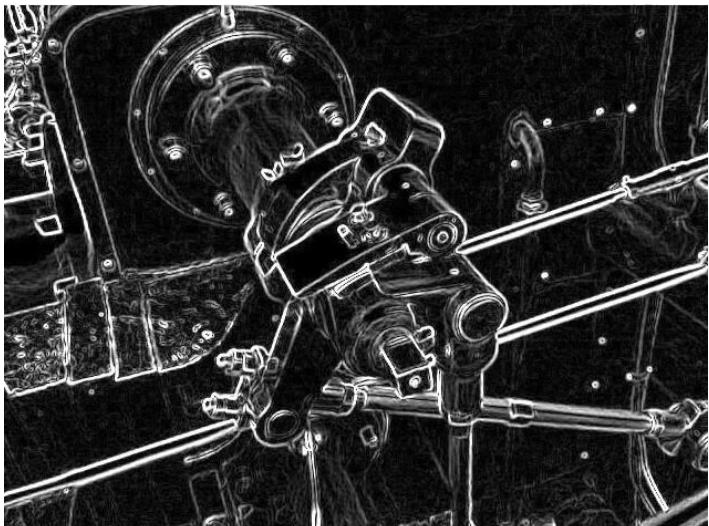
original



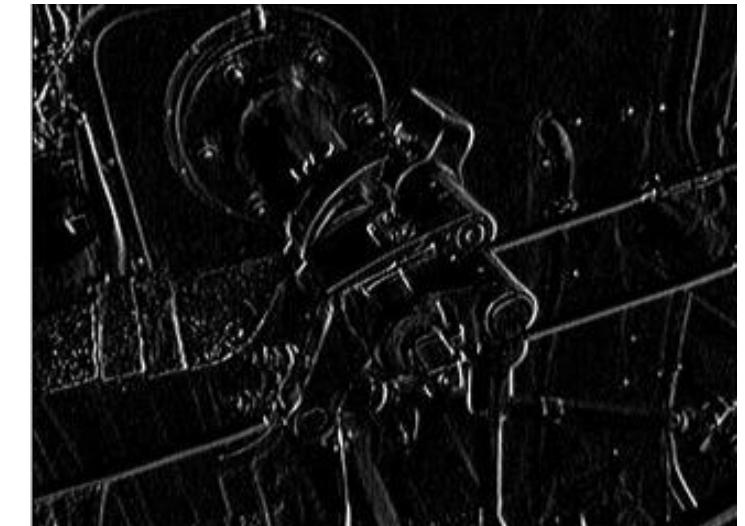
vertical derivative



gradient amplitude

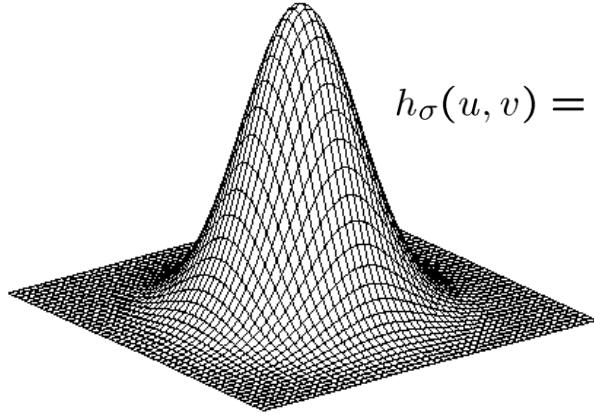


horizontal derivative



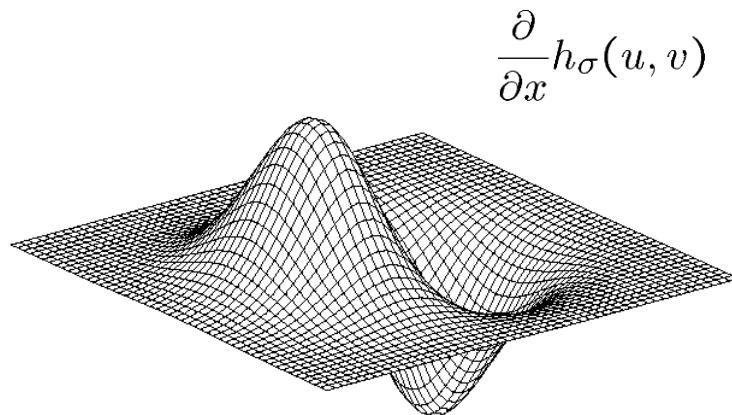
How does the gradient direction relate to these edges?

# 2D Gaussian filters



Gaussian

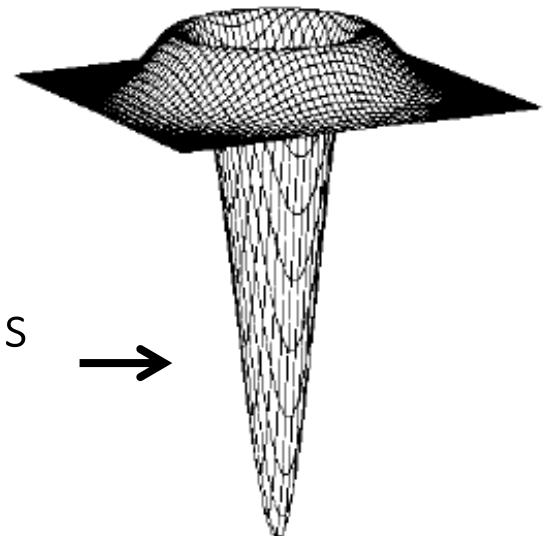
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



Derivative of Gaussian

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

$$\nabla^2 h_\sigma(u, v)$$



Laplacian of Gaussian

how does this relate to this  
lecture's cover picture?

# Laplace and LoG filtering examples



Laplacian of Gaussian filtering



Laplace filtering

# Laplacian of Gaussian vs Derivative of Gaussian

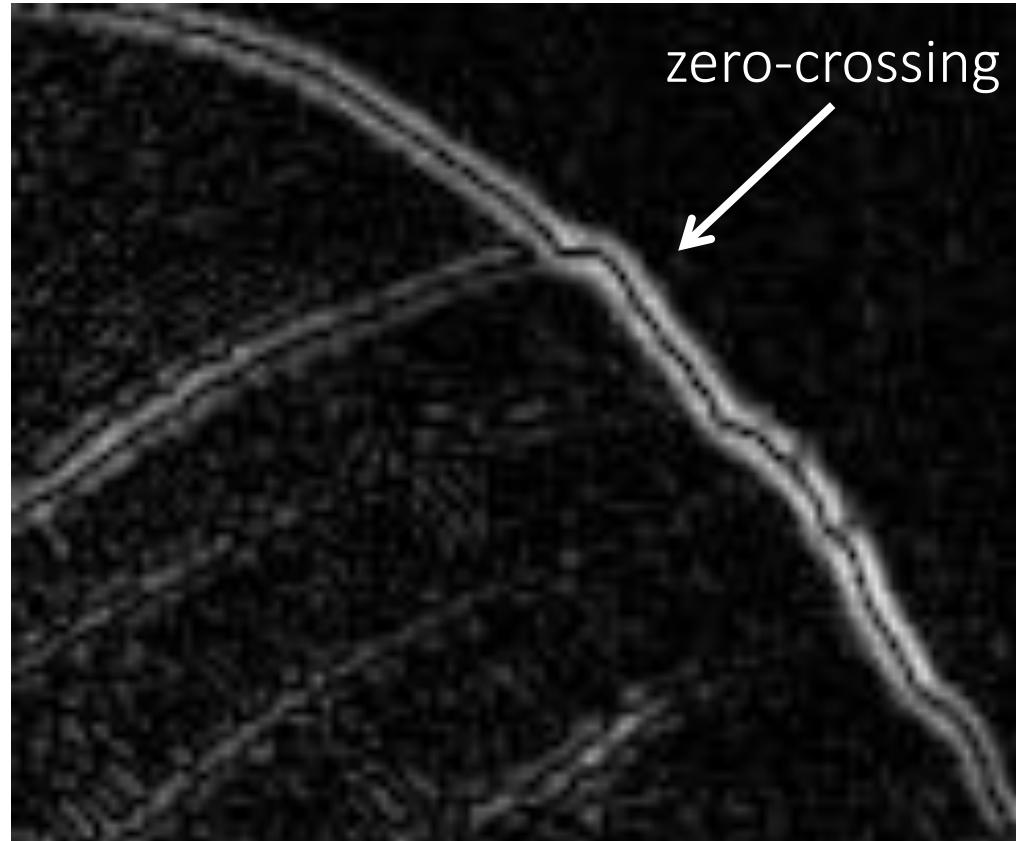


Laplacian of Gaussian filtering

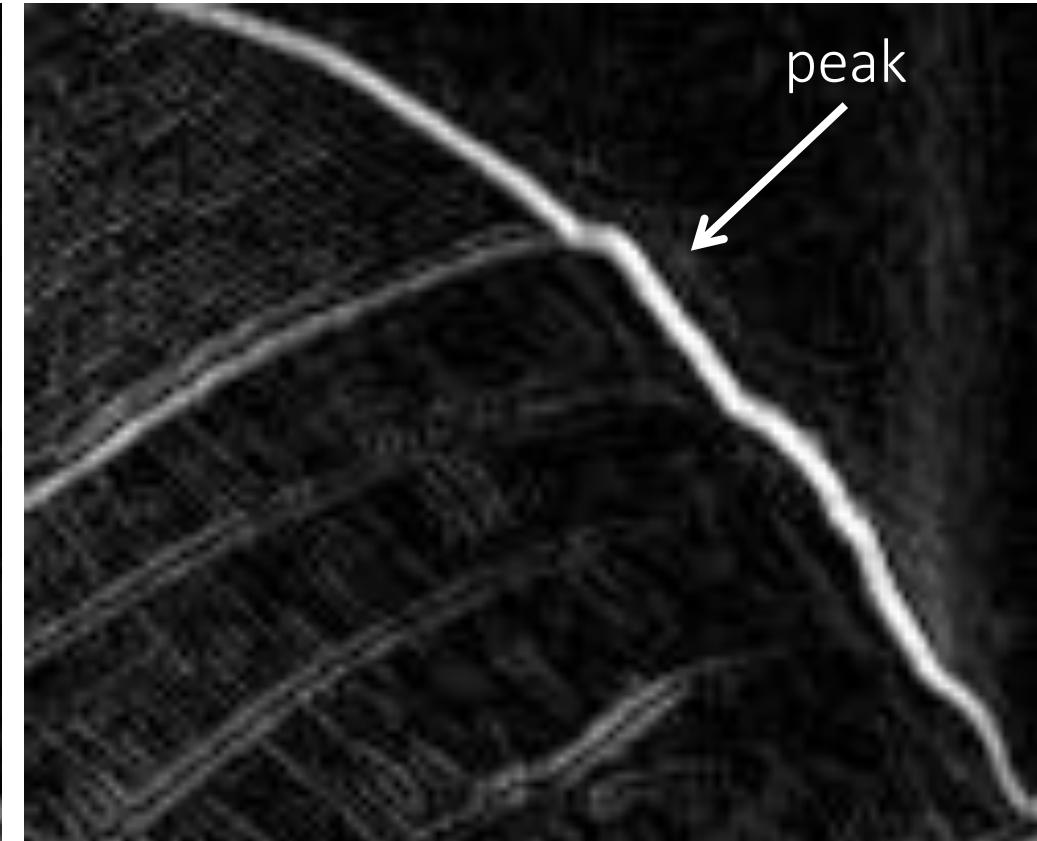


Derivative of Gaussian filtering

# Laplacian of Gaussian vs Derivative of Gaussian



Laplacian of Gaussian filtering



Derivative of Gaussian filtering

Zero crossings are more accurate at localizing edges (but not very convenient).